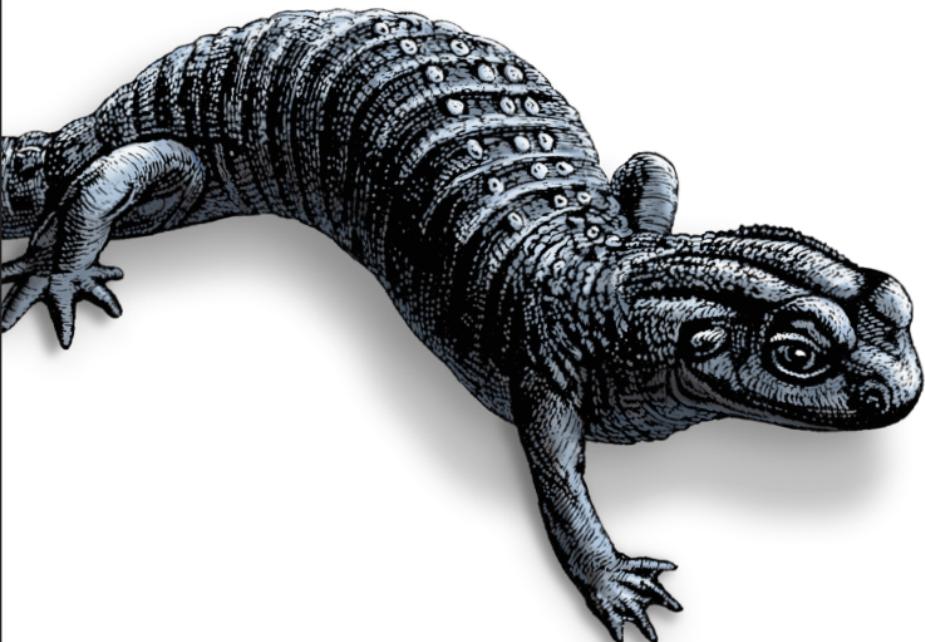


O'REILLY®

4th Edition

SQL Pocket Guide

A Guide to SQL Usage



Alice Zhao

SQL Pocket Guide

If you use SQL in your day-to-day work as a data analyst, data scientist, or data engineer, this popular pocket guide is your ideal on-the-job reference. You'll find many examples that address the language's complexities, along with key aspects of SQL used in Microsoft SQL Server, MySQL, Oracle Database, PostgreSQL, and SQLite.

In this updated edition, author Alice Zhao describes how these database management systems implement SQL syntax for both querying and making changes to a database. You'll find details on data types and conversions, regular expression syntax, window functions, pivoting and unpivoting, and more.

- Quickly look up how to perform specific tasks using SQL
- Apply the book's syntax examples to your own queries
- Update SQL queries to work in five different database management systems
- New: Connect Python and R to a relational database
- New: Look up frequently asked SQL questions in the "How Do I?" chapter

DATA

Twitter: @oreillymedia
facebook.com/oreilly

US \$29.99

CAN \$39.99

ISBN: 978-1-492-09040-3



9 781492 090403

FOURTH EDITION

SQL Pocket Guide

Alice Zhao

Beijing • Boston • Farnham • Sebastopol • Tokyo

O'REILLY®

SQL Pocket Guide

by Alice Zhao

Copyright © 2021 Alice Zhao. All rights reserved.

Printed in the United States of America.

Published by O'Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472.

O'Reilly books may be purchased for educational, business, or sales promotional use. Online editions are also available for most titles (<http://oreilly.com>). For more information, contact our corporate/institutional sales department: 800-998-9938 or corporate@oreilly.com.

Acquisitions Editor: Andy Kwan

Development Editor: Amelia Blevins and Jeff Bleiel

Production Editor: Caitlin Ghegan

Copyeditor: Piper Editorial Consulting, LLC

Proofreader: James Fraleigh

Indexer: Ellen Troutman-Zaig

Interior Designer: David Futato

Cover Designer: Karen Montgomery

Illustrator: Kate Dullea

September 2021: Fourth Edition

Revision History for the Fourth Edition

2021-08-26: First Release

See <https://oreil.ly/sqlpocketerrata> for release details.

The O'Reilly logo is a registered trademark of O'Reilly Media, Inc. *SQL Pocket Guide*, the cover image, and related trade dress are trademarks of O'Reilly Media, Inc.

The views expressed in this work are those of the author, and do not represent the publisher's views. While the publisher and the author have used good faith efforts to ensure that the information and instructions contained in this work are accurate, the publisher and the author disclaim all responsibility for errors or omissions, including without limitation responsibility for damages resulting from the use of or reliance on this work. Use of the information and instructions contained in this work is at your own risk. If any code samples or other technology this work contains or describes is subject to open source licenses or the intellectual property rights of others, it is your responsibility to ensure that your use thereof complies with such licenses and/or rights.

978-1-492-09040-3

[LSI]

Table of Contents

Preface	xi
Chapter 1: SQL Crash Course	1
What Is a Database?	1
SQL	1
NoSQL	2
Database Management Systems (DBMS)	3
A SQL Query	6
The SELECT Statement	7
Order of Execution	9
A Data Model	10
Chapter 2: Where Can I Write SQL Code?	13
RDBMS Software	14
SQLite	15
MySQL	17
Oracle	17
PostgreSQL	18
SQL Server	19

Database Tools	20
Connect a Database Tool to a Database	22
Other Programming Languages	24
Connect Python to a Database	25
Connect R to a Database	31
Chapter 3: The SQL Language	37
Comparison to Other Languages	37
ANSI Standards	39
SQL Terms	41
Keywords and Functions	42
Identifiers and Aliases	43
Statements and Clauses	45
Expressions and Predicates	47
Comments, Quotes, and Whitespace	48
Sublanguages	50
Chapter 4: Querying Basics	53
The SELECT Clause	55
Aliasing Columns	57
Qualifying Columns	59
Selecting Subqueries	61
DISTINCT	63
The FROM Clause	66
From Multiple Tables	66
From Subqueries	69
The WHERE Clause	73
Filtering on Subqueries	75
The GROUP BY Clause	78
The HAVING Clause	83
The ORDER BY Clause	85

The LIMIT Clause	88
Chapter 5: Creating, Updating, and Deleting	91
Databases	91
Display Names of Existing Databases	93
Display Name of Current Database	94
Switch to Another Database	95
Create a Database	95
Delete a Database	96
Creating Tables	97
Create a Simple Table	98
Display Names of Existing Tables	100
Create a Table That Does Not Already Exist	100
Create a Table with Constraints	101
Create a Table with Primary and Foreign Keys	105
Create a Table with an Automatically Generated Field	108
Insert the Results of a Query into a Table	110
Insert Data from a Text File into a Table	112
Modifying Tables	115
Rename a Table or Column	115
Display, Add, and Delete Columns	117
Display, Add, and Delete Rows	119
Display, Add, Modify, and Delete Constraints	120
Update a Column of Data	124
Update Rows of Data	125
Update Rows of Data with the Results of a Query	126
Delete a Table	128
Indexes	129
Create an Index to Speed Up Queries	131
Views	133
Create a View to Save the Results of a Query	135

Transaction Management	138
Double-Check Changes Before a COMMIT	139
Undo Changes with a ROLLBACK	141
Chapter 6: Data Types	143
How to Choose a Data Type	145
Numeric Data	147
Numeric Values	147
Integer Data Types	148
Decimal Data Types	150
Floating Point Data Types	151
String Data	154
String Values	154
Character Data Types	156
Unicode Data Types	159
Datetime Data	161
Datetime Values	161
Datetime Data Types	165
Other Data	172
Boolean Data	172
External Files (Images, Documents, etc.)	173
Chapter 7: Operators and Functions	179
Operators	180
Logical Operators	181
Comparison Operators	182
Math Operators	189
Aggregate Functions	191
Numeric Functions	193
Apply Math Functions	194
Generate Random Numbers	196

Round and Truncate Numbers	197
Convert Data to a Numeric Data Type	198
String Functions	199
Find the Length of a String	199
Change the Case of a String	200
Trim Unwanted Characters Around a String	201
Concatenate Strings	203
Search for Text in a String	203
Extract a Portion of a String	206
Replace Text in a String	207
Delete Text from a String	208
Use Regular Expressions	209
Convert Data to a String Data Type	217
Datetime Functions	218
Return the Current Date or Time	218
Add or Subtract a Date or Time Interval	220
Find the Difference Between Two Dates or Times	221
Extract a Part of a Date or Time	226
Determine the Day of the Week of a Date	228
Round a Date to the Nearest Time Unit	229
Convert a String to a Datetime Data Type	230
Null Functions	234
Return an Alternative Value if There Is a Null Value	235
Chapter 8: Advanced Querying Concepts	237
Case Statements	238
Display Values Based on If-Then Logic for a Single Column	239
Display Values Based on If-Then Logic for Multiple Columns	240
Grouping and Summarizing	242

GROUP BY Basics	242
Aggregate Rows into a Single Value or List	245
ROLLUP, CUBE, and GROUPING SETS	247
Window Functions	250
Rank the Rows in a Table	252
Return the First Value in Each Group	255
Return the Second Value in Each Group	256
Return the First Two Values in Each Group	257
Return the Prior Row Value	258
Calculate the Moving Average	259
Calculate the Running Total	261
Pivoting and Unpivoting	263
Break Up the Values of a Column into Multiple Columns	263
List the Values of Multiple Columns in a Single Column	265
Chapter 9: Working with Multiple Tables and Queries	269
Joining Tables	270
Join Basics and INNER JOIN	274
LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN	277
USING and NATURAL JOIN	279
CROSS JOIN and Self Join	281
Union Operators	284
UNION	285
EXCEPT and INTERSECT	289
Common Table Expressions	291
CTEs Versus Subqueries	293
Recursive CTEs	295

Chapter 10: How Do I...?	303
Find the Rows Containing Duplicate Values	303
Select Rows with the Max Value for Another Column	306
Concatenate Text from Multiple Fields into a Single Field	308
Find All Tables Containing a Specific Column Name	311
Update a Table Where the ID Matches Another Table	313
Index	317

Why SQL?

Since the last edition of *SQL Pocket Guide* was published, a lot has changed in the data world. The amount of data generated and collected has exploded, and a number of tools and jobs have been created to handle the influx of data. Through all of the changes, SQL has remained an integral part of the data landscape.

Over the past 15 years, I have worked as an engineer, consultant, analyst, and data scientist, and I have used SQL in every one of my roles. Even if my main responsibilities were focused on another tool or skill, I had to know SQL in order to access data at a company.

If there was a programming language award for best supporting actor, SQL would take home the prize.

As new technologies emerge, SQL is still top of mind when it comes to working with data. Cloud-based storage solutions like Amazon Redshift and Google BigQuery require users to write SQL queries to pull data. Distributed data processing frameworks like Hadoop and Spark have sidekicks Hive and Spark SQL, respectively, which provide SQL-like interfaces for users to analyze data.

SQL has been around for almost five decades, and it is not going away anytime soon. It is one of the oldest programming languages still being used widely today, and I am excited to share the latest and greatest with you in this book.

Goals of This Book

There are many existing SQL books out there, ranging from ones that teach beginners how to code in SQL to detailed technical specifications for database administrators. This book is not intended to cover all SQL concepts in depth, but rather to be a simple reference for when:

- You've forgotten some SQL syntax and need to look it up quickly
- You've come across a slightly different set of database tools at a new job and need to look up the nuanced differences
- You've been focusing on another coding language for a while and need a quick refresher on how SQL works

If SQL plays a large supporting role in your job, then this is the perfect pocket guide for you.

Updates to the Fourth Edition

The third edition of the *SQL Pocket Guide* by Jonathan Gennick was published in 2010, and it was well received by readers. I've made the following updates to the fourth edition:

- The syntax has been updated for Microsoft SQL Server, MySQL, Oracle Database, and PostgreSQL. IBM's Db2 has been removed due to its decrease in popularity, and SQLite has been added due to its increase in popularity.
- The third edition of this book was organized alphabetically. I've rearranged the sections in the fourth edition so that similar concepts are grouped together. There is still an

index at the end of this book that lists concepts alphabetically.

- Due to the number of data analysts and data scientists who are now using SQL in their jobs, I've added sections on how to use SQL with Python and R (popular open source programming languages), as well as a SQL crash course for those who need a quick refresher.

Frequently Asked (SQL) Questions

The last chapter of this book is called "[How Do I...?](#)" and it includes frequently asked questions by SQL beginners or those who haven't used SQL in a while.

It's a good place to start if you don't remember the exact keyword or concept that you're looking for. Example questions include:

- How do I find the rows containing duplicate values?
- How do I select rows with the max value for another column?
- How do I concatenate text from multiple fields into a single field?

Navigating This Book

This book is organized into three sections.

I. Basic Concepts

- Chapters [1](#) through [3](#) introduce basic keywords, concepts, and tools for writing SQL code.
- [Chapter 4](#) breaks down each clause of a SQL query.

II. Database Objects, Data Types, and Functions

- [Chapter 5](#) lists common ways to create and modify objects within a database.
- [Chapter 6](#) lists common data types that are used in SQL.
- [Chapter 7](#) lists common operators and functions in SQL.

III. Advanced Concepts

- Chapters [8](#) and [9](#) explain advanced querying concepts including joins, case statements, window functions, etc.
- [Chapter 10](#) walks through solutions to some of the most commonly searched for SQL questions.

Conventions Used in This Book

The following typographical conventions are used in this book:

Italic

Indicates new terms, URLs, email addresses, filenames, and file extensions.

Constant width

Used for program listings, as well as within paragraphs to refer to program elements such as variable or function names, databases, data types, environment variables, statements, and keywords.

Constant width bold

Shows commands or other text that should be typed literally by the user, or values determined by context.

TIP

This element signifies a tip or suggestion.

NOTE

This element signifies a general note.

WARNING

This element indicates a warning or caution.

Using Code Examples

If you have a technical question or a problem using the code examples, please send email to bookquestions@oreilly.com.

This book is here to help you get your job done. In general, if example code is offered with this book, you may use it in your programs and documentation. You do not need to contact us for permission unless you're reproducing a significant portion of the code. For example, writing a program that uses several chunks of code from this book does not require permission. Selling or distributing examples from O'Reilly books does require permission. Answering a question by citing this book and quoting example code does not require permission. Incorporating a significant amount of example code from this book into your product's documentation does require permission.

We appreciate, but generally do not require, attribution. An attribution usually includes the title, author, publisher, and ISBN. For example: “*SQL Pocket Guide, 4th ed.* by Alice Zhao (O'Reilly). Copyright 2021 Alice Zhao, 978-1-492-209040-3.”

If you feel your use of code examples falls outside fair use or the permission given above, feel free to contact us at permissions@oreilly.com.

O'Reilly Online Learning



For more than 40 years, *O'Reilly Media* has provided technology and business training, knowledge, and insight to help companies succeed.

Our unique network of experts and innovators share their knowledge and expertise through books, articles, and our online learning platform. O'Reilly's online learning platform gives you on-demand access to live training courses, in-depth learning paths, interactive coding environments, and a vast collection of text and video from O'Reilly and 200+ other publishers. For more information, visit <http://oreilly.com>.

How to Contact Us

Please address comments and questions concerning this book to the publisher:

O'Reilly Media, Inc.
1005 Gravenstein Highway North
Sebastopol, CA 95472
800-998-9938 (in the United States or Canada)
707-829-0515 (international or local)
707-829-0104 (fax)

We have a web page for this book, where we list errata, examples, and any additional information. You can access this page at <https://oreil.ly/jreAj>.

Email bookquestions@oreilly.com to comment or ask technical questions about this book.

For news and information about our books and courses, visit <http://oreilly.com>.

Find us on Facebook: <http://facebook.com/oreilly>.

Follow us on Twitter: <http://twitter.com/oreillymedia>.

Watch us on YouTube: <http://youtube.com/oreillymedia>.

Acknowledgments

Thank you to Jonathan Gennick for creating this pocket guide from scratch and writing the first three editions, and to Andy Kwan for trusting me to continue on with the publication.

I couldn't have completed this book without the help of my editors Amelia Blevins, Jeff Bleiel, and Caitlin Ghegan, and my technical reviewers Alicia Nevels, Joan Wang, Scott Haines, and Thomas Nield. I truly appreciate the time you've spent reading each page of this book. Your feedback has been invaluable.

To my parents—thank you for fostering my love for learning and creating. To my kids Henry and Lily—your excitement for this book warms my heart. Finally, to my husband, Ali—thank you for all of your notes on this book, for your encouragement, and for being my biggest fan.

SQL Crash Course

This short chapter is intended to quickly get you up to speed on basic SQL terminology and concepts.

What Is a Database?

Let's start with the basics. A *database* is a place to store data in an organized way. There are many ways to organize data, and as a result, there are many databases to choose from. The two categories that databases fall into are *SQL* and *NoSQL*.

SQL

SQL is short for *Structured Query Language*. Imagine you have an app that remembers all of your friend's birthdays. SQL is the most popular language you would use to talk to that app.

English: "Hey app. When is my husband's birthday?"

SQL: `SELECT * FROM birthdays
WHERE person = 'husband';`

SQL databases are often called *relational databases* because they are made up of relations, which are more commonly referred to as tables. Many tables connected to each other make up a database. [Figure 1-1](#) shows a picture of a relation in a SQL database.

The diagram illustrates a SQL relation, also known as a table. It features a rectangular grid with three columns labeled 'ID', 'Name', and 'Age'. The first row contains the column headers. Below these headers, there are three data rows: the first row has values 1, Ali, and 36; the second row has values 2, Henry, and 5; and the third row has values 3, Lily, and 2. A blue arrow points from the left towards the first data row, labeled 'Tuple (row)'. Another blue arrow points downwards from the top towards the column headers, labeled 'Attribute (column)'. A blue bracket at the bottom of the grid is labeled 'Relation (table)'.

ID	Name	Age
1	Ali	36
2	Henry	5
3	Lily	2

Figure 1-1. A relation (also known as a table) in a SQL database

The main thing to note about SQL databases is that they require predefined *schemas*. You can think of a schema as the way that data in a database is organized or structured. Let's say you'd like to create a table. Before loading any data into the table, the structure of the table must first be decided on, including things like what columns are in the table, whether those columns hold integer or decimal values, etc.

There comes a time, though, when data cannot be organized in such a structured way. Your data may have varying fields or you may need a more effective way of storing and accessing a large amount of data. That's where NoSQL comes in.

NoSQL

NoSQL stands for *not only SQL*. It will not be covered in detail in this book, but I wanted to point it out because the term has grown a lot in popularity since the 2010s and it's important to understand there are ways to store data beyond just tables.

NoSQL databases are often referred to as *non-relational databases*, and they come in all shapes and sizes. Their main characteristics are that they have dynamic schemas (meaning the schema doesn't have to be locked in up front) and they allow for horizontal scaling (meaning the data can spread across multiple machines).

The most popular NoSQL database is *MongoDB*, which is more specifically a document database. Figure 1-2 shows a picture of how data is stored in MongoDB. You'll notice that the data is no longer in a structured table and the number of fields (similar to a column) varies for each document (similar to a row).

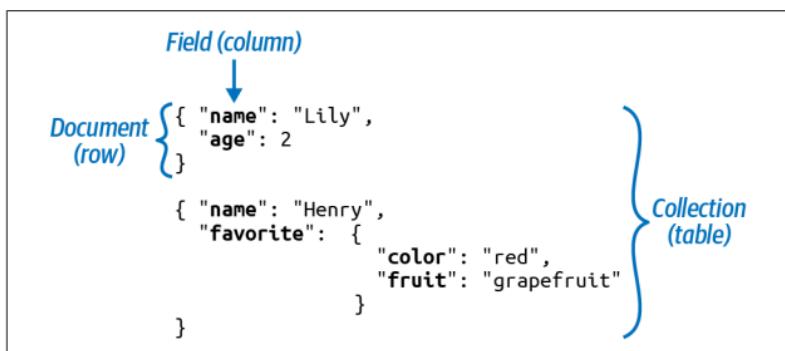


Figure 1-2. A collection (a variant of a table) in MongoDB, a NoSQL database

That all said, the focus of this book is on SQL databases. Even with the introduction of NoSQL, most companies still store the majority of their data in tables in relational databases.

Database Management Systems (DBMS)

You may have heard terms like *PostgreSQL* or *SQLite*, and be wondering how they are different from SQL. They are two types of *Database Management Systems* (DBMS), which is software used to work with a database.

This includes things like figuring out how to import data and organize it, as well as things like managing how users or other programs access the data. A *Relational Database Management System* (RDBMS) is software that is specifically for relational databases, or databases made up of tables.

Each RDBMS has a different implementation of SQL, meaning that the syntax varies slightly from software to software. For example, this is how you would output 10 rows of data in 5 different RDBMSs:

MySQL, PostgreSQL, and SQLite

```
SELECT * FROM birthdays LIMIT 10;
```

Microsoft SQL Server

```
SELECT TOP 10 * FROM birthdays;
```

Oracle Database

```
SELECT * FROM birthdays WHERE ROWNUM <= 10;
```

Googling SQL Syntax

When searching for SQL syntax online, always include the RDBMS you are working with in the search. When I first learned SQL, I could not for the life of me figure out why my copy-pasted code from the internet didn't work and this was the reason!

Do this.

Search: *create table datetime postgresql*

→ Result: *timestamp*

Search: *create table datetime microsoft sql server*

→ Result: *datetime*

Not this.

Search: *create table datetime*

→ Result: syntax could be for any RDBMS

This book covers SQL basics along with the nuances of five popular database management systems: Microsoft SQL Server, MySQL, Oracle Database, PostgreSQL and SQLite.

Some are proprietary, meaning they are owned by a company and cost money to use, and others are open source, meaning they are free for anyone to use. **Table 1-1** details the differences between the RDBMSs.

Table 1-1. RDBMS comparison table

RDBMS	Owner	Highlights
Microsoft SQL Server	Microsoft	<ul style="list-style-type: none">- Popular proprietary RDBMS- Often used alongside other Microsoft products including Microsoft Azure and the .NET framework- Common on the Windows platform- Also referred to as <i>MSSQL</i> or <i>SQL Server</i>
MySQL	Open Source	<ul style="list-style-type: none">- Popular open source RDBMS- Often used alongside web development languages like HTML/CSS/Javascript- Acquired by Oracle, though still open source
Oracle Database	Oracle	<ul style="list-style-type: none">- Popular proprietary RDBMS- Often used at large corporations given the amount of features, tools, and support available- Also referred to simply as <i>Oracle</i>
PostgreSQL	Open Source	<ul style="list-style-type: none">- Quickly growing in popularity- Often used alongside open source technologies like Docker and Kubernetes- Efficient and great for large datasets
SQLite	Open Source	<ul style="list-style-type: none">- World's most used database engine- Common on iOS and Android platforms- Lightweight and great for a small database

NOTE

Going forward in this book:

- Microsoft SQL Server will be referred to as *SQL Server*.
 - Oracle Database will be referred to as *Oracle*.
-

Installation instructions and code snippets for each RDBMS can be found in [RDBMS Software](#) in [Chapter 2](#).

A SQL Query

A common acronym in the SQL world is *CRUD*, which stands for “Create, Read, Update, and Delete.” These are the four major operations that are available within a database.

SQL Statements

People who have *read and write access* to a database are able to perform all four operations. They can create and delete tables, update data in tables, and read data from tables. In other words, they have all the power.

They write *SQL statements*, which is general SQL code that can be written to perform any of the CRUD operations. These people often have titles like *database administrator* (DBA) or *database engineer*.

SQL Queries

People who have *read access* to a database are only able to perform the read operation, meaning they can look at data in tables.

They write *SQL queries*, which are a more specific type of SQL statement. Queries are used for finding and displaying data, otherwise known as “reading” data. This action is sometimes referred to as *querying tables*. These people often have titles like *data analyst* or *data scientist*.

The next two sections are a quick-start guide for writing SQL queries, since it is the most common type of SQL code that you’ll see. More details on creating and updating tables can be found in [Chapter 5](#).

The SELECT Statement

The most basic SQL query (which will work in any SQL software) is:

```
SELECT * FROM my_table;
```

which says, show me all of the data within the table named `my_table`—all of the columns and all of the rows.

While SQL is case-insensitive (`SELECT` and `select` are equivalent), you'll notice that some words are in all caps and others are not.

- The uppercase words in the query are called *keywords*, meaning that SQL has reserved them to perform some sort of operation on the data.
- All other words are lowercase. This includes table names, column names, etc.

The uppercase and lowercase formats are not enforced, but it is a good style convention to follow for readability's sake.

Let's go back to this query:

```
SELECT * FROM my_table;
```

Let's say that instead of returning all of the data in its current state, I want to:

- Filter the data
- Sort the data

This is where you would modify the `SELECT` statement to include a few more *clauses*, and the result would look something like this:

```
SELECT *
FROM my_table
WHERE column1 > 100
ORDER BY column2;
```

More details on all of the clauses can be found in [Chapter 4](#), but the main thing to note is this: the clauses must always be listed in the same order.

Memorize This Order

All SQL queries will contain some combination of these clauses. If you remember nothing else, remember this order!

SELECT	-- columns to display
FROM	-- table(s) to pull from
WHERE	-- filter rows
GROUP BY	-- split rows into groups
HAVING	-- filter grouped rows
ORDER BY	-- columns to sort

NOTE

The -- is the start of a [comment](#) in SQL, meaning the text after it is just for documentation's sake and the code will not be executed.

For the most part, the `SELECT` and `FROM` clauses are required and all other clauses are optional. The exception is if you are [selecting a particular database function](#), then only the `SELECT` is required.

The classic mnemonic to remember the order of the clauses is:

Sweaty feet will give horrible odors.

If you don't want to think about sweaty feet each time you write a query, here's one that I made up:

Start Fridays with grandma's homemade oatmeal.

Order of Execution

The order that SQL code is executed is not something typically taught in a beginner SQL course, but I'm including it here because it's a common question I received when I taught SQL to students coming from a Python coding background.

A sensible assumption would be that the order that you *write* the clauses is the same order that the computer *executes* the clauses, but that is not the case. After a query is run, this is the order that the computer works through the data:

1. FROM
2. WHERE
3. GROUP BY
4. HAVING
5. SELECT
6. ORDER BY

Compared to the order that you actually write the clauses, you'll notice that the `SELECT` has been moved to the fifth position. The high-level takeaway here is that SQL works in this order:

1. Gathers all of the data with the `FROM`
2. Filters rows of data with the `WHERE`
3. Groups rows together with the `GROUP BY`
4. Filters grouped rows with the `HAVING`
5. Specifies columns to display with the `SELECT`
6. Rearranges the results with the `ORDER BY`

A Data Model

I'd like to spend the final section of the crash course going over a simple *data model* and point out some terms that you'll often hear in fun SQL conversations around the office.

A data model is a visualization that summarizes how all of the tables in a database are related to one another, along with some details about each table. **Figure 1-3** is a simple data model of a student grades database.

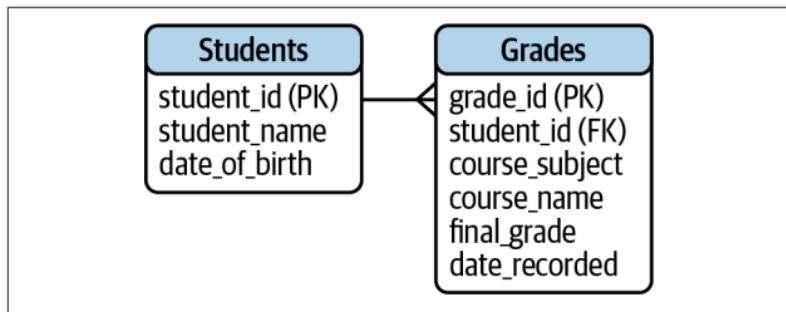


Figure 1-3. A data model of student grades

Table 1-2 lists out the technical terms that describe what's happening in the data model.

Table 1-2. Terms used to describe what's in a data model

Term	Definition	Example
Database	A database is a place to store data in an organized way.	This data model shows all of the data in the student grades database.
Table	A table is made up of rows and columns. In the data model, they are represented by rectangles.	There are two tables in the student grades database: Students and Grades.

Term	Definition	Example
Column	A table consists of multiple columns, which are sometimes referred to as attributes or fields. Each column contains a particular type of data. In the data model, all of the columns in a table are listed within each rectangle.	In the Students table, the columns are <code>student_id</code> , <code>student_name</code> , and <code>date_of_birth</code> .
Primary Key	A <i>primary key</i> uniquely identifies each row of data in a table. A primary key can be made up of one or more columns in a table. In a data model, it is flagged as pk or with a key icon.	In the Students table, the primary key is the <code>student_id</code> column, meaning that the <code>student_id</code> value is different for each row of data.
Foreign Key	A <i>foreign key</i> in a table refers to a primary key in another table. The two tables can be linked together by the common column. A table can have multiple foreign keys. In a data model, it is flagged as fk.	In the Grades table, <code>student_id</code> is a foreign key, meaning that the values in that column match up with values in the corresponding primary key column in the Students table.
Relationship	A <i>relationship</i> describes how the rows in one table map to the rows in another table. In a data model, it is represented by a line with symbols at the end points. Common types are one-to-one and one-to-many relationships.	In this data model, the two tables have a one-to-many relationship represented by the fork. One student can have many grades, or one row of the Students table maps to multiple rows in the Grades table.

More details on these terms can be found in “[Creating Tables](#)” on page 97 in [Chapter 5](#).

You might be wondering why we're spending so much time reading a data model instead of writing SQL code already! The reason is because you'll often be writing queries that link up a number of tables, so it's a good idea to first get familiar with the data model to know how they all connect.

Data models can typically be found in a documentation repository at a company. You may want to print out the data models that you frequently work with—both for easy reference and easy desk decor.

You can also write queries within an RDBMS to look up information contained in a data model, such as the tables in a database, the columns of a table, or the constraints of a table.

And That Is Your Crash Course!

The remainder of this book is intended to be a reference book and does not need to be read in order. Please use it to look up concepts, keywords, and standards.

Where Can I Write SQL Code?

This chapter covers three places where you can write SQL code:

RDBMS Software

To write SQL code, you first have to download an RDBMS like MySQL, Oracle, PostgreSQL, SQL Server, or SQLite. The nuances of each RDBMS are highlighted in “[RDBMS Software](#)” on page 14.

Database Tools

Once you’ve downloaded an RDBMS, the most basic way to write SQL code is through a *terminal window*, which is a text-only black-and-white screen. Most people prefer to use a *database tool* instead, which is a more user-friendly application that connects to an RDBMS behind the scenes.

A database tool will have a *graphical user interface* (GUI), which allows users to visually explore tables and more easily edit SQL code. “[Database Tools](#)” on page 20 goes through how to connect a database tool to an RDBMS.

Other Programming Languages

SQL can be written within many other programming languages. This chapter focuses on two in particular: Python and R. They are popular open source programming

languages used by data scientists and data analysts, who often need to write SQL code as well.

Instead of switching back and forth between Python/R and an RDBMS, you can connect Python/R directly to an RDBMS and write SQL code within Python/R. “[Other Programming Languages](#)” on page 24 walks through how to do so step by step.

RDBMS Software

This section includes installation instructions and short code snippets for the five RDBMSs that are covered in this book.

Which RDBMS to Choose?

If you are working at a company that is already using an RDBMS, you will need to use the same one.

If you are working on a personal project, you will need to decide which RDBMS to use. You can refer back to [Table 1-1](#) in [Chapter 1](#) to review the details of some popular ones.

Quick Start with SQLite

Want to start writing SQL code as soon as possible? SQLite is the fastest RDBMS to set up.

Compared to the other RDBMSs in this book, it’s less secure and can’t handle multiple users, but it provides basic SQL functionality in a compact package.

Because of this, I’ve moved SQLite up to the front of each section of this chapter since its setup is generally more straightforward than the others.

What Is a Terminal Window?

I'll often refer to a terminal window in this chapter because once you've downloaded an RDBMS, it's the most basic way to interact with the RDBMS.

A *terminal window* is an application on your computer that typically has a black background and only allows text inputs. The application name varies by operating system:

- On Windows, use the Command Prompt application.
- On macOS and Linux, use the Terminal application.

Once you open up a terminal window, you'll see a *command prompt*, which looks like a > followed by a flashing box. This means that it's ready to take in text commands from the user.

TIP

The next sections include links to download RDBMS installers for Windows, macOS, and Linux.

On macOS and Linux, an alternative to downloading an installer is to use the [Homebrew](#) package manager instead. Once you install Homebrew, you can run simple `brew install` commands from the Terminal to do all of the RDBMS installations.

SQLite

SQLite is free and the most lightweight install, meaning that it doesn't take up much space on your computer and is extremely quick to set up. For Windows and Linux, SQLite Tools can be downloaded from the [SQLite Download Page](#). macOS comes with SQLite already installed.

TIP

The simplest way to start using SQLite is to open a **terminal window** and type **sqlite3**. With this approach, however, everything is done in memory, meaning that changes will not be saved once you close SQLite.

```
> sqlite3
```

If you want your changes to be saved, you should connect to a database upon opening with the following syntax:

```
> sqlite3 my_new_db.db
```

The command prompt for SQLite looks like this:

```
sqlite>
```

Some quick code to test things out:

```
sqlite> CREATE TABLE test (id int, num int);
sqlite> INSERT INTO test VALUES (1, 100), (2, 200);
sqlite> SELECT * FROM test LIMIT 1;
```

```
1|100
```

To show databases, show tables, and exit:

```
sqlite> .databases
sqlite> .tables
sqlite> .quit
```

TIP

If you want to display column names in your output, type:

```
sqlite> .headers on
```

To hide them again, type:

```
sqlite> .headers off
```

MySQL

MySQL is free, even though it is now owned by Oracle. MySQL Community Server can be downloaded from the [MySQL Community Downloads](#) page. On macOS and Linux, alternatively, you can do the installation with Homebrew by typing `brew install mysql` in the Terminal.

The command prompt for MySQL looks like this:

```
mysql>
```

Some quick code to test things out:

```
mysql> CREATE TABLE test (id int, num int);
mysql> INSERT INTO test VALUES (1, 100), (2, 200);
mysql> SELECT * FROM test LIMIT 1;

+-----+-----+
| id   | num   |
+-----+-----+
|    1 |  100 |
+-----+-----+
1 row in set (0.00 sec)
```

To show databases, switch databases, show tables, and exit:

```
mysql> show databases;
mysql> connect another_db;
mysql> show tables;
mysql> quit
```

Oracle

Oracle is proprietary and works on Windows and Linux machines. Oracle Database Express Edition, the free edition, can be downloaded from the [Oracle Database XE Downloads page](#).

The command prompt for Oracle looks like this:

```
SQL>
```

Some quick code to test things out:

```
SQL> CREATE TABLE test (id int, num int);
SQL> INSERT INTO test VALUES (1, 100);
SQL> INSERT INTO test VALUES (2, 200);
SQL> SELECT * FROM test WHERE ROWNUM <=1;
```

ID	NUM
1	100

To show databases, show all tables (including system tables), show user-created tables, and exit:

```
SQL> SELECT * FROM global_name;
SQL> SELECT table_name FROM all_tables;
SQL> SELECT table_name FROM user_tables;
SQL> quit
```

PostgreSQL

PostgreSQL is free and often used alongside other open source technologies. PostgreSQL can be downloaded from the [PostgreSQL Downloads](#) page. On macOS and Linux, alternatively, you can do the installation with Homebrew by typing `brew install postgresql` in the Terminal.

The command prompt for PostgreSQL looks like this:

```
postgres=#
```

Some quick code to test things out:

```
postgres=# CREATE TABLE test (id int, num int);
postgres=# INSERT INTO test VALUES (1, 100),
          (2, 200);
postgres=# SELECT * FROM test LIMIT 1;
```

id	num
1	100

```
1 | 100
(1 row)
```

To show databases, switch databases, show tables, and exit:

```
postgres=# \l
postgres=# \c another_db
postgres=# \d
postgres=# \q
```

TIP

If you ever see `postgres-#`, that means that you've forgotten a semicolon at the end of a SQL statement. Type ; and you should see `postgres=#` again.

If you ever see :, that means you've been automatically switched to the vi text editor, and you can exit by typing q.

SQL Server

SQL Server is proprietary (owned by Microsoft) and works on Windows and Linux machines. It can also be installed via Docker. SQL Server Express, the free edition, can be downloaded from the [Microsoft SQL Server Downloads](#) page.

The command prompt for SQL Server looks like this:

```
1>
```

Some quick code to test things out:

```
1> CREATE TABLE test (id int, num int);
2> INSERT INTO test VALUES (1, 100), (2, 200);
3> go
1> SELECT TOP 1 * FROM test;
2> go
```

```
id      num
-----
1       100
(1 row affected)
```

To show databases, switch databases, show tables, and exit:

```
1> SELECT name FROM master.sys.databases;
2> go
1> USE another_db;
2> go
1> SELECT * FROM information_schema.tables;
2> go
1> quit
```

NOTE

In *SQL Server*, SQL code is not executed until you type the `go` command on a new line.

Database Tools

Instead of working with an RDBMS directly, most people will use a database tool to interact with a database. A database tool comes with a nice graphical user interface that allows you to point, click, and write SQL code in a user-friendly setting.

Behind the scenes, a database tool uses a *database driver*, which is software that helps the database tool talk to a database. **Figure 2-1** shows the visual differences between accessing a database directly through a terminal window versus indirectly through a database tool.

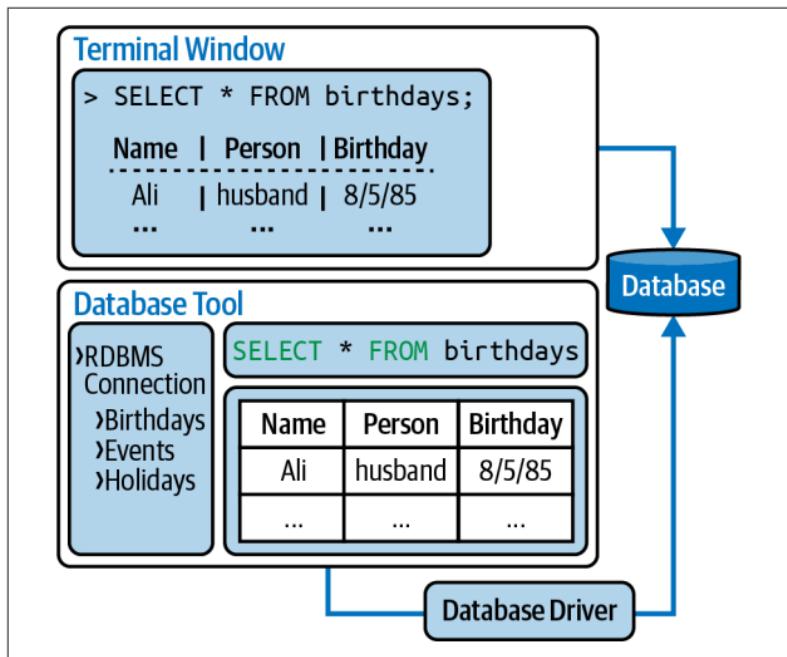


Figure 2-1. Accessing an RDBMS through a terminal window versus a database tool

There are a number of database tools available. Some work specifically with a single RDBMS, and others work with multiple RDBMSs. **Table 2-1** lists each RDBMS along with one of the most popular database tools for that particular RDBMS. All of the database tools in the table are free to download and use, and there are many other proprietary ones out there as well.

Table 2-1. Database tool comparison table

RDBMS	Database Tool	Details
SQLite	DB Browser for SQLite	- Different developer than SQLite - One of many tool options for SQLite
MySQL	MySQL Workbench	- Same developer as MySQL
Oracle	Oracle SQL Developer	- Developed by Oracle
PostgreSQL	pgAdmin	- Different contributors than PostgreSQL - Included with the PostgreSQL install

RDBMS	Database Tool	Details
SQL Server	SQL Server Management Studio	- Developed by Microsoft
Multiple	DBeaver	- One of many tool options for connecting to a variety of RDBMSs (including any of the preceding five listed)

Connect a Database Tool to a Database

When opening up a database tool, the first step is to connect to a database. This can be done in several ways:

Option 1: Create a New Database

You can create a brand-new database by writing a `CREATE` statement:

```
CREATE DATABASE my_new_db;
```

Afterward, you can create tables to populate the database. More details can be found in “[Creating Tables](#)” on page 97 in [Chapter 5](#).

Option 2: Open Up a Database File

You may have downloaded or been given a file with a `.db` extension:

```
my_new_db.db
```

This `.db` file will already contain a number of tables. You can simply open it up within a database tool and start interacting with the database.

Option 3: Connect to an Existing Database

You may want to work with a database that is either on your computer or on a *remote server*, meaning that the data is on a computer located elsewhere. This is extremely common these days with *cloud computing*, where people use servers owned by companies like Amazon, Google, or Microsoft.

Database Connection Fields

To connect to a database, you'll need to fill out the following fields within a database tool:

Host

Where the database is located.

- If the database is on your computer, then this should be *localhost* or *127.0.0.1*.
- If the database is on a remote server, then this should be the IP address of that computer example: *123.45.678.90*.

Port

How to connect to the RDBMS.

There should already be a default port number in this field, and you shouldn't change it. It will be different for each RDBMS.

- MySQL: *3306*
- Oracle: *1521*
- PostgreSQL: *5432*
- SQL Server: *1433*

Database

The name of the database you'd like to connect to.

Username

Your username for the database.

There may already be a default username in this field. If you don't remember setting up a username, keep the default value.

Password

Your password associated with the username.

If you don't remember setting up a password for your username, try leaving this field blank.

NOTE

For *SQLite*, instead of filling out these five database connection fields, you would enter in the file path of the *.db* database file you are trying to connect to.

Once you fill in the database connection fields correctly, you should have access to the database. You can now use the database tool to find the tables and fields you are interested in, and start writing SQL code.

Other Programming Languages

SQL can be written within a number of other programming languages. This chapter focuses on two popular open source ones: Python and R.

As a data scientist or data analyst, you likely do your analysis in Python or R, and also need to write SQL queries to pull data from a database.

A Basic Data Analysis Workflow

1. Write a SQL query within a database tool.
2. Export the results as a *.csv* file.
3. Import the *.csv* file into Python or R.
4. Continue doing analysis in Python or R.

The preceding approach is fine for doing a quick, one-time export. However, if you need to continuously edit your SQL query or are working with multiple queries, this can get annoying very quickly.

A Better Data Analysis Workflow

1. Connect Python or R to a database.
2. Write SQL queries within Python or R.
3. Continue doing analysis in Python or R.

This second approach allows you to do all of your querying and analysis within one tool, which is helpful if you need to tweak your queries as you are doing analysis. The remainder of this chapter provides code for each step of this second workflow.

Connect Python to a Database

It takes three steps to connect Python to a database:

1. Install a database driver for Python.
2. Set up a database connection in Python.
3. Write SQL code in Python.

Step 1: Install a database driver for Python

A database driver is software that helps Python talk to a database, and there are many driver options to choose from. **Table 2-2** includes code for how to install a popular driver for each RDBMS.

This is a one-time installation you'll need to do via either a `pip install` or a `conda install`. The following code should be run in a **terminal window**.

Table 2-2. Install a driver for Python using either pip or conda

RDBMS	Option	Code
SQLite	n/a	No install necessary (Python 3 comes with sqlite3)
MySQL	pip	<code>pip install mysql-connector-python</code>
	conda	<code>conda install -c conda-forge mysql-connector-python</code>
Oracle	pip	<code>pip install cx_Oracle</code>
	conda	<code>conda install -c conda-forge cx_oracle</code>
PostgreSQL	pip	<code>pip install psycopg2</code>
	conda	<code>conda install -c conda-forge psycopg2</code>
SQL Server	pip	<code>pip install pyodbc</code>
	conda	<code>conda install -c conda-forge pyodbc</code>

Step 2: Set up a database connection in Python

To set up a database connection, you first need to know the location and name of the database you are trying to connect to, as well as your username and password. More details can be found in “[Database Connection Fields](#)” on page 23.

Table 2-3 contains the Python code you need to run each time you plan on writing SQL code in Python. You can include it at the top of your Python script.

Table 2-3. Python code to set up a database connection

RDBMS	Code
SQLite	<pre>import sqlite3 conn = sqlite3.connect('my_new_db.db')</pre>
MySQL	<pre>import mysql.connector conn = mysql.connector.connect(host='localhost', database='my_new_db', user='alice', password='password')</pre>
Oracle	<pre># Connecting to Oracle Express Edition import cx_Oracle conn = cx_Oracle.connect(dsn='localhost/XE', user='alice', password='password')</pre>
PostgreSQL	<pre>import psycopg2 conn = psycopg2.connect(host='localhost', database='my_new_db', user='alice', password='password')</pre>
SQL Server	<pre># Connecting to SQL Server Express import pyodbc conn = pyodbc.connect(driver='{SQL Server}', host='localhost\SQLEXPRESS', database='my_new_db', user='alice', password='password')</pre>

TIP

Not all arguments are required. If you exclude an argument completely, then the default value will be used. For example, the default host is *localhost*, which is your computer. If no username and password were set up, then those arguments can be left out.

Keeping Your Passwords Safe in Python

The preceding code is fine for testing out a connection to a database, but in reality, you should not be saving your password within a script for everyone to see.

There are multiple ways to avoid doing so, including:

- generating an SSH key
- setting environment variables
- creating a configuration file

These options, however, all require additional knowledge of computers or file formats.

The recommended approach: create a separate Python file.

The most straightforward approach, in my opinion, is to save your username and password in a separate Python file, and then call that file within your database connection script. While this is less secure than the other options, it is the quickest start.

To use this approach, start by creating a *db_config.py* file with the following code:

```
usr = "alice"  
pwd = "password"
```

Import the *db_config.py* file when setting up your database connection. The following example modifies the Oracle code from [Table 2-3](#) to use the *db_config.py* values instead of hardcoded user and password values (changes are bolded):

```
import cx_Oracle
import db_config

conn = cx_Oracle.connect(dsn='localhost/XE',
    user=db_config.usr,
    password=db_config.pwd)
```

Step 3: Write SQL code in Python

Once the database connection has been established, you can start writing SQL queries within your Python code.

Write a simple query to test your database connection:

```
cursor = conn.cursor()
cursor.execute('SELECT * FROM test;')
result = cursor.fetchall()
print(result)

[(1, 100),
 (2, 200)]
```

WARNING

When using `cx_Oracle` in Python, remove the semicolon
(`;`) at the end of all queries to avoid getting an error.

Save the results of a query as a pandas dataframe:

```
# pandas must already be installed
import pandas as pd

df = pd.read_sql('''SELECT * FROM test;''', conn)
print(df)
print(type(df))

   id  num
0   1  100
1   2  200
<class 'pandas.core.frame.DataFrame'>
```

Close the connection when you are done using the database:

```
cursor.close()  
conn.close()
```

It is always good practice to close the database connection to save resources.

SQLAlchemy for Python Lovers

Another popular way to connect to a database is using the SQLAlchemy package in Python. It is an *object relational mapper* (ORM), which turns database data into Python objects, allowing you to code in pure Python instead of using SQL syntax.

Imagine you want to see all the table names in a database. (The following code is PostgreSQL-specific, but SQLAlchemy will work with any RDBMS.)

Without SQLAlchemy:

```
pd.read_sql("""SELECT tablename  
    FROM pg_catalog.pg_tables  
    WHERE schemaname='public'""", conn)
```

With SQLAlchemy:

```
conn.table_names()
```

When using SQLAlchemy, the `conn` object comes with a `table_names()` Python method, which you may find easier to remember than SQL syntax. While SQLAlchemy provides cleaner Python code, it does slow down performance due to the additional time it spends turning data into Python objects.

To use SQLAlchemy in Python:

1. You must already have a **database driver** (like `psycopg2`) installed.
2. In a terminal window, type `pip install sqlalchemy` or a `conda install -c conda-forge sqlalchemy` to install SQLAlchemy.
3. Run the following code in Python to set up a SQLAlchemy connection. (The following code is PostgreSQL-specific.)

The [SQLAlchemy documentation](#) provides code for other RDBMSs and drivers:

```
from sqlalchemy import create_engine
conn = create_engine('postgresql+psycopg2://
                      alice:password@localhost:5432/my_new_db')
```

Connect R to a Database

It takes three steps to connect R to a database:

1. Install a database driver for R
2. Set up a database connection in R
3. Write SQL code in R

Step 1: Install a database driver for R

A database driver is software that helps R talk to a database, and there are many driver options to choose from. [Table 2-4](#) includes code for how to install a popular driver for each RDBMS.

This is a one-time installation. The following code should be run in R.

Table 2-4. Install a driver for R

RDBMS	Code
SQLite	<code>install.packages("RSQLite")</code>
MySQL	<code>install.packages("RMySQL")</code>
Oracle	The <code>ROracle</code> package can be downloaded from the Oracle ROracle Downloads page. <code>setwd("folder_where_you_downloaded_ROracle")</code> <code># Update the name of the .zip file based on the latest version</code> <code>install.packages("ROracle_1.3-2.zip",</code> <code>repos=NULL)</code>

RDBMS	Code
PostgreSQL	<pre>install.packages("RPostgres")</pre>
SQL Server	<p>On Windows, the odbc (Open Database Connectivity) package is pre-installed. On macOS and Linux, it can be downloaded from the Microsoft ODBC page.</p> <pre>install.packages("odbc")</pre>

Step 2: Set up a database connection in R

To set up a database connection, you first need to know the location and name of the database you are trying to connect to, as well as your username and password. More details can be found in [“Database Connection Fields” on page 23](#).

Table 2-5 contains the R code you need to run each time you plan on writing SQL code in R. You can include it at the top of your R script.

Table 2-5. R code to set up a database connection

RDBMS	Code
SQLite	<pre>library(DBI) con <- dbConnect(RSQLite::SQLite(), "my_new_db.db")</pre>
MySQL	<pre>library(RMySQL) con <- dbConnect(RMySQL::MySQL(), host="localhost", dbname="my_new_db", user="alice", password="password")</pre>
Oracle	<pre>library(ROracle) drv <- dbDriver("Oracle") con <- dbConnect(drv, "alice", "password", dbname="my_new_db")</pre>

```
PostgreSQL library(RPostgres)
  con <- dbConnect(RPostgres::Postgres(),
                  host="localhost",
                  dbname="my_new_db",
                  user="alice",
                  password="password")

SQL Server library(DBI)
  con <- DBI::dbConnect(odbc::odbc(),
                       Driver="SQL Server",
                       Server="localhost\\SQLEXPRESS",
                       Database="my_new_db",
                       User="alice",
                       Password="password",
                       Trusted_Connection="True")
```

TIP

Not all arguments are required. If you exclude an argument completely, then the default value will be used.

- For example, the default host is *localhost*, which is your computer.
 - If no username and password were set up, then those arguments can be left out.
-

Keeping Your Passwords Safe in R

The preceding code is fine for testing out a connection to a database, but in reality, you should not be saving your password within a script for everyone to see.

There are multiple ways to avoid doing so, including:

- encrypting credentials with the `keyring` package

- creating a configuration file with the `config` package
- setting up environment variables with an `.Renviron` file
- recording the user and password as a global option in R with the `options` command

The recommended approach: prompt the user for a password.

The most straightforward approach, in my opinion, is to have RStudio prompt you for your password instead.

Instead of this:

```
con <- dbConnect(...,  
                  password="password",  
                  ...)
```

Do this:

```
install.packages("rstudioapi")  
con <- dbConnect(...,  
                  password=rstudioapi::askForPassword("Password?"),  
                  ...)
```

Step 3: Write SQL code in R

Once the database connection has been established, you can start writing SQL queries within your R code.

Show all tables in the database:

```
dbListTables(con)
```

```
[1] "test"
```

TIP

For *SQL Server*, include the schema name to limit the number of tables displayed—`dbListTables(con, schema="dbo")`. `dbo` stands for database owner and it is the default schema in SQL Server.

Take a look at the test table in the database:

```
dbReadTable(con, "test")
```

```
  id num
1  1 100
2  2 200
```

NOTE

For *Oracle*, the table name is case-sensitive. Since Oracle automatically converts table names to uppercase, you'll likely need to run the following instead: `dbReadTable(con, "TEST")`.

Write a simple query and output a dataframe:

```
df <- dbGetQuery(con, "SELECT * FROM test
                      WHERE id = 2")
print(df); class(df)

  id num
1  2 200
[1] "data.frame"
```

Close the connection when you are done using the database.

```
dbDisconnect(con)
```

It is always good practice to close the database connection to save resources.

The SQL Language

This chapter covers SQL fundamentals including its **standards**, **key terms**, and **sublanguages**, along with answers to the following questions:

- What is **ANSI SQL** and how is it different from SQL?
- What is a **keyword** versus a **clause**?
- Do **capitalization** and **whitespace** matter?
- What is there **beyond the SELECT statement**?

Comparison to Other Languages

Some people in the technology space don't consider SQL to be a real programming language.

While SQL stands for “Structured Query *Language*,” you can’t use it in the same way as some other popular programming languages like Python, Java, or C++. With those languages, you can write code to specify the exact steps that a computer should take to get a task done. This is called *imperative programming*.

In Python, if you want to sum up a list of values, you can tell the computer exactly *how* you want to do so. The following

example code goes through a list, item by item, and adds each value to a running total, to finally calculate the total sum:

```
calories = [90, 240, 165]
total = 0
for c in calories:
    total += c
print(total)
```

With SQL, instead of telling a computer exactly *how* you want to do something, you just describe *what* you want done, which in this case is to calculate the sum. Behind the scenes, SQL figures out how to optimally execute the code. This is called *declarative programming*.

```
SELECT SUM(calories)
FROM workouts;
```

The main takeaway here is that SQL is not a *general-purpose programming language* like Python, Java, or C++, which can be used for a variety of applications. Instead, SQL is a *special-purpose programming language*, specifically made for managing data in a relational database.

Extensions for SQL

At its core, SQL is a declarative language, but there are extensions that allow it to do more:

- Oracle has *procedural language SQL* (PL/SQL)
- SQL Server has *transact SQL* (T-SQL)

With these extensions, you can do things like group together SQL code into procedures and functions, and more. The syntax doesn't follow ANSI standards, but it makes SQL much more powerful.

ANSI Standards

The *American National Standards Institute* (ANSI) is an organization based in the United States that documents standards on everything from drinking water to nuts and bolts.

SQL became an ANSI standard in 1986. In 1989, they published a very detailed document of specifications (think hundreds of pages) on what a database language should be able to do and how it should be done. Every few years, the standards get updated, so that's why you'll hear terms like ANSI-89 and ANSI-92, which were different sets of SQL standards that were added in 1989 and 1992, respectively. The latest standard is ANSI SQL2016.

SQL Versus ANSI SQL Versus MySQL Versus ...

SQL is the general term for structured query language.

ANSI SQL refers to SQL code that follows the ANSI standards and will run in any relational database management system (RDBMS) software.

MySQL is one of many RDBMS options. Within MySQL, you can write both ANSI code and MySQL-specific SQL code.

Other RDBMS options include *Oracle*, *PostgreSQL*, *SQL Server*, *SQLite*, and others.

Even with the standards, no two RDBMSs are exactly the same. While some aim to be fully ANSI compliant, they are all just partially ANSI compliant. Each vendor ends up choosing which standards to implement and which additional features to build that only work within their software.

Should I follow the standards?

Most of the basic SQL code you write adheres to ANSI standards. If you find code that does something complex using simple yet unfamiliar keywords, then there's a good chance it's outside of the standards.

If you work solely within one RDBMS, like *Oracle* or *SQL Server*, it is absolutely fine to not follow the ANSI standards and take advantage of all of the features of the software.

The issue comes when you have code working in one RDBMS that you want to use in another RDBMS. Non-ANSI code likely won't run in the new RDBMS and would need to be rewritten.

Let's say you have the following query that works in *Oracle*. It does not meet ANSI standards because the `DECODE` function is only available within *Oracle* and not other software. If I copy the query over to *SQL Server*, the code will not run:

```
-- Oracle-specific code
SELECT item, DECODE (flag, 0, 'No', 1, 'Yes')
    AS Yes_or_No
FROM items;
```

The following query has the same logic, but uses a `CASE` statement instead, which is an ANSI standard. Because of this, it will work in *Oracle*, *SQL Server*, and other software:

```
-- Code that works in any RDBMS
SELECT item, CASE WHEN flag = 0 THEN 'No'
                  ELSE 'Yes' END AS Yes_or_No
FROM items;
```

Which Standard Should I Choose?

The following two code blocks perform a join using two different standards. ANSI-89 was the first widely adopted standard, followed by ANSI-92, which included some major revisions.

```
-- ANSI-89
SELECT c.id, c.name, o.date
FROM customer c, order o
WHERE c.id = o.id;

-- ANSI-92
SELECT c.id, c.name, o.date
FROM customer c INNER JOIN order o
ON c.id = o.id;
```

If you're writing new SQL code, I would recommend either using the latest standard (which is currently ANSI SQL2016) or the syntax provided in the documentation of the RDBMS you are working in.

However, it's important to be aware of the earlier standards because you will likely come across older code if your company has been around for a few decades.

SQL Terms

Here is a block of SQL code that shows the number of sales each employee closed in 2021. We'll be using this code block to highlight a number of SQL terms.

```
-- Sales closed in 2021
SELECT e.name, COUNT(s.sale_id) AS num_sales
FROM employee e
LEFT JOIN sales s ON e.emp_id = s.emp_id
WHERE YEAR(s.sale_date) = 2021
AND s.closed IS NOT NULL
GROUP BY e.name;
```

Keywords and Functions

Keywords and functions are terms built into SQL.

Keywords

A *keyword* is text that already has some meaning in SQL. All the keywords in the code block are bolded here:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales  
FROM employee e  
    LEFT JOIN sales s ON e.emp_id = s.emp_id  
WHERE YEAR(s.sale_date) = 2021  
    AND s.closed IS NOT NULL  
GROUP BY e.name;
```

SQL Is Case-Insensitive

Keywords are typically capitalized for readability. However, SQL is case-insensitive, meaning that an uppercase WHERE and a lowercase where mean the same thing when the code is run.

Functions

A *function* is a special type of keyword. It takes in zero or more inputs, does something to the inputs, and returns an output. In SQL, a function is usually followed by parentheses, but not always. The two functions in the code block are bolded here:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales  
FROM employee e  
    LEFT JOIN sales s ON e.emp_id = s.emp_id  
WHERE YEAR(s.sale_date) = 2021  
    AND s.closed IS NOT NULL  
GROUP BY e.name;
```

There are four categories of functions: numeric, string, date-time, and other:

- `COUNT()` is a numeric function. It takes in a column and returns the number of non-null rows (rows that have a value).
- `YEAR()` is a date function. It takes in a column of a date or datetime data type, extracts the years, and returns the values as a new column.

A list of common functions can be found in [Table 7-2](#).

Identifiers and Aliases

Identifiers and aliases are terms that the user defines.

Identifiers

An *identifier* is the name of a database object, such as a table or a column. All identifiers in the code block are bolded here:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales  
FROM employee e  
    LEFT JOIN sales s ON e.emp_id = s.emp_id  
WHERE YEAR(s.sale_date) = 2021  
    AND s.closed IS NOT NULL  
GROUP BY e.name;
```

Identifiers should start with a letter (a-z or A-Z), followed by any combination of letters, numbers, and underscores (_). Some software will allow additional characters such as @, #, and \$.

For readability's sake, identifiers are typically lowercase while keywords are uppercase, although the code will run regardless of case.

TIP

As a best practice, identifiers should not be given the same name as an existing keyword. For example, you wouldn't want to name a column COUNT because that is already a keyword in SQL.

If you still choose to do so, you can avoid confusion by enclosing the identifier in double quotes. So instead of naming a column COUNT, you can name it "COUNT", but it is best to use a completely different name altogether like num_sales.

MySQL uses backticks (``) to enclose identifiers instead of double quotes ("").

Aliases

An *alias* renames a column or a table temporarily, only for the duration of the query. In other words, the new alias names will be displayed in the results of the query, but the original column names will remain unchanged in the tables you are querying from. All the aliases in the code block are bolded here:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales
  FROM employee e
    LEFT JOIN sales s ON e.emp_id = s.emp_id
   WHERE YEAR(s.sale_date) = 2021
     AND s.closed IS NOT NULL
  GROUP BY e.name;
```

The standard is to use AS when renaming columns (AS num_sales) and no additional text when renaming tables (e). Technically, though, either syntax works for both columns and tables.

In addition to columns and tables, aliases are also useful if you'd like to temporarily name a **subquery**.

Statements and Clauses

These are ways to refer to subsets of SQL code.

Statements

A *statement* starts with a keyword and ends with a semicolon. This entire code block is called a SELECT statement because it starts with the keyword SELECT.

```
SELECT e.name, COUNT(s.sale_id) AS num_sales
FROM employee e
    LEFT JOIN sales s ON e.emp_id = s.emp_id
WHERE YEAR(s.sale_date) = 2021
    AND s.closed IS NOT NULL
GROUP BY e.name;
```

TIP

Many **database tools** that provide a graphical user interface do not require a semicolon (;) at the end of a statement.

The SELECT statement is the most popular type of SQL statement, and is often called a query instead because it finds data in a database. Other types of statements are covered in “[Sublanguages](#)” on page 50.

Clauses

A *clause* is a way to refer to a particular section of a statement. Here is our original SELECT statement:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales
FROM employee e
    LEFT JOIN sales s ON e.emp_id = s.emp_id
WHERE YEAR(s.sale_date) = 2021
    AND s.closed IS NOT NULL
GROUP BY e.name;
```

This statement contains four main clauses:

- SELECT clause

```
SELECT e.name, COUNT(s.sale_id) AS num_sales
```

- FROM clause

```
FROM employee e  
LEFT JOIN sales s ON e.emp_id = s.emp_id
```

- WHERE clause

```
WHERE YEAR(s.sale_date) = 2021  
AND s.closed IS NOT NULL
```

- GROUP BY clause

```
GROUP BY e.name;
```

In conversation, you'll often hear people refer to a section of a statement like "take a look at the tables in the `FROM` clause." It's a helpful way to zoom in on a particular section of the code.

NOTE

This statement actually has more clauses than the four listed. In grammar, a clause is a part of a sentence that contains a subject and a verb. So you could refer to the following:

```
LEFT JOIN sales s ON e.emp_id = s.emp_id
```

as the `LEFT JOIN` clause if you want to get even more specific about the section of the code that you are referring to.

The six most popular clauses start with `SELECT`, `FROM`, `WHERE`, `GROUP BY`, `HAVING`, and `ORDER BY` and are covered in detail in [Chapter 4](#).

Expressions and Predicates

These are combinations of **functions**, **identifiers**, and more.

Expressions

An *expression* can be thought of as a formula that results in a value. An expression in the code block was:

```
COUNT(s.sale_id)
```

This expression includes a function (COUNT) and an identifier (s.sale_id). Together, they make an expression that says to count the number of sales.

Other examples of expressions are:

- s.sale_id + 10 is a numeric expression that incorporates basic math operations.
- CURRENT_DATE is a datetime expression, simply a single function, that returns the current date.

Predicates

A *predicate* is a logical comparison that results in one of three values: TRUE/FALSE/UNKNOWN. They are sometimes called *conditional statements*. The three predicates in the code block are bolded here:

```
SELECT e.name, COUNT(s.sale_id) AS num_sales  
FROM employee e  
LEFT JOIN sales s ON e.emp_id = s.emp_id  
WHERE YEAR(s.sale_date) = 2021  
AND s.closed IS NOT NULL  
GROUP BY e.name;
```

Some things you'll notice from these examples are:

- The equal sign (=) is the most popular **operator** to compare values.

- The `NULL` stands for no value. When checking to see if a field has no value, instead of writing `= NULL`, you would write `IS NULL`.

Comments, Quotes, and Whitespace

These are punctuation marks with meaning in SQL.

Comments

A *comment* is text that is ignored when the code is run, like the following.

```
-- Sales closed in 2021
```

It is useful to include comments in your code so that other reviewers of your code (including your future self!) can quickly understand the intent of the code without reading all of it.

To comment out:

- A single line of text:

```
-- These are my comments
```

- Multiple lines of text:

```
/* These are  
my comments */
```

Quotes

There are two types of quotes you can use in SQL, the single quote and the double quote.

```
SELECT "This column"  
FROM my_table  
WHERE name = 'Bob';
```

Single Quotes: Strings

Take a look at 'Bob'. Single quotes are used when referring to a string value. You will see far more single quotes in practice compared to double quotes.

Double Quotes: Identifiers

Take a look at "This column". Double quotes are used when referring to an **identifier**. In this case, because there is a space in between This and column, the double quotes are necessary for This column to be interpreted as a column name. Without the double quotes, SQL would throw an error due to the space. That said, it is best practice to use _ instead of spaces when naming columns to avoid using the double quotes.

NOTE

MySQL uses backticks (``) to enclose identifiers instead of double quotes ("").

Whitespace

SQL does not care about the number of spaces between terms. Whether it's one space, a tab, or a new line, SQL will execute the query from the first keyword all the way to the semicolon at the end of the statement. The following two queries are equivalent.

```
SELECT * FROM my_table;  
  
SELECT *  
  FROM my_table;
```

NOTE

For simple SQL queries, you may see code all written on one line. For longer queries that have dozens or even hundreds of lines, you'll see new lines for new clauses, tabs when listing many columns or tables, etc.

The end goal is to have readable code, so you'll need to decide how you want to space out your code (or follow your company's guidelines) so that it looks clean and can be quickly skimmed.

Sublanguages

There are many types of statements that can be written within SQL. They all fall under one of five sublanguages, which are detailed in [Table 3-1](#).

Table 3-1. SQL sublanguages

Sublanguage	Description	Common Commands	Reference Sections
Data Query Language (DQL)	This is the language that most people are familiar with. These statements are used to retrieve information from a database object, such as a table, and are often referred to as SQL queries.	SELECT	The majority of this book is dedicated to DQL
Data Definition Language (DDL)	This is the language used to define or create a database object, such as a table or an index.	CREATE ALTER DROP	Creating, Updating, and Deleting
Data Manipulation Language (DML)	This is the language used to manipulate or modify data in a database.	INSERT UPDATE DELETE	Creating, Updating, and Deleting

Sublanguage	Description	Common Commands	Reference Sections
Data Control Language (DCL)	This is the language used to control access to data in a database, which are sometimes referred to as permissions or privileges.	GRANT REVOKE	Not covered
Transaction Control Language (TCL)	This is the language used to manage transactions in a database, or apply permanent changes to a database.	COMMIT ROLLBACK	Transaction Management

While most data analysts and data scientists will write DQL SELECT statements to query tables, it is important to know that database administrators and data engineers will also write code in these other sublanguages to maintain a database.

The SQL Language Summary

- ANSI SQL is standardized SQL code that works across all database software. Many RDBMSs have extensions that don't meet the standards but add functionality to their software.
- Keywords are terms that are reserved in SQL and have a special meaning.
- Clauses refer to particular sections of a statement. Common clauses are `SELECT`, `FROM`, `WHERE`, `GROUP BY`, `HAVING`, and `ORDER BY`.
- Capitalization and whitespace do not matter in SQL for execution, but there are best practices for readability.
- In addition to `SELECT` statements, there are commands for defining objects, manipulating data, and more.

Querying Basics

A *query* is a nickname for a SELECT statement, which consists of six main **clauses**. Each section of this chapter covers a clause in detail:

1. **SELECT**
2. **FROM**
3. **WHERE**
4. **GROUP BY**
5. **HAVING**
6. **ORDER BY**

The last section of this chapter covers the **LIMIT** clause, which is supported by *MySQL*, *PostgreSQL*, and *SQLite*.

The code examples in this chapter reference four tables:

`waterfall`

waterfalls in Michigan's Upper Peninsula

`owner`

owners of the waterfalls

county

counties where the waterfalls are located

tour

tours that consist of multiple waterfall stops

Here is a sample query that uses the six main clauses. It is followed by the query results, which are also known as the *result set*.

```
-- Tours with 2 or more public waterfalls
SELECT t.name AS tour_name,
       COUNT(*) AS num_waterfalls
FROM tour t LEFT JOIN waterfall w
  ON t.stop = w.id
WHERE w.open_to_public = 'y'
GROUP BY t.name
HAVING COUNT(*) >= 2
ORDER BY tour_name;
```

tour_name	num_waterfalls
M-28	6
Munising	6
US-2	4

```
-----
```

tour_name	num_waterfalls
M-28	6
Munising	6
US-2	4

To *query* a database means to retrieve data from a database, typically from a table or multiple tables.

NOTE

It is also possible to query a *view* instead of a table. Views look like tables and are derived from tables, but they themselves do not hold any data. More on views can be found in “Views” on page 133 in Chapter 5.

The SELECT Clause

The `SELECT` clause specifies the columns that you want a statement to return.

In the `SELECT` clause, the `SELECT` keyword is followed by a list of column names and/or **expressions** that are separated by commas. Each column name and/or expression then becomes a column in the results.

Selecting Columns

The simplest `SELECT` clause lists one or more column names from the tables in the `FROM` clause:

```
SELECT id, name  
FROM owner;  
  
id      name  
-----  
1 Pictured Rocks  
2 Michigan Nature  
3 AF LLC  
4 MI DNR  
5 Horseshoe Falls
```

Selecting All Columns

To return all columns from a table, you can use a single asterisk rather than write out each column name:

```
SELECT *  
FROM owner;  
  
id      name          phone        type  
-----  
1 Pictured Rocks   906.387.2607  public  
2 Michigan Nature  517.655.5655  private  
3 AF LLC           906.228.6561  public  
4 MI DNR           906.387.2635  private  
5 Horseshoe Falls
```

WARNING

The asterisk is a helpful shortcut when testing out queries because it can save you quite a bit of typing. However, it is risky to use the asterisk in production code because the columns in a table may change over time, causing your code to fail when there are fewer or more columns than expected.

Selecting Expressions

In addition to simply listing columns, you can also list more complex **expressions** within the `SELECT` clause to return as columns in the results.

The following statement includes an expression to calculate a 10% drop in population, rounded to zero decimal places:

```
SELECT name, ROUND(population * 0.9, 0)  
FROM county;
```

name	ROUND(population * 0.9, 0)
Alger	8876
Baraga	7871
Ontonagon	7036
...	

Selecting Functions

Expressions in the `SELECT` list typically refer to columns in the tables that you are pulling from, but there are exceptions. For example, a common function that doesn't refer to any tables is the one to return the current date:

```
SELECT CURRENT_DATE;
```

CURRENT_DATE
2021-12-01

The preceding code works in *MySQL*, *PostgreSQL*, and *SQLite*. Equivalent code that works in other RDBMSs can be found in “[Datetime Functions](#)” on page 218 in [Chapter 7](#).

NOTE

The majority of queries include both a `SELECT` and a `FROM` clause, but only the `SELECT` clause is required when using particular database functions, such as `CURRENT_DATE`.

It is also possible to include expressions within the `SELECT` clause that are *subqueries* (a query nested inside another query). More details can be found in “[Selecting Subqueries](#)” on page 61.

Aliasing Columns

The purpose of a *column alias* is to give a temporary name to any column or expression listed in the `SELECT` clause. That temporary name, or column alias, is then displayed as a column name in the results.

Note that this is not a permanent name change because the column names in the original tables remain unchanged. The alias only exists within the query.

This code displays three columns.

```
SELECT id, name,  
       ROUND(population * 0.9, 0)  
  FROM county;  
  
id      name      ROUND(population * 0.9, 0)  
-----  
2      Alger          8876  
6      Baraga         7871  
7      Ontonagon      7036  
...
```

Let's say we want to rename the column names in the results. `id` is too ambiguous and we'd like to give it a more descriptive name. `ROUND(population * 0.9, 0)` is too long and we'd like to give it a simpler name.

To create a column alias, you follow a column name or expression with either (1) an alias name or (2) the `AS` keyword and an alias name.

```
-- alias_name  
SELECT id county_id, name,  
       ROUND(population * 0.9, 0) estimated_pop  
FROM county;
```

or:

```
-- AS alias_name  
SELECT id AS county_id, name,  
       ROUND(population * 0.90, 0) AS estimated_pop  
FROM county;
```

county_id	name	estimated_pop
2	Alger	8876
6	Baraga	7871
7	Ontonagon	7036
...		

Both options are used in practice when creating aliases. Within the `SELECT` clause, the second option is more popular because the `AS` keyword makes it visually easier to differentiate column names and aliases among a long list of column names.

NOTE

Older versions of *PostgreSQL* require the use of `AS` when creating a column alias.

Although column aliases are not required, they are highly recommended when working with expressions to give sensible names to the columns in the results.

Aliases with case sensitivity and punctuation

As can be seen with the column aliases `county_id` and `estimated_pop`, the convention is to use lowercase letters with underscores in place of spaces when naming column aliases.

You can also create aliases containing uppercase letters, spaces, and punctuation using the double quote syntax, as shown in this example:

```
SELECT id AS "Waterfall #",
       name AS "Waterfall Name"
    FROM waterfall;

Waterfall #  Waterfall Name
-----
1 Munising Falls
2 Tannery Falls
3 Alger Falls
...
```

Qualifying Columns

Let's say you write a query that pulls data from two tables and they both contain a column called `name`. If you were to just include `name` in the `SELECT` clause, the code wouldn't know which table you were referring to.

To solve this problem, you can *qualify* a column name by its table name. In other words, you can give a column a prefix to specify which table it belongs to using *dot notation*, as in `table_name.column_name`.

The following example queries a single table, so while it isn't necessary to qualify the columns here, this is shown for demonstration's sake. This is how you would qualify a column by its table name:

```
SELECT owner.id, owner.name  
FROM owner;
```

TIP

If you get an error in SQL referencing an *ambiguous column name*, it means that multiple tables in your query have a column of the same name and you haven't specified which table/column combination you are referring to. You can resolve the error by qualifying the column name.

Qualifying tables

If you qualify a column name by its table name, you can also qualify that table name by its **database** or **schema** name. The following query retrieves data specifically from the `owner` table within the `sqlbook` schema:

```
SELECT sqlbook.owner.id, sqlbook.owner.name  
FROM sqlbook.owner;
```

The preceding code is lengthy since `sqlbook.owner` is repeated multiple times. To save on typing, you can provide a *table alias*. The following example gives the alias `o` to the table `owner`:

```
SELECT o.id, o.name  
FROM sqlbook.owner o;
```

or:

```
SELECT o.id, o.name  
FROM owner o;
```

Column Aliases Versus Table Aliases

Column aliases are defined within the `SELECT` clause to rename a column in the results. It is common to include `AS`, although not required.

```
-- Column alias  
SELECT num AS new_col  
FROM my_table;
```

Table aliases are defined within the FROM clause to create a temporary nickname for a table. It is common to exclude AS, although including AS also works.

```
-- Table alias  
SELECT *  
FROM my_table mt;
```

Selecting Subqueries

A *subquery* is a query that is nested inside another query. Subqueries can be located within various clauses, including the SELECT clause.

In the following example, in addition to the id, name, and population, let's say we also want to see the average population of all the counties. By including a subquery, we are creating a new column in the results for the average population.

```
SELECT id, name, population,  
       (SELECT AVG(population) FROM county)  
             AS average_pop  
FROM county;
```

id	name	population	average_pop
2	Alger	9862	18298
6	Baraga	8746	18298
7	Ontonagon	7818	18298
...			

A few things to note here:

- A subquery must be surrounded by parentheses.
- When writing a subquery within the SELECT clause, it is highly recommended that you specify a **column alias**,

which in this case was `average_pop`. That way, the column has a simple name in the results.

- There is only one value in the `average_pop` column that is repeated across all rows. When including a subquery within the `SELECT` clause, the result of the subquery must return a single column and either zero or one row, as shown in the following subquery to calculate the average population.

```
SELECT AVG(population) FROM county;
```

```
AVG(population)
```

```
-----
```

```
18298
```

- If the subquery returned zero rows, then the new column would be filled with `NULL` values.

Noncorrelated Versus Correlated Subqueries

The preceding example is a *noncorrelated subquery*, meaning that the subquery does not refer to the outer query. The subquery can be run on its own independent of the outer query.

The other type of subquery is called a *correlated subquery*, which is one that does refer to values in the outer query. This often significantly slows down processing time, so it's best to rewrite the query using a `JOIN` instead. What follows is an example of a correlated subquery along with more efficient code.

Performance issues with correlated subqueries

The following query returns the number of waterfalls for each owner. Note the `o.id = w.owner_id` step in the subquery references the `owner` table in the outer query, making it a correlated subquery.

```
SELECT o.id, o.name,
       (SELECT COUNT(*) FROM waterfall w
        WHERE o.id = w.owner_id) AS num_waterfalls
  FROM owner o;
```

id	name	num_waterfalls
1	Pictured Rocks	3
2	Michigan Nature	3
3	AF LLC	1
4	MI DNR	1
5	Horseshoe Falls	0

A better approach would be to rewrite the query with a `JOIN`. That way, the tables are first joined together and then the rest of the query is run, which is much faster than rerunning a subquery for each row of data. More on joins can be found in “Joining Tables” on page 270 in Chapter 9.

```
SELECT o.id, o.name,
       COUNT(w.id) AS num_waterfalls
  FROM owner o LEFT JOIN waterfalls w
    ON o.id = w.owner_id
 GROUP BY o.id, o.name
```

id	name	num_waterfalls
1	Pictured Rocks	3
2	Michigan Nature	3
3	AF LLC	1
4	MI DNR	1
5	Horseshoe Falls	0

DISTINCT

When a column is listed in the `SELECT` clause, by default, all of the rows are returned. To be more explicit, you can include the `ALL` keyword, but it is purely optional. The following queries return each `type/open_to_public` combination.

```
SELECT o.type, w.open_to_public
FROM owner o
JOIN waterfall w ON o.id = w.owner_id;
```

or:

```
SELECT ALL o.type, w.open_to_public
FROM owner o
JOIN waterfall w ON o.id = w.owner_id;
```

type	open_to_public
public	y
public	y
public	y
private	y
public	y

If you want to remove duplicate rows from the results, you can use the DISTINCT keyword. The following query returns a list of unique type/open_to_public combinations.

```
SELECT DISTINCT o.type, w.open_to_public
FROM owner o
JOIN waterfall w ON o.id = w.owner_id;
```

type	open_to_public
public	y
private	y

COUNT and DISTINCT

To count the number of unique values within a *single column*, you can combine the COUNT and DISTINCT keywords within the SELECT clause. The following query returns the number of unique type values.

```
SELECT COUNT(DISTINCT type) AS unique
FROM owner;
```

```
unique
```

```
-----
```

```
2
```

To count the number of unique combinations of *multiple columns*, you can wrap a DISTINCT query up as a **subquery**, and then do a COUNT on the subquery. The following query returns the number of unique type/open_to_public combinations.

```
SELECT COUNT(*) AS num_unique
FROM (SELECT DISTINCT o.type, w.open_to_public
      FROM owner o JOIN waterfall w
      ON o.id = w.owner_id) my_subquery;
```

```
num_unique
```

```
-----
```

```
2
```

MySQL and PostgreSQL support the use of the COUNT(DISTINCT) syntax on multiple columns. The following two queries are equivalent to the preceding query, without needing a subquery:

```
-- MySQL equivalent
SELECT COUNT(DISTINCT o.type, w.open_to_public)
      AS num_unique
      FROM owner o JOIN waterfall w
      ON o.id = w.owner_id;
```

```
-- PostgreSQL equivalent
```

```
SELECT COUNT(DISTINCT (o.type, w.open_to_public))
      AS num_unique
      FROM owner o JOIN waterfall w
      ON o.id = w.owner_id;
```

```
num_unique
```

```
-----
```

```
2
```

The FROM Clause

The `FROM` clause is used to specify the source of the data you want to retrieve. The simplest case is to name a single table or view in the `FROM` clause of query.

```
SELECT name  
FROM waterfall;
```

You can **qualify** a table or view with either a **database** or **schema** name using the dot notation. The following query retrieves data specifically from the `waterfall` table within the `sqlbook` schema:

```
SELECT name  
FROM sqlbook.waterfall;
```

From Multiple Tables

Instead of retrieving data from one table, you'll often want to pull together data from multiple tables. The most common way to do this is using a `JOIN` clause within the `FROM` clause. The following query retrieves data from both the `waterfall` and `tour` tables and displays a single results table.

```
SELECT *  
FROM waterfall w JOIN tour t  
ON w.id = t.stop;
```

id	name	...	name	stop	...
-----	-----	...	-----	-----	-----
1	Munising Falls		M-28		1
1	Munising Falls		Munising		1
2	Tannery Falls		Munising		2
3	Alger Falls		M-28		3
3	Alger Falls		Munising		3
...					

Let's break down each part of the code block.

Table aliases

```
waterfall w JOIN tour t
```

The waterfall and tour tables are given **table aliases** w and t, which are temporary names for the tables within the query. Table aliases are not required in a JOIN clause, but they are very helpful for shortening table names that need to be referenced within the ON and SELECT clauses.

JOIN ... ON ...

```
waterfall w JOIN tour t  
ON w.id = t.stop
```

These two tables are pulled together with the JOIN keyword. A JOIN clause is always followed by an ON clause, which specifies how the tables should be linked together. In this case, the id of the waterfall in the waterfall table must match the stop with the waterfall in the tour table.

NOTE

You may see the FROM, JOIN, and ON clauses on different lines or indented. This is not required, but helpful for readability's sake, especially when you are joining many tables together.

Results table

A query always results in a single table. The waterfall table has 12 columns and the tour table has 3 columns. After joining these tables together, the results table has 15 columns.

<code>id</code>	<code>name</code>	<code>...</code>	<code>name</code>	<code>stop</code>	<code>...</code>
1	Munising Falls		M-28	1	
1	Munising Falls		Munising	1	
2	Tannery Falls		Munising	2	
3	Alger Falls		M-28	3	
3	Alger Falls		Munising	3	
<code>...</code>					

You'll notice that there are two columns called `name` in the results table. The first is from the `waterfall` table, and the second is from the `tour` table. To refer to them in the `SELECT` clause, you would need to **qualify** the column names.

```
SELECT w.name, t.name
FROM waterfall w JOIN tour t
    ON w.id = t.stop;
```

<code>name</code>	<code>name</code>
Munising Falls	M-28
Munising Falls	Munising
Tannery Falls	Munising
<code>...</code>	

To differentiate the two columns, you would also want to **alias** the column names.

```
SELECT w.name AS waterfall_name,
      t.name AS tour_name
FROM waterfall w JOIN tour t
    ON w.id = t.stop;
```

<code>waterfall_name</code>	<code>tour_name</code>
Munising Falls	M-28
Munising Falls	Munising
Tannery Falls	Munising
Alger Falls	M-28
Alger Falls	Munising
<code>...</code>	

JOIN variations

In the preceding example, if a waterfall isn't listed in any tour, then it would not appear in the results table. If you wanted to see all waterfalls in the results, you would need to use a different type of join.

JOIN Defaults to INNER JOIN

This example uses a simple `JOIN` keyword to pull together data from two tables, although it is best practice to explicitly state the type of join you are using. `JOIN` on its own defaults to an **INNER JOIN**, meaning that only records that are in both tables are returned in the results.

There are a variety of join types used in SQL, which are covered in more detail in “[Joining Tables](#)” on page 270 in [Chapter 9](#).

From Subqueries

A subquery is a query that is nested inside another query. Subqueries within the `FROM` clause should be standalone `SELECT` statements, meaning that they do not reference the outer query at all and can be run on their own.

NOTE

A subquery within the `FROM` clause is also known as a *derived table* because the subquery ends up essentially acting like a table for the duration of the query.

The following query lists all publicly owned waterfalls, with the subquery portion bolded.

```
SELECT w.name AS waterfall_name,
       o.name AS owner_name
  FROM (SELECT * FROM owner WHERE type = 'public') o
        JOIN waterfall w
      ON o.id = w.owner_id;

waterfall_name  owner_name
-----
Little Miners    Pictured Rocks
Miners Falls     Pictured Rocks
Munising Falls   Pictured Rocks
Wagner Falls     MI DNR
```

It is important to understand the order in which the query is executed.

Step 1: Execute the subquery

The contents of the subquery are first executed. You can see that this results in a table of only public owners:

```
SELECT * FROM owner WHERE type = 'public';

id      name           phone      type
-----
1 Pictured Rocks  906.387.2607  public
4 MI DNR          906.228.6561  public
```

Going back to the original query, you'll notice that the subquery is immediately followed by the letter `o`. This is the temporary name, or alias, that we are assigning to the results of the subquery.

NOTE

Aliases are required for subqueries within the `FROM` clause in *MySQL*, *PostgreSQL*, and *SQL Server*, but not in *Oracle* and *SQLite*.

Step 2: Execute the entire query

Next, you can think of the letter `o` taking the place of the subquery. The query is now executed as usual.

```
SELECT w.name AS waterfall_name,
       o.name AS owner_name
  FROM o JOIN waterfall w
    ON o.id = w.owner_id;
```

waterfall_name	owner_name
Little Miners	Pictured Rocks
Miners Falls	Pictured Rocks
Munising Falls	Pictured Rocks
Wagner Falls	MI DNR

Subqueries Versus the WITH Clause

An alternative to writing a subquery is to write a common table expression (CTE) using a `WITH` clause instead. The advantage of the `WITH` clause is that the subquery is named up front, which makes for cleaner code and also the ability to reference the subquery multiple times.

```
WITH o AS (SELECT * FROM owner
            WHERE type = 'public')

SELECT w.name AS waterfall_name,
       o.name AS owner_name
  FROM o JOIN waterfall w
    ON o.id = w.owner_id;
```

The `WITH` clause is supported by *MySQL 8.0+* (2018 and later), *PostgreSQL*, *Oracle*, *SQL Server*, and *SQLite*. “[Common Table Expressions](#)” on page 291 in [Chapter 9](#) includes more examples of this technique.

Why Use a Subquery in the FROM Clause?

The main advantage of using subqueries is that you can turn a large problem into smaller ones. Here are two examples:

Example 1: Multiple steps to get to results

Let's say you want to find the average number of stops on a tour. First, you'd have to find the number of stops on each tour, and then average the results.

The following query finds the number of stops on each tour:

```
SELECT name, MAX(stop) as num_stops  
FROM tour  
GROUP BY name;
```

name	num_stops
M-28	11
Munising	6
US-2	14

You could then turn the query into a subquery and write another query around it to find the average:

```
SELECT AVG(num_stops) FROM  
(SELECT name, MAX(stop) as num_stops  
FROM tour  
GROUP BY name) tour_stops;  
  
AVG(num_stops)  
-----  
10.33333333333333
```

Example 2: Table in FROM clause is too large

The original goal was to list all publicly owned waterfalls. This can actually be done without a subquery and with a JOIN instead:

```
SELECT w.name AS waterfall_name,  
       o.name AS owner_name  
  FROM owner o  
        JOIN waterfall w ON o.id = w.owner_id  
 WHERE o.type = 'public';
```

```
waterfall_name  owner_name
-----
Little Miners  Pictured Rocks
Miners Falls   Pictured Rocks
Munising Falls Pictured Rocks
Wagner Falls   MI DNR
```

Let's say that the query takes a really long time to run. This can happen when you join massive tables together (think tens of millions of rows). There are multiple ways you could rewrite the query to speed it up, and one of them is to use a subquery.

Since we are only interested in public owners, we can first write a subquery that filters out all of the private owners. The smaller `owner` table would then be joined with the `waterfall` table, which would take less time and produce the same results.

```
SELECT w.name AS waterfall_name,
       o.name AS owner_name
  FROM (SELECT * FROM owner
        WHERE type = 'public') o
        JOIN waterfall w ON o.id = w.owner_id;
```

```
waterfall_name  owner_name
-----
Little Miners  Pictured Rocks
Miners Falls   Pictured Rocks
Munising Falls Pictured Rocks
Wagner Falls   MI DNR
```

These are just two of the many examples of how subqueries can be used to break down a larger query into smaller steps.

The WHERE Clause

The `WHERE` clause is used to restrict query results to only rows of interest, or simply put, it is the place to filter data. Rarely will you want to display all rows from a table, but rather rows that match specific criteria.

TIP

When exploring a table with millions of rows, you never want to do a `SELECT * FROM my_table;` because it will take an unnecessarily long time to run.

Instead, it's a good idea to filter down the data. Two common ways to do this are:

Filter by a column within the WHERE clause

Better yet, filter by a column that is already **indexed** to make the retrieval even faster.

```
SELECT *
  FROM my_table
 WHERE year_id = 2021;
```

*Show the top few rows of data with the LIMIT clause (or WHERE ROWNUM <= 10 in Oracle or SELECT TOP 10 * in SQL Server)*

```
SELECT *
  FROM my_table
 LIMIT 10;
```

The following query finds all waterfalls that do not contain *Falls* in the name. More on the `LIKE` keyword can be found in [Chapter 7](#).

```
SELECT id, name
  FROM waterfall
 WHERE name NOT LIKE '%Falls%';
```

id	name
7	Little Miners
14	Rapid River Fls

The bolded section is often referred to as a conditional statement or a **predicate**. The predicate makes a logical comparison for each row of data that results in TRUE/FALSE/UNKNOWN.

The waterfall table has 16 rows. For each row, it checks if the waterfall name contains *Falls* or not. If it doesn't contain *Falls*, then the name NOT LIKE '%Falls%' predicate is TRUE, and the row is returned in the results, which was the case for the two preceding rows.

Multiple Predicates

It is also possible to combine multiple predicates with *operators* like AND or OR. The following example shows waterfalls without *Falls* in its name and that also don't have an owner:

```
SELECT id, name
FROM waterfall
WHERE name NOT LIKE '%Falls%'
      AND owner_id IS NULL;

id      name
-----
14    Rapid River Fls
```

More details on operators can be found in [Operators in Chapter 7](#).

Filtering on Subqueries

A subquery is a query nested inside another query, and the WHERE clause is a common place to find one. The following example retrieves publicly accessible waterfalls located in Alger County:

```
SELECT w.name
FROM   waterfall w
WHERE  w.open_to_public = 'y'
      AND w.county_id IN (
          SELECT c.id FROM county c
          WHERE c.name = 'Alger');
```

```
name
-----
Munising Falls
Tannery Falls
Alger Falls
...
```

NOTE

Unlike subqueries within the SELECT clause or the FROM clause, subqueries in the WHERE clause do not require an alias. In fact, you will get an error if you include an alias.

Why use a subquery in the WHERE clause?

The original goal was to retrieve publicly accessible waterfalls located in Alger County. If you were to write this query from scratch, you would likely start with the following:

```
SELECT w.name
FROM   waterfall w
WHERE  w.open_to_public = 'y';
```

At this point, you have all waterfalls that are publicly accessible. The final touch is to find ones that are specifically in Alger County. You know that the waterfall table doesn't have a county name column, but the county table does.

You have two options to pull the county name into the results. You can either (1) write a subquery within the WHERE clause that specifically pulls the Alger County information or (2) join together the waterfall and county tables:

```
-- Subquery in WHERE clause
SELECT w.name
FROM   waterfall w
WHERE  w.open_to_public = 'y'
      AND w.county_id IN (
```

```
SELECT c.id FROM county c  
WHERE c.name = 'Alger');
```

or:

```
-- JOIN clause  
SELECT w.name  
FROM waterfall w INNER JOIN county c  
ON w.county_id = c.id  
WHERE w.open_to_public = 'y'  
AND c.name = 'Alger';
```

```
name  
-----  
Munising Falls  
Tannery Falls  
Alger Falls  
...
```

The two queries produce the same results. The advantage of the first approach is that subqueries are often easier to understand than joins. The advantage of the second approach is that joins typically execute faster than subqueries.

Working > Optimizing

When writing SQL code, there are often multiple ways to do the same thing.

Your top priority should be to write *working* code. If it takes a long time to run or it's ugly, it doesn't matter...it works!

The next step, if you have time, is to *optimize* the code by improving the performance by perhaps rewriting it with a JOIN, making it more readable with indentations and capitalizations, etc.

Don't stress about writing the most optimized code up front, but rather writing code that works. Writing elegant code comes with experience.

Other Ways to Filter Data

The WHERE clause is not the only place within a SELECT statement to filter rows of data.

- FROM clause: When joining together tables, the ON clause specifies how they should be linked together. This is where you can include conditions to restrict rows of data returned by the query. See [Joining Tables in Chapter 9](#) for more details.
- HAVING clause: If there are aggregations within the SELECT statement, the HAVING clause is where you specify how the aggregations should be filtered. See [“The HAVING Clause” on page 83](#) for more details.
- LIMIT clause: To display a specific number of rows, you can use the LIMIT clause. In *Oracle*, this is done with WHERE ROWNUM and in *SQL Server*, this is done with SELECT TOP. See [“The LIMIT Clause” on page 88](#) in this chapter for more details.

The GROUP BY Clause

The purpose of the GROUP BY clause is to collect rows into groups and summarize the rows within the groups in some way, ultimately returning just one row per group. This is sometimes referred to as “slicing” the rows into groups and “rolling up” the rows in each group.

The following query counts the number of waterfalls along each of the tours:

```
SELECT t.name AS tour_name,
       COUNT(*) AS num_waterfalls
  FROM waterfall w INNER JOIN tour t
    ON w.id = t.stop
 GROUP BY t.name;
```

tour_name num_waterfalls

M-28	6
Munising	6
US-2	4

There are two parts to focus on here:

- *The collecting of rows*, which is specified within the GROUP BY clause
- *The summarizing of rows* within groups, which is specified within the SELECT clause

Step 1: The collecting of rows

In the GROUP BY clause:

```
GROUP BY t.name
```

we state that we would like to look at all of the rows of data and put the M-28 tour waterfalls into a group, all of the Munising tour waterfalls into a group, and so on. Behind the scenes, the data is being grouped like this:

tour_name	waterfall_name
M-28	Munising Falls
M-28	Alger Falls
M-28	Scott Falls
M-28	Canyon Falls
M-28	Agate Falls
M-28	Bond Falls
Munising	Munising Falls
Munising	Tannery Falls
Munising	Alger Falls
Munising	Wagner Falls
Munising	Horseshoe Falls
Munising	Miners Falls

US-2	Bond Falls
US-2	Fumee Falls
US-2	Kakabika Falls
US-2	Rapid River Fls

Step 2: The summarizing of rows

In the SELECT clause,

```
SELECT t.name AS tour_name,  
       COUNT(*) AS num_waterfalls
```

we state that for each group, or each tour, we want to count the number of rows of data in the group. Because each row represents a waterfall, this would result in the total number of waterfalls along each tour.

The COUNT() function here is more formally known as an *aggregate function*, or a function that summarizes many rows of data into a single value. More aggregate functions can be found in “Aggregate Functions” on page 191 in Chapter 7.

WARNING

In this example, COUNT(*) returns the number of waterfalls along each tour. However, this is only because each row of data in the waterfall and tour tables represent a single waterfall.

If a single waterfall was listed on multiple rows, COUNT(*) would result in a larger value than expected. In this case, you could potentially use COUNT(DISTINCT waterfall_name) instead to find the unique waterfalls. More details can be found in [COUNT and DISTINCT](#).

The key takeaway is that it is important to manually double-check the results of the aggregate function to make sure it is summarizing the data in the way that you intended.

Now that the groups have been created with the GROUP BY clause, the aggregate function will be applied once to each group:

tour_name	COUNT(*)
M-28	6
M-28	
Munising	6
Munising	
US-2	4
US-2	
US-2	
US-2	

Any columns to which an aggregate function has not been applied, which in this case is the tour_name column, are now collapsed into one value:

tour_name	COUNT(*)
M-28	6
Munising	6
US-2	4

NOTE

This collapsing of many detail rows into one aggregate row means that when using a GROUP BY clause, the SELECT clause should *only* contain:

- All columns listed in the GROUP BY clause: t.name
- Aggregations: COUNT(*)

```
SELECT t.name AS tour_name,  
       COUNT(*) AS num_waterfalls  
...  
GROUP BY t.name;
```

Not doing so could either result in an error message or return inaccurate values.

GROUP BY In Practice

These are the steps you should take when using a GROUP BY:

1. Figure out what column(s) you want to use to separate out, or group, your data (i.e., tour name).
2. Figure out how you'd like to summarize the data within each group (i.e. count the waterfalls within each tour).

When you've decided on those:

1. In the SELECT clause, list the column(s) you want to group by (i.e., tour name) and the aggregation(s) you want to calculate within each group (i.e., count of waterfalls).
2. In the GROUP BY clause, list all columns that are not aggregations (i.e., tour name).

For more complex grouping situations including ROLLUP, CUBE, and GROUPING SETS, go to “[Grouping and Summarizing](#)” on page 242 in Chapter 8.

The HAVING Clause

The HAVING clause places restrictions on the rows returned from a GROUP BY query. In other words, it allows you to filter on the results after a GROUP BY has been applied.

NOTE

A HAVING clause always immediately follows a GROUP BY clause. Without a GROUP BY clause, there can be no HAVING clause.

This is a query that lists the number of waterfalls on each tour using a GROUP BY clause:

```
SELECT      t.name AS tour_name,
            COUNT(*) AS num_waterfalls
  FROM      waterfall w INNER JOIN tour t
            ON w.id = t.stop
GROUP BY  t.name;
```

tour_name	num_waterfalls
M-28	6
Munising	6
US-2	4

Let's say we only want to list the tours that have exactly six stops. To do so, you would add a HAVING clause after the GROUP BY clause:

```
SELECT      t.name AS tour_name,
            COUNT(*) AS num_waterfalls
  FROM      waterfall w INNER JOIN tour t
            ON w.id = t.stop
GROUP BY  t.name
HAVING    COUNT(*) = 6;
```

tour_name	num_waterfalls
M-28	6
Munising	6

WHERE Versus HAVING

The purpose of both clauses is to filter data. If you are trying to:

- Filter on particular columns, write your conditions within the WHERE clause
- Filter on aggregations, write your conditions within the HAVING clause

The contents of a WHERE and HAVING clause cannot be swapped:

- Never put a condition with an aggregation in the WHERE clause. You will get an error.
- Never put a condition in the HAVING clause that does not involve an aggregation. Those conditions are evaluated much more efficiently in the WHERE clause.

You'll notice that the HAVING clause refers to the aggregation COUNT(*),

```
SELECT COUNT(*) AS num_waterfalls  
...  
HAVING COUNT(*) = 6;
```

and not the alias,

```
# code will not run  
SELECT COUNT(*) AS num_waterfalls  
...  
HAVING num_waterfalls = 6;
```

The reason for this is because of the **order of execution** of the clauses. The SELECT clause is written before the HAVING clause. However, the SELECT clause is actually executed *after* the HAVING clause.

That means that the alias `num_waterfalls` in the `SELECT` clause does not exist at the time the `HAVING` clause is being executed. The `HAVING` clause must refer to the raw aggregation `COUNT(*)` instead.

NOTE

MySQL and *SQLite* are exceptions, and allow aliases (`num_waterfalls`) in the `HAVING` clause.

The ORDER BY Clause

The `ORDER BY` clause is used to specify how you want the results of a query to be sorted.

The following query returns a list of owners and waterfalls, without any sorting:

```
SELECT COALESCE(o.name, 'Unknown') AS owner,
       w.name AS waterfall_name
  FROM waterfall w
 LEFT JOIN owner o ON w.owner_id = o.id;
```

owner	waterfall_name
Pictured Rocks	Munising Falls
Michigan Nature	Tannery Falls
AF LLC	Alger Falls
MI DNR	Wagner Falls
Unknown	Horseshoe Falls
...	

The COALESCE Function

The COALESCE function replaces all NULL values in a column with a different value. In this case, it turned the NULL values in the `o.name` column into the text Unknown.

If the COALESCE function were not used here, all waterfalls without owners would have been left out of the results. Instead, they are now marked as having an Unknown owner, and can be sorted on and included in the results.

More details can be found in [Chapter 7](#).

The following query returns the same list, but first sorted alphabetically by owner, and then by waterfall:

```
SELECT COALESCE(o.name, 'Unknown') AS owner,
       w.name AS waterfall_name
  FROM waterfall w
 LEFT JOIN owner o ON w.owner_id = o.id
 ORDER BY owner, waterfall_name;
```

owner	waterfall_name
AF LLC	Alger Falls
MI DNR	Wagner Falls
Michigan Nature	Tannery Falls
Michigan Nature	Twin Falls #1
Michigan Nature	Twin Falls #2
...	

The default sort is in ascending order, meaning text will go from A to Z and numbers will go from lowest to highest. You can use the keywords ASCENDING and DESCENDING (which can be abbreviated as ASC and DESC) to control the sort on each column.

The following is a modification of the previous sort, but this time, it sorts owner names in reverse order:

```
SELECT COALESCE(o.name, 'Unknown') AS owner,  
       w.name AS waterfall_name  
     ...  
ORDER BY owner DESC, waterfall_name ASC;
```

owner	waterfall_name
-----	-----
Unknown	Agate Falls
Unknown	Bond Falls
Unknown	Canyon Falls
...	

You can sort by columns and expressions that are not in your SELECT list:

```
SELECT COALESCE(o.name, 'Unknown') AS owner,  
       w.name AS waterfall_name  
  FROM waterfall w  
    LEFT JOIN owner o ON w.owner_id = o.id  
ORDER BY o.id DESC, w.id;
```

owner	waterfall_name
-----	-----
MI DNR	Wagner Falls
AF LLC	Alger Falls
Michigan Nature	Tannery Falls
...	

You can also sort by numeric column position:

```
SELECT COALESCE(o.name, 'Unknown') AS owner,  
       w.name AS waterfall_name  
     ...  
ORDER BY 1 DESC, 2 ASC;
```

owner	waterfall_name
-----	-----
Unknown	Agate Falls
Unknown	Bond Falls

...

Because the rows of a SQL table are unordered, if you don't include an `ORDER BY` clause in a query, each time you execute the query, the results could be displayed in a different order.

ORDER BY Cannot Be Used in a Subquery

Of the six main clauses, only the `ORDER BY` clause cannot be used in a subquery. Unfortunately, you can't force the rows of a subquery to be ordered.

To avoid this issue, you would need to rewrite your query with different logic to avoid using an `ORDER BY` clause within the subquery, and only include an `ORDER BY` clause in the outer query.

The LIMIT Clause

When quickly viewing a table, it is best practice to return a limited number of rows instead of the entire table.

MySQL, PostgreSQL, and SQLite support the `LIMIT` clause. *Oracle* and *SQL Server* use different syntax with the same functionality:

```
-- MySQL, PostgreSQL, and SQLite
SELECT *
FROM owner
LIMIT 3;

-- Oracle
SELECT *
FROM owner
WHERE ROWNUM <= 3;

-- SQL Server
SELECT TOP 3 *
FROM owner;
```

<code>id</code>	<code>name</code>	<code>phone</code>	<code>type</code>

1	Pictured Rocks	906.387.2607	public
2	Michigan Nature	517.655.5655	private
3	AF LLC		private

Another way to limit the number of rows returned is to filter on a column within the **WHERE** clause. The filtering will execute even faster if the column is **indexed**.

Creating, Updating, and Deleting

The majority of this book covers how to read data from a database with SQL queries. Reading is one of the four basic database operations out of create, read, update, and delete (CRUD).

This chapter focuses on the remaining three operations for **Databases**, **Tables**, **Indexes**, and **Views**. In addition, the **Transaction Management** section covers how to execute multiple commands as a single unit.

Databases

A *database* is a place to store data in an organized way.

Within a database, you can create *database objects*, which are things that store or reference data. Common database objects include **tables**, **constraints**, **indexes**, and **views**.

A *data model* or a *schema* describes how database objects are organized within a database.

Figure 5-1 shows a database that contains many tables. The specifics around how the tables are defined (i.e., the `Sales` table contains five columns) and how they connect with one another (i.e., the `customer_id` column in the `Sales` table matches the

`customer_id` column in the `Customer` table) are all a part of the *schema* of the database.

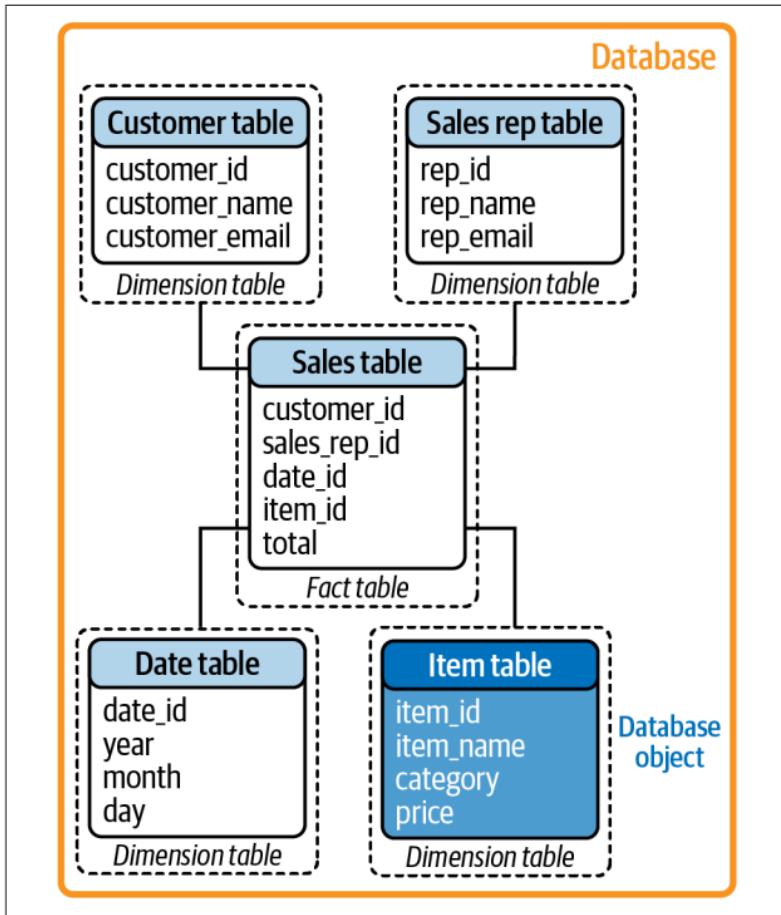


Figure 5-1. A database containing a star schema

The tables in [Figure 5-1](#) are arranged in a *star schema*, which is a basic way of organizing tables in a database. The star schema includes a *fact table* in the center and is surrounded by *dimension tables* (also known as *lookup tables*). The fact table records transactions made (sales in this case) along with IDs of additional information, which are fully detailed out in the dimension tables.

Data Model Versus Schema

When designing a database, you would first come up with a *data model*, which is how you want your database organized at a high level. It could look like [Figure 5-1](#) and include table names, how they are connected to one another, etc.

When you are ready to take action, you would create a *schema*, which is the implementation of the data model in a database. Within the software you are working in, you would specify the tables, constraints, primary and foreign keys, etc.

NOTE

The definition of a schema varies for some RDBMSs.

In *MySQL*, a schema is the same thing as a database and the two terms can be used interchangeably.

In *Oracle*, a schema consists of the database objects owned by a particular user, so the terms *schema* and *user* are used interchangeably.

Display Names of Existing Databases

All database objects reside in databases, so a good first step is to see what databases currently exist. [Table 5-1](#) shows the code to display the names of all existing databases in each RDBMS.

Table 5-1. Code to display names of existing databases

RDBMS	Code
MySQL	SHOW databases;
Oracle	SELECT * FROM global_name;
PostgreSQL	\l
SQL Server	SELECT name FROM master.sys.databases;
SQLite	.database (or look in the file browser for .db files)

NOTE

SQLite: For most RDBMS software, databases are located within the RDBMS. However, for SQLite, databases are stored outside of SQLite as *.db* files. To use a database, you would specify a *.db* file name when launching SQLite:

```
> sqlite3 existing_db.db
```

Display Name of Current Database

You may want to confirm the database you are currently in before writing any queries. **Table 5-2** shows the code to display the name of the database you are currently in for each RDBMS.

Table 5-2. Code to display name of current database

RDBMS	Code
MySQL	SELECT database();
Oracle	SELECT * FROM global_name;
PostgreSQL	SELECT current_database();
SQL Server	SELECT db_name();
SQLite	.database

NOTE

You may have noticed that the current database code is the same as the existing database code for Oracle and SQLite.

An instance of *Oracle* can only connect to a single database at a time, and you typically don't switch databases.

With *SQLite*, you can only open up and work with a single database file at a time.

Switch to Another Database

You may want to use data in another database or switch to a newly created database. **Table 5-3** shows the code to switch to another database in each RDBMS.

Table 5-3. Code to switch to another database

RDBMS	Code
MySQL, SQL Server	USE another_db;
Oracle	You typically don't switch databases (see earlier note), but to switch users, you would type: connect another_user
PostgreSQL	\c another_db
SQLite	.open another_db

Create a Database

If you have CREATE privileges, you can create a new database. If not, you may only be able to work within an existing database. **Table 5-4** shows the code to create a database in each RDBMS.

Table 5-4. Code to create a database

RDBMS	Code
MySQL, Oracle, PostgreSQL, SQL Server	CREATE DATABASE my_new_db;
SQLite	> sqlite3 my_new_db.db

NOTE

Oracle: There are some additional steps (regarding instances, environment variables, etc.) surrounding the CREATE DATABASE statement in Oracle, which can be found in the Oracle [documentation](#).

SQLite: The > symbol is not a character that you actually type. It just signifies that this is command line code, not SQL code.

Delete a Database

If you have DELETE privileges, you can delete a database. [Table 5-5](#) shows the code to delete a database in each RDBMS.

WARNING

If you delete a database, you will lose all of the data in the database. *There is no undo*, unless a backup has been created. I recommend not running this command unless you are 100% sure you don't need the database.

Table 5-5. Table 5-5. Code to delete a database

RDBMS	Code
MySQL,	DROP DATABASE my_new_db;
Oracle,	
PostgreSQL,	
SQL Server	
SQLite	Delete the .db file in the file browser

NOTE

Oracle: There are some additional steps (regarding mounting, etc.) surrounding the `DROP DATABASE` statement in Oracle, which can be found in the Oracle [documentation](#).

In some RDBMSs, you can't drop a database you are currently in. You would have to first switch to another database, such as the default one, before dropping the database:

- In *PostgreSQL*, the default database is `postgres`:

```
\c postgres
DROP DATABASE my_new_db;
```

- In *SQL Server*, the default database is `master`:

```
USE master;
go
DROP DATABASE my_new_db;
go
```

Creating Tables

Tables consist of rows and columns, and store all of the data in a database. In SQL, there are a few additional requirements for tables:

- Each row of a table should be unique
- All data in a column must be of the same data type (integer, text, etc.)

NOTE

In *SQLite*, the data in a column *does not* have to all be of the same data type. SQLite is more flexible in that each value has a data type associated with it, rather than an entire column.

To be compatible with other RDBMSs, SQLite does support columns having data type assignments. These *type affinities* are recommended data types for the columns, and are not required.

Create a Simple Table

It takes two steps to create a table in SQL. You must first define the structure of a table before loading data into it:

1. Create a table.

The following code creates an empty table called `my_simple_table` with three columns: `id`, `country`, and `name`. All values in the first column (`id`) must be integers, and the other two columns (`country`, `name`) can contain up to 2 and up to 15 characters:

```
CREATE TABLE my_simple_table (
    id INTEGER,
    country VARCHAR(2),
    name VARCHAR(15)
);
```

More data types in addition to `INTEGER` and `VARCHAR` are listed in [Chapter 6](#).

2. Insert rows.

- a. Insert a single row of data.

The following code inserts one row of data into columns `id`, `country`, and `name`:

```
INSERT INTO my_simple_table (id, country, name)
VALUES (1, 'US', 'Sam');
```

- b. Insert multiple rows of data.

Table 5-6 shows how to insert multiple rows of data into a table in each RDBMS, instead of one row at a time.

Table 5-6. Code to insert multiple rows of data

RDBMS	Code
MySQL,	INSERT INTO my_simple_table
PostgreSQL,	(id, country, name)
SQL Server,	VALUES (2, 'US', 'Selena'),
SQLite	(3, 'CA', 'Shawn'),
	(4, 'US', 'Sutton');
Oracle	INSERT ALL
	INTO my_simple_table (id, country, name)
	VALUES (2, 'US', 'Selena')
	INTO my_simple_table (id, country, name)
	VALUES (3, 'CA', 'Shawn')
	INTO my_simple_table (id, country, name)
	VALUES (4, 'US', 'Sutton')
	SELECT * FROM dual;

After inserting the data, the table would look like this:

```
SELECT * FROM my_simple_table;
```

id	country	name
1	US	Sam
2	US	Selena
3	CA	Shawn
4	US	Sutton

When inserting rows of data, the order of the values must match the order of the column names exactly.

Values in any columns omitted from the column list will take on their default value of `NULL`, unless another **default value** is specified.

NOTE

You need CREATE privileges to create a table. If you get an error when running the preceding code, you do not have the permission to do so and need to talk to your database administrator.

Display Names of Existing Tables

Before creating a table, you may want to see if the table name already exists. **Table 5-7** shows the code to display the names of existing tables in the database for each RDBMS.

Table 5-7. Code to display names of existing tables

RDBMS	Code
MySQL	SHOW tables;
Oracle	-- All tables, including system tables SELECT table_name FROM all_tables; -- All user created tables SELECT table_name FROM user_tables;
PostgreSQL	\dt
SQL Server	SELECT table_name FROM information_schema.tables;
SQLite	.tables

Create a Table That Does Not Already Exist

In *MySQL*, *PostgreSQL*, and *SQLite*, you can check for existing tables using the `IF NOT EXISTS` keywords when creating a table:

```
CREATE TABLE IF NOT EXISTS my_simple_table (
    id INTEGER,
    country VARCHAR(2),
    name VARCHAR(15)
);
```

If the table name does not exist, a new table will get created. If the table name already exists, without the IF NOT EXISTS, you would get an error message. With the IF NOT EXISTS, no new table gets created and you would avoid an error message.

If you want to replace an existing table, there are two approaches to doing so:

- You could use **DROP TABLE** to completely delete the existing table, and then create a new one.
- You could *truncate* the existing table, meaning you keep the schema (aka structure) of the table, but clear out the data inside of it. This can be done by using **DELETE FROM** to delete data from the table.

Create a Table with Constraints

A *constraint* is a rule that specifies what data can be inserted into a table. The following code creates two tables and multiple constraints (bolded):

```
CREATE TABLE another_table (
    country VARCHAR(2) NOT NULL,
    name VARCHAR(15) NOT NULL,
    description VARCHAR(50),
    CONSTRAINT pk_another_table
        PRIMARY KEY (country, name)
);

CREATE TABLE my_table (
    id INTEGER NOT NULL,
    country VARCHAR(2) DEFAULT 'CA'
        CONSTRAINT chk_country
            CHECK (country IN ('CA','US')),
    name VARCHAR(15),
    cap_name VARCHAR(15),
    CONSTRAINT pk
        PRIMARY KEY (id),
    CONSTRAINT fk1
```

```
    FOREIGN KEY (country, name)
        REFERENCES another_table (country, name),
    CONSTRAINT unq_country_name
        UNIQUE (country, name),
    CONSTRAINT chk_upper_name
        CHECK (cap_name = UPPER(name))
);
```

The `CONSTRAINT` keyword names the constraint for future reference and is optional. You should avoid using the same name for both a column and a constraint.

For quick access to the constraint sections: `NOT NULL`, `DEFAULT`, `CHECK`, `UNIQUE`, `PRIMARY KEY`, `FOREIGN KEY`.

Constraint: Not allowing NULL values in a column with NOT NULL

In a SQL table, cells without a value are replaced with the term `NULL`. For each column, you can specify whether `NULL` values are allowed or not:

```
CREATE TABLE my_table (
    id INTEGER NOT NULL,
    country VARCHAR(2) NULL,
    name VARCHAR(15)
);
```

The `NOT NULL` constraint on the `id` column means that the column will not allow `NULL` values. In other words, there can be no missing values inserted into the column, or else you will get an error message.

The `NULL` constraint on the `country` column means that the column will allow `NULL` values. If you are inserting data into the table and exclude the `country` column, then no value will be inserted, and the cell will be replaced with a `NULL` value.

By not specifying `NULL` or `NOT NULL`, the `name` column defaults to `NULL`, meaning it will allow `NULL` values.

Constraint: Setting default values in a column with DEFAULT

When inserting data into a table, missing values get replaced with the term NULL. To replace missing values with another value, you can use the DEFAULT constraint. The following code turns any missing country value into CA:

```
CREATE TABLE my_table (
    id INTEGER,
    country VARCHAR(2) DEFAULT 'CA',
    name VARCHAR(15)
);
```

Constraint: Restricting values in a column with CHECK

You can restrict the values allowed in a column by using the CHECK constraint. The following code only allows values of CA and US in the country column.

You can place the CHECK keyword immediately after the column name and data type:

```
CREATE TABLE my_table (
    id INTEGER,
    country VARCHAR(2) CHECK
        (country IN ('CA', 'US')),
    name VARCHAR(15)
);
```

Or you can place the CHECK keyword after all the column names and data types:

```
CREATE TABLE my_table (
    id INTEGER,
    country VARCHAR(2),
    name VARCHAR(15),
    CHECK (country IN ('CA','US'))
);
```

You can also include logic that checks multiple columns:

```
CREATE TABLE my_table (
    id INTEGER,
    country VARCHAR(2),
```

```
    name VARCHAR(15),
CONSTRAINT chk_id_country
  CHECK (id > 100 AND country IN ('CA','US'))
);
```

Constraint: Requiring unique values in a column with UNIQUE

You can require the values of a column to be unique by using the **UNIQUE** constraint.

You can place the **UNIQUE** keyword immediately after the column name and data type:

```
CREATE TABLE my_table (
  id INTEGER UNIQUE,
  country VARCHAR(2),
  name VARCHAR(15)
);
```

Or you can place the **UNIQUE** keyword after all the column names and data types:

```
CREATE TABLE my_table (
  id INTEGER,
  country VARCHAR(2),
  name VARCHAR(15),
  UNIQUE (id)
);
```

You can also include logic that forces the combination of multiple columns to be unique. The following code requires unique **country/name** combinations, meaning that one row can include CA/Emma and another can include US/Emma:

```
CREATE TABLE my_table (
  id INTEGER,
  country VARCHAR(2),
  name VARCHAR(15),
  CONSTRAINT unq_country_name
  UNIQUE (country, name)
);
```

Create a Table with Primary and Foreign Keys

Primary keys and foreign keys are special types of constraints that uniquely identify rows of data.

Specify a primary key

A *primary key* uniquely identifies each row of data in a table. A primary key can be made up of one or more columns in a table. Every table should have a primary key.

You can place the PRIMARY KEY keywords immediately after the column name and data type:

```
CREATE TABLE my_table (
    id INTEGER PRIMARY KEY,
    country VARCHAR(2),
    name VARCHAR(15)
);
```

Or you can place the PRIMARY KEY keywords after all the column names and data types:

```
CREATE TABLE my_table (
    id INTEGER,
    country VARCHAR(2),
    name VARCHAR(15),
    PRIMARY KEY (id)
);
```

To specify a primary key consisting of multiple columns (also known as a *composite key*):

```
CREATE TABLE my_table (
    id INTEGER NOT NULL,
    country VARCHAR(2),
    name VARCHAR(15) NOT NULL,
    CONSTRAINT pk_id_name
    PRIMARY KEY (id, name)
);
```

By creating a PRIMARY KEY, the constraints that you are putting on the column(s) are that they cannot include NULL values (**NOT NULL**) and the values must be unique (**UNIQUE**).

Primary Key Best Practices

Every table should have a primary key. This ensures that every row can be uniquely identified.

It is recommended that primary keys consist of ID columns, like (country_id, name_id) instead of (country, name). Technically, multiple rows could have the same country and name combination. By adding columns that contain unique IDs (101, 102, etc.), the combination of country_id and name_id is guaranteed to be unique.

Primary keys should be immutable, meaning that they can't be changed. This allows for a particular row in a table to always be identified by the same primary key.

Specify a foreign key

A *foreign key* in a table refers to a primary key in another table. The two tables can be linked together by the common column. A table can have zero or more foreign keys.

Figure 5-2 shows a data model of two tables: the customers table, which has a primary key of id, and the orders table, which has a primary key of o_id. From the viewpoint of customers, its order_id column matches with values of the o_id column, making order_id a foreign key because it refers to a primary key in another table.

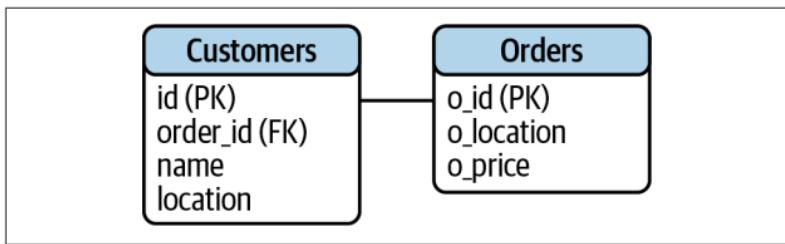


Figure 5-2. Two tables with primary keys and a foreign key

To specify a foreign key, use the following steps:

1. Locate the table you plan to reference and identify the primary key.

In this case, we will be referencing `orders`, specifically the `o_id` column:

```
CREATE TABLE orders (
    o_id INTEGER PRIMARY KEY,
    o_location VARCHAR(20),
    o_price DECIMAL(6,2)
);
```

2. Create a table with a foreign key that references the primary key in the other table.

In this case, we are creating the `customers` table where the `order_id` column references the `o_id` primary key in the `orders` table:

```
CREATE TABLE customers (
    id INTEGER PRIMARY KEY,
    order_id INTEGER,
    name VARCHAR(15),
    location VARCHAR(20),
    FOREIGN KEY (order_id)
    REFERENCES orders (o_id)
);
```

To specify a foreign key consisting of multiple columns, the primary key must consist of multiple columns as well:

```
CREATE TABLE orders (
    o_id INTEGER,
    o_location VARCHAR(20),
    o_price DECIMAL(6,2),
    PRIMARY KEY (o_id, o_location)
);

CREATE TABLE customers (
    id INTEGER PRIMARY KEY,
    order_id INTEGER,
    name VARCHAR(15),
    location VARCHAR(20),
    CONSTRAINT fk_id_name
    FOREIGN KEY (order_id, location)
    REFERENCES orders (o_id, o_location)
);
```

NOTE

The foreign key (`order_id`) and primary key it references (`o_id`) must both be the same data type.

Create a Table with an Automatically Generated Field

If you plan to load a dataset without a unique ID column, you may want to create a column that automatically generates a unique ID. The code in [Table 5-8](#) automatically generates sequential numbers (1, 2, 3, etc.) in the `u_id` column, in each RDBMS.

Table 5-8. Code to automatically generate a unique ID

RDBMS	Code
MySQL	<pre>CREATE TABLE my_table (u_id INTEGER PRIMARY KEY AUTO_INCREMENT, country VARCHAR(2), name VARCHAR(15));</pre>
Oracle	<pre>CREATE TABLE my_table (u_id INTEGER GENERATED BY DEFAULT ON NULL AS IDENTITY, country VARCHAR2(2), name VARCHAR2(15));</pre>
PostgreSQL	<pre>CREATE TABLE my_table (u_id SERIAL, country VARCHAR(2), name VARCHAR(15));</pre>
SQL Server	<pre>-- u_id to begin at 1 and increment by 1 CREATE TABLE my_table (u_id INTEGER IDENTITY(1,1), country VARCHAR(2), name VARCHAR(15));</pre>
SQLite	<pre>CREATE TABLE my_table (u_id INTEGER PRIMARY KEY AUTOINCREMENT, country VARCHAR(2), name VARCHAR(15));</pre>

NOTE

In *Oracle*, VARCHAR2 is typically used instead of VARCHAR. They are identical in terms of functionality, but VARCHAR may one day be modified, so it's safer to use VARCHAR2.

SQLite recommends against using AUTOINCREMENT unless absolutely necessary because it uses additional computing resources. The code will still run without error.

Insert the Results of a Query into a Table

Instead of manually typing values to insert into a new table, you may want to load a new table with data from existing table(s).

Here is a table:

```
SELECT * FROM my_simple_table;
```

id	country	name
1	US	Sam
2	US	Selena
3	CA	Shawn
4	US	Sutton

Create a new table with two columns:

```
CREATE TABLE new_table_two_columns (
    id INTEGER,
    name VARCHAR(15)
);
```

Insert the results from a query into the new table:

```
INSERT INTO new_table_two_columns
    (id, name)
SELECT id, name
FROM   my_simple_table
WHERE  id < 3;
```

The new table would then look like:

```
SELECT * FROM new_table_two_columns;

id  name
--- -----
1  Sam
2  Selena
```

You can also insert values from an existing table and either add or modify other values along the way.

Create a new table with four columns:

```
CREATE TABLE new_table_four_columns (
    id INTEGER,
    name VARCHAR(15),
    new_num_column INTEGER,
    new_text_column VARCHAR(30)
);
```

Insert the results from a query into the new table and fill in values for the new columns:

```
INSERT INTO new_table_four_columns
    (id, name, new_num_column, new_text_column)
SELECT id, name, 2017, 'stargazing'
FROM my_simple_table
WHERE id = 2;
```

Insert the results from a query into the new table and change a value in the row (id in this case):

```
INSERT INTO new_table_four_columns
    (id, name, new_num_column, new_text_column)
SELECT 3, name, 2017, 'wolves'
FROM my_simple_table
WHERE id = 2;
```

The new table would then look like:

```
SELECT * FROM new_table_four_columns;

id  name      new_num_column  new_text_column
--- -----
2  Selena    2017           stargazing
3  Selena    2017           wolves
```

Insert Data from a Text File into a Table

You may want to load data from a *text file* (data stored in plain text without special formatting) into a table. A common type of text file is a *.csv* file (comma separated values). Text files can be opened up in applications outside of an RDBMS including Excel, Notepad,TextEdit, etc.

The following code shows how to load the *my_data.csv* file into a table.

Contents of the *my_data.csv* file:

```
unique_id,canada_us,artist_name
5,"CA","Celine"
6,"CA","Michael"
7,"US","Stefani"
8,,,"Olivia"
...
...
```

Create a table:

```
CREATE TABLE new_table (
    id INTEGER,
    country VARCHAR(2),
    name VARCHAR(15)
);
```

The code in **Table 5-9** loads the *my_data.csv* file into the *new_table* table for each RDBMS. When loading data, you can specify additional details about the data, such as:

- The data is separated by commas (,)

- Text values are enclosed in double quotes ("")
- Each new row is on a new line (\n)
- The first row of the text file (which contains the header) should be ignored

Table 5-9. Code to insert data from a .csv file

RDBMS	Code
MySQL	<pre>LOAD DATA LOCAL INFILE '<file_path>/my_data.csv' INTO TABLE new_table FIELDS TERMINATED BY ',' ENCLOSED BY """ LINES TERMINATED BY '\n' IGNORE 1 ROWS;</pre>
Oracle	<p>While this can be done at the command line using <code>sqlldr</code>, the better approach is to load data through a graphical user interface like SQL*Loader or SQL Developer instead.</p>
PostgreSQL	<pre>\copy new_table FROM '<file_path>/my_data.csv' DELIMITER ',' CSV HEADER</pre>
SQL Server	<pre>BULK INSERT new_table FROM '<file_path>/my_data.csv' WITH (FORMAT = 'CSV', FIELDTERMINATOR = ',', FIELDQUOTE = """", ROWTERMINATOR = '\n', FIRSTROW = 2, TABLOCK); </pre>
SQLite	<pre>.mode csv .import <file_path>/my_data.csv new_table --skip 1</pre>

After inserting the data, the table would look like this:

```
SELECT * FROM new_table;
```

id	country	name
5	CA	Celine
6	CA	Michael
7	US	Stefani
8	NULL	Olivia
...		

Example Filepath to Desktop

If *my_data.csv* is on your Desktop, this is what the file path would look like for each operating system:

- Linux: /home/*my_username/Desktop/my_data.csv*
- MacOS: /Users/*my_username/Desktop/my_data.csv*
- Windows: C:\Users\my_username\Desktop\my_data.csv

NOTE

If MySQL gives you an error that says that loading local data is disabled, you can enable it by updating the global variable `local_infile`, quitting and restarting MySQL:

```
SET GLOBAL local_infile=1;  
quit
```

Missing Data and NULL Values

Each RDBMS interprets missing data from a *.csv* file in a different way. When the following line in a *.csv* file:

```
8,,,"Olivia"
```

is inserted into a SQL table, the missing value between 8 and Olivia would get replaced with:

- A NULL value in *PostgreSQL* and *SQL Server*
- An empty string (' ') in *MySQL* and *SQLite*

In *MySQL* and *SQLite*, you can use \N in a .csv file to represent a NULL value in a SQL table. When the following line in a .csv file,

```
8,\N,"Olivia"
```

is inserted into a *MySQL* table, the \N would get replaced with a NULL value in the table.

When it is inserted into a *SQLite* table, the \N would be hardcoded into the table. You could then run the code,

```
UPDATE new_table  
SET country = NULL  
WHERE country = '\N';
```

to replace the \N placeholders with NULL values in the table.

Modifying Tables

This section covers how to change the table name, columns, constraints, and data in a table.

NOTE

You need ALTER privileges to modify a table. If you get an error when running the code in this section, you do not have the permission to do so and need to talk to your database administrator.

Rename a Table or Column

After you've created a table, you can still rename the table and the columns of the table.

WARNING

If you modify a table, the table will be permanently changed. *There is no undo*, unless there has been a backup created. Double-check your statements before executing them.

Rename a table

The code in **Table 5-10** shows how to rename a table in each RDBMS.

Table 5-10. Code to rename a table

RDBMS	Code
MySQL, Oracle, PostgreSQL, SQLite	<code>ALTER TABLE old_table_name RENAME TO new_table_name;</code>
SQL Server	<code>EXEC sp_rename 'old_table_name', 'new_table_name';</code>

Rename a column

The code in **Table 5-11** shows how to rename a column in each RDBMS.

Table 5-11. Code to rename a column

RDBMS	Code
MySQL,	<code>ALTER TABLE my_table</code>
Oracle,	<code>RENAME COLUMN old_column_name</code>
PostgreSQL,	<code>TO new_column_name;</code>
SQLite	
SQL Server	<code>EXEC sp_rename 'my_table.old_column_name', 'new_column_name', 'COLUMN';</code>

Display, Add, and Delete Columns

After you've created a table, you can view, add, and delete columns from the table.

Display the columns of a table

The code in [Table 5-12](#) shows how to display the columns of a table in each RDBMS.

Table 5-12. Code to display the columns of a table

RDBMS	Code
MySQL,	DESCRIBE my_table;
Oracle	
PostgreSQL	\d my_table
SQL Server	SELECT column_name FROM information_schema.columns WHERE table_name = 'my_table';
SQLite	pragma table_info(my_table);

Add a column to a table

The code in [Table 5-13](#) shows how to add a column to a table in each RDBMS.

Table 5-13. Code to add a column to a table

RDBMS	Code
MySQL,	ALTER TABLE my_table
PostgreSQL	ADD new_num_column INTEGER, ADD new_text_column VARCHAR(30);
Oracle	ALTER TABLE my_table ADD (new_num_column INTEGER, new_text_column VARCHAR(30));
SQL Server	ALTER TABLE my_table ADD new_num_column INTEGER, new_text_column VARCHAR(30);

RDBMS	Code
SQLite	<pre>ALTER TABLE my_table ADD new_num_column INTEGER; ALTER TABLE my_table ADD new_text_column VARCHAR(30);</pre>

Delete a column from a table

The code in [Table 5-14](#) shows how to delete a column from a table in each RDBMS.

NOTE

If a column has any constraints, you must first [delete the constraints](#) before deleting the column.

Table 5-14. Code to delete a column from a table

RDBMS	Code
MySQL,	<pre>ALTER TABLE my_table</pre>
PostgreSQL	<pre>DROP COLUMN new_num_column, DROP COLUMN new_text_column;</pre>
Oracle	<pre>ALTER TABLE my_table DROP COLUMN new_num_column; ALTER TABLE my_table DROP COLUMN new_text_column;</pre>
SQL Server	<pre>ALTER TABLE my_table DROP COLUMN new_num_column, new_text_column;</pre>
SQLite	Refer to the manual modifications steps for SQLite

Manual Modifications in SQLite

SQLite does not support some table modifications, such as deleting columns or adding/modifying/deleting constraints.

As a workaround, you can either use a [graphical user interface](#) to generate code to modify a table, or you can manually create a new table and copy over data (see following steps).

1. Create a new table with the columns and constraints you want.

```
CREATE TABLE my_table_2 (
    id INTEGER NOT NULL,
    country VARCHAR(2),
    name VARCHAR(30)
);
```

2. Copy data from the old table to the new table.

```
INSERT INTO my_table_2
SELECT id, country, name
FROM my_table;
```

3. Confirm that the data is in the new table.

```
SELECT * FROM my_table_2;
```

4. Delete the old table.

```
DROP TABLE my_table;
```

5. Rename the new table.

```
ALTER TABLE my_table_2 RENAME TO my_table;
```

Display, Add, and Delete Rows

After you've created a table, you can view, add, and delete rows from the table.

Display rows of a table

To display the rows of a table, simply write a `SELECT` statement:

```
SELECT * FROM my_table;
```

Add rows to a table

Use `INSERT INTO` to add rows of data to a table:

```
INSERT INTO my_table  
    (id, country, name)  
VALUES (9, 'US', 'Charlie');
```

Delete rows from a table

Use `DELETE FROM` to delete rows of data from a table:

```
DELETE FROM my_table  
WHERE id = 9;
```

Omit the `WHERE` clause to remove all rows from a table:

```
DELETE FROM my_table;
```

Deleting rows from a table is also known as *truncating*, which removes all of the data in a table without changing the table definition. So while the column names and constraints of the table still exist, it is now empty.

To get rid of a table completely, you can [drop the table](#).

Display, Add, Modify, and Delete Constraints

A *constraint* is a rule that specifies what data can be inserted into a table. More on the various types of constraints can be found earlier in this chapter in the [Create a Table with Constraints](#) section.

Display the constraints of a table

The code in [Table 5-15](#) shows how to display the constraints of a table in each RDBMS.

Table 5-15. Code to display the constraints of a table

RDBMS	Code
MySQL	SHOW CREATE TABLE my_table;
Oracle	SELECT * FROM user_cons_columns WHERE table_name = 'MY_TABLE';
PostgreSQL	\d my_table
SQL Server	-- List constraints (except default ones) SELECT table_name, constraint_name, constraint_type FROM information_schema.table_constraints WHERE table_name = 'my_table'; -- List all default constraints SELECT OBJECT_NAME(parent_object_id), COL_NAME(parent_object_id, parent_column_id), definition FROM sys.default_constraints WHERE OBJECT_NAME(parent_object_id) = 'my_table';
SQLite	.schema my_table

NOTE

Oracle stores table names and column names in all caps, unless you surround the column name with **double quotes**. When referring to a table name or a column name in a SQL statement, you must write the name in all caps (MY_TABLE).

Add a constraint

Let's start with the following CREATE TABLE statement:

```
CREATE TABLE my_table (
    id INTEGER NOT NULL,
    country VARCHAR(2) DEFAULT 'CA',
    name VARCHAR(15),
    lower_name VARCHAR(15)
);
```

The code in **Table 5-16** adds a constraint that makes sure that the `lower_name` column is a lowercase version of the `name` column in each RDBMS.

Table 5-16. Code to add a constraint

RDBMS	Code
MySQL,	ALTER TABLE my_table
PostgreSQL,	ADD CONSTRAINT chk_lower_name
SQL Server	CHECK (lower_name = LOWER(name));
Oracle	ALTER TABLE my_table ADD (CONSTRAINT chk_lower_name CHECK (lower_name = LOWER(name)));
SQLite	Refer to the manual modifications steps for SQLite

Modify a constraint

Let's start with the following CREATE TABLE statement:

```
CREATE TABLE my_table (
    id INTEGER NOT NULL,
    country VARCHAR(2) DEFAULT 'CA',
    name VARCHAR(15),
    lower_name VARCHAR(15)
);
```

The code in **Table 5-17** modifies the following constraints:

- Changes the `country` column from defaulting to CA to defaulting to NULL
- Changes the `name` column from allowing 15 characters to allowing 30 characters

Table 5-17. Code to modify constraints in a table

RDBMS	Code
MySQL	<pre>ALTER TABLE my_table MODIFY country VARCHAR(2) NULL, MODIFY name VARCHAR(30);</pre>
Oracle	<pre>ALTER TABLE my_table MODIFY (country DEFAULT NULL, name VARCHAR2(30));</pre>
PostgreSQL	<pre>ALTER TABLE my_table ALTER country DROP DEFAULT, ALTER name TYPE VARCHAR(30);</pre>
SQL Server	<pre>ALTER TABLE my_table ALTER COLUMN country VARCHAR(2) NULL; ALTER TABLE my_table ALTER COLUMN name VARCHAR(30) NULL;</pre>
SQLite	Refer to the manual modifications steps for SQLite

Delete a constraint

The code in **Table 5-18** shows how to delete a constraint from a table in each RDBMS.

Table 5-18. Code to delete a constraint from a table

RDBMS	Code
MySQL	<code>ALTER TABLE my_table DROP CHECK chk_lower_name;</code>
Oracle, PostgreSQL, SQL Server	<code>ALTER TABLE my_table DROP CONSTRAINT chk_lower_name;</code>
SQLite	Refer to the manual modifications steps for SQLite

NOTE

In MySQL, CHECK can be replaced with DEFAULT, INDEX (for UNIQUE constraints), PRIMARY KEY, and FOREIGN KEY. To delete a NOT NULL constraint, you would **MODIFY** the constraint instead.

Update a Column of Data

Use `UPDATE ... SET ...` to update the values in a column of data.

Here is a table:

```
SELECT *  
FROM my_table;
```

id	country	name	awards
2	CA	Celine	5
3	CA	Michael	4
4	US	Stefani	9

Preview the change you'd like to make:

```
SELECT LOWER(name)  
FROM my_table;
```

LOWER(name)

```
-----  
celine  
michael  
stefani
```

Update the values in a column of data:

```
UPDATE my_table  
SET name = LOWER(name);  
  
SELECT * FROM my_table;
```

id	country	name	awards
2	CA	celine	5
3	CA	michael	4
4	US	stefani	9

Update Rows of Data

Use `UPDATE .. SET .. WHERE ..` to update values in a row or multiple rows of data.

Here is a table:

```
SELECT *  
FROM my_table;  
  
id country name awards  
--- -----  
2 CA Celine 5  
3 CA Michael 4  
4 US Stefani 9
```

Preview the change you'd like to make:

```
SELECT awards + 1  
FROM my_table  
WHERE country = 'CA';  
  
awards + 1
```

6

5

Update the values in multiple rows of data:

```
UPDATE my_table  
SET awards = awards + 1  
WHERE country = 'CA';
```

```
SELECT * FROM my_table;
```

id	country	name	awards
2	CA	Celine	6
3	CA	Michael	5
4	US	Stefani	9

WARNING

It is very important to include a `WHERE` clause along with the `SET` clause when you are updating specific rows of data. Without the `WHERE` clause, the entire table would be updated.

Update Rows of Data with the Results of a Query

Instead of manually typing values to update a table, you can set a new value based on the results of a query.

Here is a table:

```
SELECT * FROM my_table;
```

id	country	name	awards
2	CA	Celine	5
3	CA	Michael	4
4	US	Stefani	9

Preview the change you'd like to make:

```
SELECT MIN(awards) FROM my_table;
```

```
MIN(awards)
```

```
-----
```

```
4
```

Update values based on a query:

```
UPDATE my_table  
SET awards = (SELECT MIN(awards) FROM my_table)  
WHERE country = 'CA';
```

```
SELECT * FROM my_table;
```

id	country	name	awards
---	-----	-----	-----
2	CA	Celine	4
3	CA	Michael	4
4	US	Stefani	9

NOTE

MySQL does not allow you to update a table with a query on the same table. In the preceding example, you cannot have UPDATE my_table and FROM my_table. The statement will run if you query FROM another_table.

The results of the query must always return one column and either zero or one row. If zero rows are returned, then the value is set to NULL.

Delete a Table

When you no longer need a table, you can delete it using a `DROP TABLE` statement:

```
DROP TABLE my_table;
```

In *MySQL*, *PostgreSQL*, *SQL Server*, and *SQLite*, you can also add `IF EXISTS` to avoid an error message if the table doesn't exist:

```
DROP TABLE IF EXISTS my_table;
```

WARNING

If you drop a table, you will lose all of the data in the table. *There is no undo*, unless there has been a backup created. I recommend not running this command unless you are 100% sure you don't need the table.

Delete a table with foreign key references

If other tables have foreign keys that reference the table you are dropping, you will need to delete the foreign key constraints in the other tables along with the table you are dropping.

The code in [Table 5-19](#) shows how to delete a table with foreign key references in each RDBMS.

Table 5-19. Code to delete a table with foreign key references

RDBMS	Code
Oracle	<code>DROP TABLE my_table CASCADE CONSTRAINTS;</code>
PostgreSQL	<code>DROP TABLE my_table CASCADE;</code>
MySQL, SQL Server, SQLite	There is no <code>CASCADE</code> keyword, so you must manually delete any foreign key constraints that reference the table before dropping the table.

WARNING

It is dangerous to use CASCADE without knowing exactly what you are deleting. Please proceed with caution. I recommend not running this command unless you are 100% sure you don't need the constraints.”

Indexes

Imagine you have a table with 10 million rows. You write a query on the table to return values that were logged on 2021-01-01:

```
SELECT *
FROM my_table
WHERE log_date = '2021-01-01';
```

This query would take a long time to run. The reason is because behind the scenes, every single row is checked to see if the `log_date` matches `2021-01-01` or not. That is 10 million checks.

To speed this up, you could create an *index* on the `log_date` column. This is something you would do one time, and all future queries could benefit from it.

Book Index Versus SQL Index Comparison

To better understand how a SQL index works, it's helpful to use an analogy. [Table 5-20](#) compares the index at the end of this book with an index in a SQL table.

Table 5-20. Book index versus SQL index comparison

	Book	SQL Table
Terms	A book has many <i>pages</i> . Each page has <i>attributes</i> including the word count, concepts covered, etc.	A table has many <i>rows</i> . Each row has <i>columns</i> , including <code>customer_id</code> , <code>log_date</code> , etc.
Scenario	You are reading this book and want to find all pages about the concept <i>subqueries</i> .	You are querying a table and want to find all rows where the <code>log_date</code> is <code>2021-01-01</code> .
The slow approach	You could start from page 1 and flip through every page of this book to see if <i>subqueries</i> are mentioned or not. This would take a long time.	You could start from row 1 and scan through every row to see whether the <code>log_date</code> is <code>2021-01-01</code> or not. This would take a long time.
Create an index	An index was created for all concepts in this book. Each concept is listed in the index along with the page numbers that talk about the concept.	An index was created on the <code>log_date</code> column in the table. Every <code>log_date</code> is listed in the index along with the row numbers that contain the <code>log_date</code> .
The fast approach	To find pages about <i>subqueries</i> , you can go to the index to quickly find the page numbers that reference <i>subqueries</i> and go to those pages.	To find rows with a <code>log_date</code> of <code>2021-01-01</code> , your query uses the index to quickly find the row numbers that contain the date and return those rows.

When the same query is run on `my_table` (that now has the `log_date` column indexed):

```
SELECT *
FROM my_table
WHERE log_date = '2021-01-01';
```

the query will run much faster because instead of checking each row in the table, it sees the `log_date` of `2021-01-01`, goes to the index, and quickly pulls all rows that have that `log_date`.

TIP

It's a good idea to create an index on a few columns that you filter on often. For example, the primary key column, the date column, etc.

You wouldn't want to create an index for too many columns, though, because it does take up space. Also, any time rows are added or removed, the index would need to be rebuilt, which is time consuming.

Create an Index to Speed Up Queries

The following code creates a new index called `my_index` on the `log_date` column in the `my_table` table:

```
CREATE INDEX my_index ON my_table (log_date);
```

NOTE

When creating an index in *Oracle*, you must uppercase the column name and surround it in quotes:

```
CREATE INDEX my_index ON my_table  
('LOG_DATE');
```

Oracle automatically creates an index for PRIMARY KEY and UNIQUE columns when a table is created.

Indexes can take a long time to create. However, it's a one-time task that's worth it in the long run for many faster queries in the future.

You can also create a multicolumn index or a *composite index*. The following code creates an index on two columns: `log_date` and `team`:

```
CREATE INDEX my_index ON my_table (log_date, team);
```

The order of the columns matters here. If you write a query that filters on:

- Both columns: the index will make the query fast
- The first column (`log_date`): the index will make the query fast
- The second column (`team`): the index will not help because it first organizes data by the `log_date` and then the `team` column

NOTE

You need CREATE privileges to create an index. If you get an error when running the preceding code, you do not have the permission to do so and need to talk to your database administrator.

Delete an index

The code in [Table 5-21](#) shows how to delete an index in each RDBMS.

Table 5-21. Code to delete an index

RDBMS	Code
MySQL, SQL Server	<code>DROP INDEX my_index ON my_table;</code>
Oracle, PostgreSQL, SQLite	<code>DROP INDEX my_index;</code>

WARNING

Dropping an index cannot be undone. Be 100% sure you want to delete an index before dropping it.

On the bright side, there is no data loss. The data in the table is untouched, and you can always recreate the index.

Views

Imagine you have a long and complex SQL query that includes many joins, filters, aggregations, etc. The results of the query are useful to you and something that you want to reference again at a later point.

This is a great situation to create a *view*, or give a name to the output of a query. Remember that the output of a query is a single table, so a view looks just like a table. The difference is that the view doesn't actually hold any data like a table, but just references the data instead.

NOTE

Sometimes database administrators (DBAs) will create views to restrict access to tables. Imagine there is a `customer` table. Most people should only be able to read the data in the table, and not make changes to it.

The DBA can create a `customer` view that includes data identical to the `customer` table. Now, everyone can query the `customer view`, and only the DBA would be able to edit the data within the `customer table`.

The following code is a complex query that we don't want to write over and over again:

```
-- Number of waterfalls owned by each owner
SELECT o.id, o.name,
       COUNT(w.id) AS num_waterfalls
FROM owner o LEFT JOIN waterfall w
      ON o.id = w.owner_id
GROUP BY o.id, o.name;
```

id	name	num_waterfalls
1	Pictured Rocks	3

2 Michigan Nature	3
3 AF LLC	1
4 MI DNR	1
5 Horseshoe Falls	0

Let's say that we want to find the average number of waterfalls that an owner owns. We could do this using either a subquery or a view:

```
-- Subquery Approach
SELECT AVG(num_waterfalls) FROM
(SELECT o.id, o.name,
    COUNT(w.id) AS num_waterfalls
FROM owner o LEFT JOIN waterfall w
    ON o.id = w.owner_id
GROUP BY o.id, o.name) my_subquery;
```

```
AVG(num_waterfalls)
-----
1.6
```

```
-- View Approach
CREATE VIEW owner_waterfalls_vw AS
SELECT o.id, o.name,
    COUNT(w.id) AS num_waterfalls
FROM owner o LEFT JOIN waterfall w
    ON o.id = w.owner_id
GROUP BY o.id, o.name;
```

```
SELECT AVG(num_waterfalls)
FROM owner_waterfalls_vw;
```

```
AVG(num_waterfalls)
-----
1.6
```

NOTE

You need CREATE privileges to create a view. If you get an error when running the preceding code, you do not have the permission to do so and need to talk to your database administrator.

Subqueries Versus Views

Both subqueries and views represent the results of a query, which can then go on to be queried themselves.

- A *subquery* is temporary. It only exists for the duration of the query and is great for one-time use.
- A *view* is saved. Once a view is created, you can continue to write queries that reference the view.

Create a View to Save the Results of a Query

Use `CREATE VIEW` to save the results of a query as a view. The view can then be queried, just like a table.

Using this query:

```
SELECT *
FROM my_table
WHERE country = 'US';
```

id	country	name
1	US	Anna
2	US	Emily
3	US	Molly

Create a view:

```
CREATE VIEW my_view AS  
SELECT *  
FROM my_table  
WHERE country = 'US';
```

Query the view:

```
SELECT * FROM my_view;
```

id	country	name
---	-----	-----
1	US	Anna
2	US	Emily
3	US	Molly

Display existing views

The code in [Table 5-22](#) shows how to display all existing views in each RDBMS.

Table 5-22. Code to display existing views

RDBMS	Code
MySQL	SHOW FULL TABLES WHERE table_type = 'VIEW';
Oracle	SELECT view_name FROM user_views;
PostgreSQL	SELECT table_name FROM information_schema.views WHERE table_schema NOT IN ('information_schema', 'pg_catalog');
SQL Server	SELECT table_name FROM information_schema.views;
SQLite	SELECT name FROM sqlite_master WHERE type = 'view';

Update a view

To update a view is another way of saying to overwrite a view. The code in [Table 5-23](#) shows how to update a view in each RDBMS.

Table 5-23. Code to update a view

RDBMS	Code
MySQL,	CREATE OR REPLACE VIEW my_view AS
Oracle,	SELECT *
PostgreSQL	FROM my_table WHERE country = 'CA';
SQL Server	CREATE OR ALTER VIEW my_view AS SELECT * FROM my_table WHERE country = 'CA';
SQLite	DROP VIEW IF EXISTS my_view; CREATE VIEW my_view AS SELECT * FROM my_table WHERE country = 'CA';

Delete a view

When you no longer need a view, you can delete it using a `DROP VIEW` statement:

```
DROP VIEW my_view;
```

WARNING

Dropping a view cannot be undone. Be 100% sure you want to delete a view before dropping it.

On the bright side, there is no data loss. The data is still in the original table, and you can always recreate the view.

Transaction Management

A *transaction* allows you to more safely update a database. It consists of a sequence of operations that are executed as a single unit. Either all of the operations are executed or none of them are, which is also known as *atomicity*.

The following code kicks off a transaction before making any changes to the tables. After the statements are run, no updates are permanently made to the database until the changes are committed:

```
START TRANSACTION;

INSERT INTO page_views (user_id, page)
    VALUES (525, 'home');
INSERT INTO page_views (user_id, page)
    VALUES (525, 'contact us');
DELETE FROM new_users WHERE user_id = 525;
UPDATE page_views SET page = 'request info'
    WHERE page = 'contact us';

COMMIT;
```

Why is it safer to use a transaction?

After starting a transaction:

All four statements are treated as one unit.

Imagine you run the first three statements, and while you're doing that, someone else edits the database in a way that your fourth statement doesn't run. This is problematic because for you to update the database properly, all four statements need to run together. The transaction does just that—it requires all four statements to act as one unit, so either all of them run or none of them run.

You can undo your changes if needed.

After starting the transaction, you can run each of the statements and see how they would affect the tables. If everything looks right, you can end the transaction and

lock in your changes with a `COMMIT`. If something looks wrong and you want to return things back to the way they were before the transaction, you can do so with a `ROLL BACK`.

In general, if you are updating a database, it is good practice to use a transaction.

The following sections cover two scenarios in which using a transaction is helpful—one ending in a `COMMIT` to confirm changes and the other ending in a `ROLLBACK` to undo changes.

Double-Check Changes Before a `COMMIT`

Imagine you want to delete some rows of data, but you want to double-check that the correct rows are going to get deleted before you permanently remove them from the table.

The following code shows the step-by-step process for how you would use a transaction in SQL to do so.

1. Start a transaction.

```
-- MySQL and PostgreSQL  
START TRANSACTION;  
or  
BEGIN;  
  
-- SQL Server and SQLite  
BEGIN TRANSACTION;
```

In *Oracle*, you are essentially always in a transaction. A transaction begins when you execute your first SQL statement. After a transaction has ended (with a `COMMIT` or `ROLL BACK`), another one begins when the next SQL statement is executed.

2. View the table you plan to change.

You are in transaction mode at this point, meaning no changes will be made to the database.

```
SELECT * FROM books;
```

id	title
1	Becoming
2	Born a Crime
3	Bossypants

3. Test the change and see how it affects the table.

You want to delete all multiword book titles. The following SELECT statement lets you view all the multiword book titles in the table.

```
SELECT * FROM books WHERE title LIKE '% %';
```

id	title
2	Born a Crime

The following DELETE statement uses the same WHERE clause to now delete the multiword book titles in the table.

```
DELETE FROM books WHERE title LIKE '% %';
```

```
SELECT * FROM books;
```

id	title
1	Becoming
3	Bossypants

You are still in transaction mode at this point, so the change has not been made permanent.

4. Confirm the change with COMMIT.

Use COMMIT to lock in the changes. After this step, you are no longer in transaction mode.

```
COMMIT;
```

WARNING

You cannot undo (aka rollback) a transaction once it has been committed.

Undo Changes with a ROLLBACK

Transactions are especially useful to test out changes and undo them if necessary.

1. Start a transaction.

```
-- MySQL and PostgreSQL  
START TRANSACTION;  
or  
BEGIN;  
  
-- SQL Server and SQLite  
BEGIN TRANSACTION;
```

In *Oracle*, you are essentially always in a transaction. A transaction begins when you execute your first SQL statement. After a transaction has ended (with a COMMIT or ROLL BACK), another one begins when the next SQL statement is executed.

2. View the table you plan to change.

You are in transaction mode at this point, meaning no changes will be made to the database.

```
SELECT * FROM books;
```

id	title
1	Becoming
2	Born a Crime
3	Bossypants

3. Test the change and see how it affects the table.

You want to delete all multiword book titles. The following `DELETE` statement accidentally deletes the entire table (you've forgotten a space in '`%%`'). You didn't want this to happen!

```
DELETE FROM books WHERE title LIKE '%%';
```

```
SELECT * FROM books;
```

id	title

It's a good thing you're still in transaction mode at this point, so the change has not been made permanent.

4. Undo the change with `ROLLBACK`.

Instead of `COMMIT`, `ROLLBACK` the changes. The table will not be deleted. After this step, you are no longer in transaction mode and can continue on with your other statements.

```
ROLLBACK;
```

CHAPTER 6

Data Types

In a SQL table, each column can only include values of a single data type. This chapter covers commonly used data types, as well as how and when to use them.

The following statement specifies three columns along with the data type for each column: `id` holds integer values, `name` holds values containing up to 30 characters, and `dt` holds date values:

```
CREATE TABLE my_table (
    id INT,
    name VARCHAR(30),
    dt DATE
);
```

`INT`, `VARCHAR`, and `DATE` are just three of the many data types in SQL. [Table 6-1](#) lists four categories of data types, along with common subcategories. Data type syntax varies widely by RDBMS, and the differences are detailed out in each section of this chapter.

Table 6-1. Data types in SQL

Numeric	String	Datetime	Other
Integer (123)	Character	Date	Boolean
Decimal (1.23)	('hello')	('2021-12-01')	(TRUE)
Floating Point (1.23e10)	Unicode ('西瓜')	Time ('2:21:00') Datetime ('2021-12-01 2:21:00')	Binary (images, documents, etc.)

Table 6-2 lists example values of each data type to show how they are represented in SQL. These values are often referred to as *literals* or *constants*.

Table 6-2. Literals in SQL

Category	Subcategory	Example Values
Numeric	Integer	123 +123 -123
	Decimal	123.45 +123.45 -123.45
	Floating Point	123.45E+23 123.45e-23
	Character	'Thank you!' 'The combo is 39-6-27.'
	Unicode	N'Amélie' N'♥♥♥'
	Date	'2022-10-15' '15-OCT-2022' (<i>Oracle</i>)
Datetime	Time	'10:30:00' '10:30:00.123456' '10:30:00 -6:00'
	Datetime	'2022-10-15 10:30:00' '15-OCT-2022 10:30:00' (<i>Oracle</i>)

Category	Subcategory	Example Values
Other	Boolean	TRUE FALSE
	Binary (example values are displayed as hexadecimal)	X 'AB12' (<i>MySQL, PostgreSQL</i>) x 'AB12' (<i>MySQL, PostgreSQL</i>) 0xAB12 (<i>MySQL, SQL Server, SQLite</i>)

The NULL Literal

Cells with no value are represented by the `NULL` keyword (aka the `NULL` literal), which is case insensitive (`NULL = Null = null`).

You will often see `null` values in a table, but `null` itself is not a data type. Any numeric, string, datetime, or other column can include `null` values within the column.

How to Choose a Data Type

When deciding on a data type for a column, it is important to balance storage size and flexibility.

Table 6-3 shows a few examples of integer data types. Note that each data type allows for a different range of values and requires a different amount of storage space.

Table 6-3. A sample of integer data types

Data Type	Range of Values Allowed	Storage Size
INT	-2,147,483,648 to 2,147,483,647	4 bytes
SMALLINT	-32,768 to 32,767	2 bytes
TINYINT	0 to 255	1 byte

Imagine you have a column of data that contains the number of students in a classroom:

15
25
50
70
100

This column contains numeric data—more specifically, integers. You could choose any of the three integer data types in [Table 6-3](#) to assign to this column.

The case for INT

If storage space isn't an issue, then INT is a simple and solid choice that works across all RDBMSs.

The case for TINYINT

Since all values are between 0 and 255, choosing TINYINT would save on storage space.

The case for SMALLINT

If there may be higher student counts inserted into the column at a later point, SMALLINT allows for more flexibility while still using less space than INT.

There is no single right answer here. The best data type for a column depends on both the storage space and flexibility required.

TIP

If you've already created a table but want to change the data type for a column, you can do so by modifying the column's constraint with an ALTER TABLE statement. More details can be found under [Modifying a Constraint](#) in Chapter 5.

Numeric Data

This section introduces numeric values to give you an idea of how they are represented in SQL, and then goes into detail on integer, decimal, and floating point data types.

Columns with numeric data can be input into numeric functions such as `SUM()` and `ROUND()`, which are covered in the [Numeric Functions](#) section in [Chapter 7](#).

Numeric Values

Numeric values include integers, decimal numbers, and floating point numbers.

Integer values

Numbers without a decimal are treated as integers. The `+` is optional.

123 +123 -123

Decimal values

Decimals include a decimal point and are stored as exact values. The `+` is optional.

123.45 +123.45 -123.45

Floating point values

Floating point values use scientific notation.

123.45E+23 123.45e-23

These values are interpreted as 123.45×10^{23} and 123.45×10^{-23} , respectively.

NOTE

Oracle allows for a trailing F, f, D, or d to indicate FLOAT or DOUBLE (more precise FLOAT value):

123F	+123f	-123.45D	123.45d
------	-------	----------	---------

Integer Data Types

The following code creates an integer column:

```
CREATE TABLE my_table (
    my_integer_column INT
);

INSERT INTO my_table VALUES
(25),
(-525),
(2500252);

SELECT * FROM my_table;
```

my_integer_column
25
-525
2500252

Table 6-4 lists the integer data type options for each RDBMS.

Table 6-4. Integer data types

RDBMS	Data Type	Range of Values Allowed	Storage Size
MySQL	TINYINT	-128 to 127 0 to 255 (unsigned)	1 byte
	SMALLINT	-32,768 to 32,767 0 to 65,535 (unsigned)	2 bytes
	MEDIUMINT	-8,388,608 to 8,388,607 0 to 16,777,215 (unsigned)	3 bytes
	INT or INTEGER	-2,147,483,648 to 2,147,483,647 0 to 4,294,967,295 (unsigned)	4 bytes
	BIGINT	- 2^{63} to $2^{63} - 1$ 0 to $2^{64} - 1$ (unsigned)	8 bytes
Oracle	NUMBER	- 10^{125} to $10^{125} - 1$	1 to 22 bytes
PostgreSQL	SMALLINT	-32,768 to 32,767	2 bytes
	INT or INTEGER	-2,147,483,648 to 2,147,483,647	4 bytes
	BIGINT	- 2^{63} to $2^{63} - 1$	8 bytes
SQL Server	TINYINT	0 to 255	1 byte
	SMALLINT	-32,768 to 32,767	2 bytes
	INT or INTEGER	-2,147,483,648 to 2,147,483,647	4 bytes
	BIGINT	- 2^{63} to $2^{63} - 1$	8 bytes
SQLite	INTEGER	- 2^{63} to $2^{63} - 1$ (if larger, will switch to a REAL data type)	1, 2, 3, 4, 6, or 8 bytes

NOTE

MySQL allows for both signed ranges (positive and negative integers) and unsigned ranges (positive integers only). The default is the signed range. To specify an unsigned range:

```
CREATE TABLE my_table (
    my_integer_column INT UNSIGNED
);
```

PostgreSQL has a **SERIAL** data type that creates an **auto-incrementing integer** (1, 2, 3, etc.) in a column. **Table 6-5** lists the **SERIAL** options, each with a different range.

Table 6-5. Serial options in PostgreSQL

Data Type	Range of Values Generated	Storage Size
SMALLSERIAL	1 to 32,767	2 bytes
SERIAL	1 to 2,147,483,647	4 bytes
BIGSERIAL	1 to 9,223,372,036,854,775,807	8 bytes

Decimal Data Types

Decimal numbers are also known as *fixed point* numbers. They include a decimal point and are stored as an exact value. Monetary data (like 799.95) is often stored as a decimal number.

The following code creates a decimal column:

```
CREATE TABLE my_table (
    my_decimal_column DECIMAL(5,2)
);

INSERT INTO my_table VALUES
(123.45),
(-123),
(12.3);

SELECT * FROM my_table;
```

```

+-----+
| my_decimal_column |
+-----+
|          123.45 |
|       -123.00 |
|          12.30 |
+-----+

```

When defining the data type DECIMAL(5,2):

- 5 is the maximum number of *total digits* that are stored.
This is called the *precision*.
- 2 is the number of digits to the *right of the decimal point*.
This is called the *scale*.

Table 6-6 lists the decimal data type options for each RDBMS.

Table 6-6. Decimal data types

RDBMS	Data Type	Max Digits Allowed	Default
MySQL	DECIMAL or NUMERIC	Total: 65 After decimal point: 30	DECIMAL(10,0)
Oracle	NUMBER	Total: 38 After decimal point: -84 to 127 (negative means before the decimal point)	0 digits after decimal point
PostgreSQL	DECIMAL or NUMERIC	Before decimal point: 131,072 After decimal point: 16,383	DECIMAL(30,6)
SQL Server	DECIMAL or NUMERIC	Total: 38 After decimal point: 38	DECIMAL(18,0)
SQLite	NUMERIC	No inputs	No default

Floating Point Data Types

Floating point numbers are a computer science concept. When a number has many digits, either before or after a decimal

point, instead of storing all the digits, floating point numbers only store a limited number of them to save on space.

- Number: 1234.56789
- Floating point notation: 1.23×10^3

You'll notice that the decimal point "floated" over a few spaces to the left and that an *approximate* value (1.23) was stored, instead of the full original value (1234.56789).

There are two floating point data types:

- *Single precision*: number is represented by at least 6 digits, with a full range of around 1E-38 to 1E+38
- *Double precision*: number is represented by at least 15 digits, with a full range of around 1E-308 to 1E+308 The following code creates both a single precision (FLOAT) and a double precision (DOUBLE) floating point column:

```
CREATE TABLE my_table (
    my_float_column FLOAT,
    my_double_column DOUBLE
);

INSERT INTO my_table VALUES
(123.45, 123.45),
(-12345.6789, -12345.6789),
(1234567.890123456789, 1234567.890123456789);

SELECT * FROM my_table;
```

my_float_column	my_double_column
123.45	123.45
-12345.7	-12345.6789
1234570	1234567.8901234567

WARNING

Because floating point data stores approximate values, comparisons and calculations may be slightly off from what you would expect.

If your data will always have the same number of decimal digits, it is better to use a fixed point data type like DECIMAL to store exact values instead of a floating point data type.

Table 6-7 lists the floating point data type options for each RDBMS.

Table 6-7. Floating point data types

RDBMS	Data Type	Input Range	Storage Size
MySQL	FLOAT	0 to 23 bits	4 bytes
	FLOAT	24 to 53 bits	8 bytes
	DOUBLE	0 to 53 bits	8 bytes
Oracle	BINARY_FLOAT	No inputs	4 bytes
	BINARY_DOUBLE	No inputs	8 bytes
PostgreSQL	REAL	No inputs	4 bytes
	DOUBLE PRECISION	No inputs	8 bytes
SQL Server	REAL	No inputs	4 bytes
	FLOAT	1 to 24 bits	4 bytes
	FLOAT	25 to 53 bits	8 bytes
SQLite	REAL	No inputs	8 bytes

NOTE

Oracle's FLOAT data type is NOT a floating point number. Instead, FLOAT is equivalent to NUMERIC, which is a decimal number. For a floating point data type, you should use BINARY_FLOAT or BINARY_DOUBLE instead.

Bits versus Bytes versus Digits

1 *bit* is the smallest unit of storage. It can have a value of 0 or 1.

1 *byte* consists of 8 bits. Example byte: 10101010.

Each character is represented by a byte. The *digit* 7 = 00000111 in byte form.

String Data

This section introduces string values to give you an idea of how they are represented in SQL, and then goes into detail on character and unicode data types.

Columns with string data can be input into string functions such as LENGTH() and REGEXP() (regular expression), which are covered in the [String Functions](#) section in [Chapter 7](#).

String Values

String values are sequences of characters including letters, numbers, and special characters.

String basics

The standard is to enclose string values in single quotes:

'This is a string.'

Use two adjacent single quotes when you need to embed a single quote in a string:

'You ''re welcome.'

SQL will treat the two adjacent single quotes as a single quote within the string and return:

```
'You're welcome.'
```

TIP

As a best practice, single quotes ('') should be used to enclose string values, while double quotes ("") should be used for identifiers (names of tables, columns, etc.).

Alternatives to single quotes

If your text contains many single quotes and you want to use a different character to denote a string, Oracle and PostgreSQL allow you to do so.

Oracle allows you to preface a string with a Q or q, followed by any character, then the string and finally the character again:

```
Q'[This is a string.]'  
q'[This is a string.]'  
Q'|This is a string.|'
```

PostgreSQL allows you to surround text with two dollar signs and an optional tag name:

```
$$This is a string.$$  
$mytag$This is a string.$mytag$
```

Escape sequences

MySQL and PostgreSQL support *escape sequences*, or a special sequence of text that has meaning. [Table 6-8](#) lists common escape sequences.

Table 6-8. Common escape sequences

Escape Sequence	Description
\'	Single quote
\t	Tab
\n	New line
\r	Carriage return
\b	Backspace
\\"	Backslash

MySQL allows you to include escape sequences within a string using the \ character:

```
SELECT 'hello', 'he\'llo', '\thello';  
+-----+-----+-----+  
| hello | he'llo |     hello |  
+-----+-----+-----+
```

PostgreSQL allows you to include escape sequences in strings if the overall string is prefaced with an E or e:

```
SELECT 'hello', E'he\'llo', e'\thello';  
-----+-----+-----+  
hello | he'llo |     hello
```

Escape sequences only apply to strings enclosed by single quotes and not strings enclosed by dollar signs.

Character Data Types

The most common way to hold string values is to use character data types. The following code creates a variable character column allowing for up to 50 characters:

```
CREATE TABLE my_table (  
    my_character_column VARCHAR(50)  
)
```

```
INSERT INTO my_table VALUES
  ('Here is some text.'),
  ('And some numbers - 1 2 3 4 5'),
  ('And some punctuation! :)');
```

```
SELECT * FROM my_table;
```

my_character_column
Here is some text.
And some numbers - 1 2 3 4 5
And some punctuation! :)

There are three main character data types:

VARCHAR (*variable character*)

This is the most popular string data type. If the data type is VARCHAR(50), then the column will allow up to 50 characters. In other words, the string length is variable.

CHAR (*character*)

If the data type is CHAR(5), then each value in the column will have exactly 5 characters. In other words, the string length is fixed. Data will be right-padded with spaces to be exactly the length specified. For example, 'hi' would be stored as 'hi '.

TEXT

Unlike VARCHAR and CHAR, TEXT requires no inputs, meaning you do not have to specify a length for the text. It is useful for storing long strings, like a paragraph or more of text.

Table 6-9 lists the character data type options for each RDBMS.

Table 6-9. Character data types

RDBMS	Data Type	Input Range	Default	Storage Size
MySQL	CHAR	0 to 255 characters	CHAR(1)	Varies
	VARCHAR	0 to 65,535 characters	Input required	Varies
	TINY TEXT	No inputs	No inputs	255 bytes
	TEXT	No inputs	No inputs	65,535 bytes
	MEDIUM TEXT	No inputs	No inputs	16,777,215 bytes
	LARGE TEXT	No inputs	No inputs	4,294,967,295 bytes
Oracle	CHAR	1 to 2,000 characters	CHAR(1)	Varies
	VAR CHAR2	1 to 4,000 characters	Input required	Varies
	LONG	No inputs	No inputs	2 GB
PostgreSQL	CHAR	1 to 10,485,760 characters	CHAR(1)	Varies
	VARCHAR	1 to 10,485,760 characters	Input required	Varies
	TEXT	No inputs	No inputs	Varies
SQL Server	CHAR	1 to 8,000 bytes	Input required	Varies
	VARCHAR	1 to 8,000 bytes, or max	Input required	Varies, or up to 2 GB
	TEXT	No inputs	No inputs	2,147,483,647 bytes
SQLite	TEXT	No inputs	No inputs	Varies

NOTE

Oracle's VARCHAR2 is typically used instead of VARCHAR. They are identical in terms of functionality, but VARCHAR may one day be modified, so it's safer to use VARCHAR2.

Unicode Data Types

Character data types are typically stored as *ASCII* data, but can also be stored as *Unicode* data if a larger library of characters is needed.

ASCII Versus Unicode Encoding

There are many ways to *encode* data, or in other words, turn data into 0's and 1's for a computer to understand. The default encoding that SQL uses is called *ASCII (American Standard Code for Information Interchange)*.

With ASCII, there are $2^8 = 128$ characters that are turned into a series of eight 0's and 1's. For example, the ! character maps to 00100001. These eight 0's and 1's are known as a *byte* of data.

There are other encoding types beyond ASCII, such as *UTF (Unicode Transformation Format)*. With Unicode, there are 2^{21} characters:

- The first 2^8 characters are the same as ASCII (! = 1000001).
- Other characters include Asian characters, math symbols, emojis, etc.
- Not all characters have been assigned values yet.

The following code shows the difference between the VARCHAR and NVARCHAR (Unicode) data types:

```
CREATE TABLE my_table (
    ascii_text VARCHAR(10),
    unicode_text NVARCHAR(10)
);

INSERT INTO my_table VALUES
    ('abc', 'abc'),
    (N'赵欣婉', N'赵欣婉');

SELECT * FROM my_table;
```

```
+-----+-----+
| ascii_text | unicode_text |
+-----+-----+
| abc       | abc        |
| ???       | 赵欣婉    |
+-----+-----+
```

NOTE

When inserting Unicode data from a [text file](#) into an NVARCHAR column, the Unicode values in the text file do not need the N prefix.

Table 6-10 lists the Unicode data type options for each RDBMS.

Table 6-10. Unicode data types

RDBMS	Data Type	Description
MySQL	NCHAR	Like CHAR, but for Unicode data
	NVARCHAR	Like VARCHAR, but for Unicode data
Oracle	NCHAR	Like CHAR, but for Unicode data
	NVARCHAR2	Like VARCHAR2, but for Unicode data

RDBMS	Data Type	Description
PostgreSQL	CHAR	CHAR supports Unicode data
	VARCHAR	VARCHAR supports Unicode data
SQL Server	NCHAR	Like CHAR, but for Unicode data
	NVARCHAR	Like VARCHAR, but for Unicode data
SQLite	TEXT	TEXT supports Unicode data

Datetime Data

This section introduces datetime values to give you an idea of how they are represented in SQL, and then goes into detail on the datetime data types in each RDBMS.

Columns with datetime data can be input into datetime functions such as `DATEDIFF()` and `EXTRACT()`, which are covered in the [Datetime Functions](#) section in [Chapter 7](#).

Datetime Values

Datetime values can come in the form of dates, times or datetimes.

Date values

A date column should have date values in the format `YYYY-MM-DD`. In *Oracle*, the default format is `DD-MON-YYYY`.

October 15th, 2022 is written as:

```
'2022-10-15'
```

In *Oracle*, October 15th, 2022 is written as:

```
'15-OCT-2022'
```

When referencing a date value in a query, you must preface the string with either a `DATE` or `CAST` keyword to tell SQL it is a date, as shown in [Table 6-11](#).

Table 6-11. Referencing a date in a query

RDBMS	Code
MySQL	<code>SELECT DATE '2021-02-25';</code> <code>SELECT DATE('2021-02-25');</code> <code>SELECT CAST('2021-02-25' AS DATE);</code>
Oracle	<code>SELECT DATE '2021-02-25' FROM dual;</code> <code>SELECT CAST('25-FEB-2021' AS DATE) FROM dual;</code>
PostgreSQL	<code>SELECT DATE '2021-02-25';</code> <code>SELECT DATE('2021-02-25');</code> <code>SELECT CAST('2021-02-25' AS DATE);</code>
SQL Server	<code>SELECT CAST('2021-02-25' AS DATE);</code>
SQLite	<code>SELECT DATE('2021-02-25');</code>

NOTE

In *Oracle*, the date format after the DATE keyword is different than the date format within the CAST function.

Also, in *Oracle*, when doing a calculation or looking up a system variable that only contains a SELECT clause, you need to add FROM dual to the end of the query. dual is a dummy table that holds a single value.

```
SELECT DATE '2021-02-25' FROM dual;
SELECT CURRENT_DATE FROM dual;
```

If a column contains dates of a different format, such as MM/DD/YY, you can apply a **string to date function** for SQL to recognize it as a date.

Time values

A time column should have time values in the format hh:mm:ss. 10:30 a.m. is written as:

```
'10:30:00'
```

You can also include more granular seconds, up to six decimal places:

```
'10:30:12.345678'
```

You can also add a time zone. Central Standard Time is also known as UTC-06:00:

```
'10:30:12.345678 -06:00'
```

When referencing a time value in a query, you must preface the string with either a TIME or CAST keyword to tell SQL it is a time, as shown in [Table 6-12](#).

Table 6-12. Referencing a time in a query

RDBMS	Code
MySQL	<pre>SELECT TIME '10:30'; SELECT TIME('10:30'); SELECT CAST('10:30' AS TIME);</pre>
Oracle	<pre>SELECT TIME '10:30:00' FROM dual; SELECT CAST('10:30' AS TIME) FROM dual;</pre>
PostgreSQL	<pre>SELECT TIME '10:30'; SELECT CAST('10:30' AS TIME);</pre>
SQL Server	<pre>SELECT CAST('10:30' AS TIME);</pre>
SQLite	<pre>SELECT TIME('10:30');</pre>

NOTE

In *Oracle*, the time format after the TIME keyword must include seconds as well.

If a column contains times of a different format, such as *mmss*, you can apply a [string to time function](#) for SQL to recognize it as a time.

Date and time values

A datetime column should have datetime values in the format *YYYY-MM-DD hh:mm:ss*. In *Oracle*, the default format is *DD-MON-YYYY hh:mm:ss*.

October 15, 2022 at 10:30 a.m. is written as:

```
'2022-10-15 10:30'
```

In *Oracle*, October 15, 2022 at 10:30 a.m. is written as:

```
'15-OCT-2022 10:30'
```

When referencing a datetime value in a query, you must prefix the string with either a DATETIME, TIMESTAMP, or CAST keyword to tell SQL it is a datetime, as shown in [Table 6-13](#).

Table 6-13. Referencing a datetime in a query

RDBMS	Code
MySQL	<pre>SELECT TIMESTAMP '2021-02-25 10:30'; SELECT TIMESTAMP('2021-02-25 10:30'); SELECT CAST('2021-02-25 10:30' AS DATETIME);</pre>
Oracle	<pre>SELECT TIMESTAMP '2021-02-25 10:30:00' FROM dual; SELECT CAST('25-FEB-2021 10:30' AS TIMESTAMP) FROM dual;</pre>
PostgreSQL	<pre>SELECT TIMESTAMP '2021-02-25 10:30'; SELECT CAST('2021-02-25 10:30' AS TIMESTAMP);</pre>
SQL Server	<pre>SELECT CAST('2021-02-25 10:30' AS DATETIME);</pre>
SQLite	<pre>SELECT DATETIME('2021-02-25 10:30');</pre>

NOTE

In *MySQL*, the keyword is **TIMESTAMP**, but the data type is **DATETIME** within the **CAST** function.

In *Oracle*, the date format after the **TIMESTAMP** keyword is different than the date format within the **CAST** function. Also, the time format after the **TIMESTAMP** keyword must include seconds, but it is not required within the **CAST** function.

If a column contains datetimes of a different format, such as *MM/DD/YY mm:ss*, you can apply a **string to date or string to time function** for SQL to recognize it as a datetime.

Datetime Data Types

There are many ways to store datetime values. Because the data types vary so widely, in this section, there is a separate subsection for each RDBMS.

MySQL datetime data types

The following code creates five different datetime columns:

```
CREATE TABLE my_table (
    dt DATE,
    tm TIME,
    dttm DATETIME,
    ts TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
    yr YEAR
);

INSERT INTO my_table (dt, tm, dttm, yr)
VALUES ('21-7-4', '6:30',
        2021, '2021-12-25 7:00:01');
```

dt	tm	dttm	yr

```

| 2021-07-04 | 06:30:00 | 2021-12-25 07:00:01 |
+-----+-----+-----+
+-----+-----+
| ts           | yr   |
+-----+-----+
| 2021-01-29 12:56:20 | 2021 |
+-----+-----+

```

Table 6-14 lists common datetime data type options in MySQL.

Table 6-14. MySQL datetime data types

Data Type	Format	Range
DATE	YYYY-MM-DD	1000-01-01 to 9999-12-31
TIME	hh:mm:ss	-838:59:59 to 838:59:59
DATETIME	YYYY-MM-DD hh:mm:ss	1000-01-01 00:00:00 to 9999-12-31 23:59:59
TIMESTAMP	YYYY-MM-DD hh:mm:ss	1970-01-01 00:00:01 UTC to 2038-01-19 03:14:07 UTC
YEAR	YYYY	0000 to 9999

NOTE

Both DATETIME and TIMESTAMP store dates and times. The difference is that DATETIME doesn't have a time zone attached to it, whereas TIMESTAMP stores Unix values (a specific point in time) and is often used to note when a record is created or updated.

Oracle datetime data types

The following code creates four different datetime columns:

```

CREATE TABLE my_table (
    dt DATE,
    ts TIMESTAMP,

```

```

ts_tz TIMESTAMP WITH TIME ZONE,
ts_lc TIMESTAMP WITH LOCAL TIME ZONE
);

INSERT INTO my_table VALUES (
  '4-Jul-21', '4-Jul-21 6:30',
  '4-Jul-21 6:30:45AM CST', '4-Jul-21 6:30'
);

```

DT TS

04-JUL-21 04-JUL-21 06.30.00.000000 AM

TS_TZ

04-JUL-21 06.30.45.000000 AM CST

TS_LC

04-JUL-21 06.30.00.000000 AM

Table 6-15 lists common datetime data type options in Oracle.

Table 6-15. Oracle datetime data types

Data Type	Description
DATE	Can store either just the date or the date and time if the NLS_DATE_FORMAT is updated
TIMESTAMP	Like DATE, but adds fractional seconds (the default is six digits, but can go up to nine digits after the decimal point)
TIMESTAMP WITH TIME ZONE	Like TIMESTAMP, but adds the time zone
TIMESTAMP WITH LOCAL TIME ZONE	Like TIMESTAMP WITH TIME ZONE, but adjusts based on the user's local time zone

Check the datetime formats in Oracle

The following code checks the current date and timestamp formats:

```
SELECT value
FROM nls_session_parameters
WHERE parameter in ('NLS_DATE_FORMAT',
                     'NLS_TIMESTAMP_FORMAT');
```

```
VALUE
-----
DD-MON-RR
DD-MON-RR HH.MI.SSXFF AM
```

To change the date or timestamp format, you can alter the `NLS_DATE_FORMAT` or `NLS_TIMESTAMP_FORMAT` parameter.

The following code changes the current `NLS_DATE_FORMAT` = `DD-MON-RR` to include time as well:

```
ALTER SESSION
SET NLS_DATE_FORMAT = 'YYYY-MM-DD HH:MI:SS';
```

Other common symbols for date and time, such as `YYYY` for year and `HH` for hour can be found in [Table 7-27: Datetime format specifiers](#).

PostgreSQL datetime data types

The following code creates five different datetime columns:

```
CREATE TABLE my_table (
    dt DATE,
    tm TIME,
    tm_tz TIME WITH TIME ZONE,
    ts TIMESTAMP,
    ts_tz TIMESTAMP WITH TIME ZONE
);

INSERT INTO my_table VALUES (
    '2021-7-4', '6:30', '6:30 CST',
```

```

'2021-12-25 7:00:01', '2021-12-25 7:00:01 CST'
);

      dt      |      tm      |      tm_tz      |
-----+-----+-----+
2021-07-04 | 06:30:00 | 06:30:00-06 |

      ts      |      ts_tz      |
-----+-----+
2021-12-25 07:00:01 | 2021-12-25 07:00:01-06

```

Table 6-16 lists common datetime data type options in PostgreSQL.

Table 6-16. PostgreSQL datetime data types

Data Type	Format	Range
DATE	YYYY-MM-DD	4713 BC to 5874897 AD
TIME	hh:mm:ss	00:00:00 to 24:00:00
TIME WITH TIME ZONE	hh:mm:ss+tz	00:00:00+1459 to 24:00:00-1459
TIMESTAMP	YYYY-MM-DD hh:mm:ss	4713 BC to 294276 AD
TIMESTAMP WITH TIME ZONE	YYYY-MM-DD hh:mm:ss+tz	4713 BC to 294276 AD

SQL Server datetime data types

The following code creates six different datetime columns:

```

CREATE TABLE my_table (
    dt DATE,
    tm TIME,
    dttm_sm SMALLDATETIME,
    dttm DATETIME,
    dttm2 DATETIME2,
    dttm_off DATETIMEOFFSET
);

INSERT INTO my_table VALUES (
    '2021-7-4', '6:30', '2021-12-25 7:00:01',

```

```

'2021-12-25 7:00:01', '2021-12-25 7:00:01',
'2021-12-25 7:00:01-06:00'
);

dt          tm
-----
2021-07-04   06:30:00.000000000

dttm_sm
-----
2021-12-25 07:00:00

dttm
-----
2021-12-25 07:00:01.000

dttm2
-----
2021-12-25 07:00:01.0000000

dttm_off
-----
2021-12-25 07:00:01.0000000 -06:00

```

Table 6-17 lists common datetime data type options in SQL Server.

Table 6-17. SQL Server datetime data types

Data Type	Format	Range
DATE	YYYY-MM-DD	0001-01-01 to 9999-12-31
TIME	hh:mm:ss	00:00:00.0000000 to 23:59:59.9999999
SMALLDATETIME	YYYY-MM-DD hh:mm:ss	<i>Date:</i> 1900-01-01 to 2079-06-06 <i>Time:</i> 0:00:00 through 23:59:59
DATETIME	YYYY-MM-DD hh:mm:ss	<i>Date:</i> 1753-01-01 to 9999-12-31 <i>Time:</i> 00:00:00 to 23:59:59.999
DATETIME2	YYYY-MM-DD hh:mm:ss	<i>Date:</i> 0001-01-01 to 9999-12-31 <i>Time:</i> 00:00:00 to 23:59:59.9999999

Data Type	Format	Range
DATETIMEOFFSET	YYYY-MM-DD hh:mm:ss +hh:mm	Time zone offset ranges from -12:00 to +14:00

SQLite datetime data types

SQLite doesn't have a datetime data type. Instead, TEXT, REAL, or INTEGER can be used to store datetime values.

NOTE

Even though there aren't specific datetime data types in SQLite, datetime functions including DATE(), TIME(), and DATETIME() allow you to work with dates and times in SQLite.

More details can be found in the [Datetime Functions](#) section in [Chapter 7](#).

The following code shows three ways to store datetime values in SQLite:

```
CREATE TABLE my_table (
    dt_text TEXT,
    dt_real REAL,
    dt_integer INTEGER
);

INSERT INTO my_table VALUES (
    '2021-12-25 7:00:01',
    '2021-12-25 7:00:01',
    '2021-12-25 7:00:01'
);

dt_text|dt_real
2021-12-25 7:00:01|2021-12-25 7:00:01
```

```
dt_integer  
2021-12-25 7:00:01
```

Table 6-18 lists the datetime data type options in SQLite.

Table 6-18. Table 6-18. SQLite datetime data types

Data Type	Description
TEXT	Stored as a string in the format <i>YYYY-MM-DD HH:MM:SS.SSS</i>
REAL	Stored as a Julian day number, which is the number of days since noon in Greenwich on November 24, 4714 BC
INTEGER	Stored as Unix time, which is the number of seconds since 1970-01-01 00:00:00 UTC

Other Data

There are many other data types in SQL, including ones that are specific to each RDBMS.

Some of them fall into one of the existing categories of data types, but capture more detailed data, like the numeric type MONEY or the datetime type INTERVAL.

Others capture more complex data, like geospatial data that notes a particular location on earth or web data stored in JSON/XML formats.

This section covers two additional data types: Boolean data and data from external files.

Boolean Data

The two Boolean values are TRUE and FALSE. They are case insensitive and should be written without quotes:

```
SELECT TRUE, True, FALSE, False;
```

```
+-----+-----+-----+-----+
|     1 |     1 |      0 |      0 |
+-----+-----+-----+-----+
```

Boolean data types

MySQL, *PostgreSQL*, and *SQLite* support Boolean data types. The following code creates a Boolean column:

```
CREATE TABLE my_table (
    my_boolean_column BOOLEAN
);

INSERT INTO my_table VALUES
    (TRUE),
    (false),
    (1);

SELECT * FROM my_table;

+-----+
| my_boolean_column |
+-----+
|              1 |
|              0 |
|              1 |
+-----+
```

Oracle and *SQL Server* don't have Boolean data types, but there are workarounds:

- In *Oracle*, use the CHAR(1) data type to hold values 'T' and 'F' or the NUMBER(1) data type to hold values 1 and 0.
- In *SQL Server*, use the BIT data type, which holds 1, 0, and NULL values.

External Files (Images, Documents, etc.)

If you plan to include images (.jpg, .png, etc.) or documents (.doc, .pdf, etc.) in a column of data, there are two approaches to do so: store links to the files (more common) or store the files as binary values.

Approach 1: Store links to the files

This is typically the recommended approach if your files are over 1 MB each. For reference, the average iPhone photo is a few MB.

The files would be stored outside of the database, putting less load on the database, and often resulting in better performance.

Steps to store links to files:

1. Note the path names of the files on the file system (/Users/images/img_001.jpg).
2. Create a column that stores strings, like **VAR CHAR(100)**.
3. Insert the path names into the column.

Approach 2: Store the files as binary values

This is typically the recommended approach if your files are smaller in size.

The files would be stored inside of the database, which makes things like backing up the data more straightforward.

Steps to store binary values:

1. Convert the files to binary (if you open up a binary file, it will look like a random sequence of characters, such as Z"/~jhJcE Ät, ÷mfPfôrà).
2. Create a column that stores binary values, like **BLOB**.
3. Insert the binary values into the column.

Binary and hexadecimal values

Binary data represents the raw values that a computer interprets. It is often displayed in a more compact, human-readable form called *hexadecimal*.

- Character: a
- Equivalent binary value: 01100001
- Equivalent hexadecimal value: 61

Hexadecimals convert 1's and 0's into a number system of 16 symbols (0-9 and A-F). Hexadecimals are proceeded by X, x, or 0x:

```
SELECT X'AF12', x'AF12', 0xAF12;
```

0xAF12	0xAF12	0xAF12
--------	--------	--------

MySQL supports all three formats. PostgreSQL supports the first two formats. SQL Server and SQLite support the third format.

In Oracle, while you can't easily display a hexidecimal value, you can use the TO_NUMBER function to display a hexidecimal value as a number instead: `SELECT TO_NUMBER('AF12', 'XXXX') FROM dual;` with the X standing for hexidecimal notation.

Binary data types

The following code creates a binary data column:

```
CREATE TABLE my_table (
    my_binary_column BLOB
);

INSERT INTO my_table VALUES
    ('a'),
    ('aaa'),
    ('ae$ iou');

SELECT * FROM my_table;
```

my_binary_column
0x61
0x616161
0x61652420696F75

In *MySQL*, *Oracle*, and *SQLite*, the most common binary data type is **BLOB**.

In *PostgreSQL*, use **bytea** instead.

In *SQL Server*, use **VARBINARY** (such as **VARBINARY(100)**) instead.

NOTE

In *Oracle* and *SQL Server*, the string `ae$ iou` isn't automatically recognized as a binary value and needs to first be converted into one before getting inserted into a table.

```
-- Oracle
SELECT RAWTOHEX('ae$ iou') FROM dual;

-- SQL Server
SELECT CONVERT(VARBINARY, 'ae$ iou');
```

Table 6-19 lists the binary data type options for each RDBMS.

Table 6-19. Binary data types

RDBMS	Data Type	Description	Input Range	Storage Size
MySQL	BINARY	Fixed length binary string where values are right-padded with 0's to get to the exact size	0 to 255 bytes	Varies
	VARBINARY	Variable length binary string	0 to 65,535 bytes	Varies
	TINYBLOB	Tiny Binary Large OBject	No inputs	255 bytes
	BLOB	Binary Large OBject	No inputs	65,535 bytes
	MEDIUMBLOB	Medium Binary Large OBject	No inputs	16,777,215 bytes
	LARGEBLOB	Large Binary Large OBject	No inputs	4,294,967,295 bytes
	RAW	Variable length binary string	1 to 32,767 bytes	Varies
Oracle	LONG RAW	Larger RAW	No inputs	2 GB
	BLOB	Larger LONG RAW	No inputs	4 GB
PostgreSQL	BYTEA	Variable length binary string	No inputs	1 or 4 bytes plus the actual binary string

RDBMS	Data Type	Description	Input Range	Storage Size
SQL Server	BINARY	Fixed length binary string where values are left padded with 0's to get to the exact size	1 to 8,000 bytes	Varies
	VARBINARY	Variable length binary string	1 to 8,000 bytes, or max	Varies, or up to 2 GB
SQLite	BLOB	Binary Large OBject	No inputs	Stored exactly as it was input

Operators and Functions

Operators and functions are used to perform calculations, comparisons, and transformations within a SQL statement. This chapter provides code examples for commonly used operators and functions.

The following query highlights five operators (+, =, OR, BETWEEN, AND) and two functions (UPPER, YEAR):

```
-- Pay increases for employees
SELECT name, pay_rate + 5 AS new_pay_rate
FROM employees
WHERE UPPER(title) = 'ANALYST'
      OR YEAR(start_date) BETWEEN 2016 AND 2018;
```

Operators Versus Functions

Operators are symbols or keywords that perform a calculation or comparison. Operators can be found within the SELECT, ON, WHERE, and HAVING clauses of a query.

Functions take in zero or more inputs, apply a calculation or transformation, and output a value. Functions can be found within the SELECT, WHERE and HAVING clauses of a query.

In addition to SELECT statements, operators and functions can also be used in INSERT, UPDATE, and DELETE statements.

This chapter includes one section on Operators and five sections on functions: **Aggregate Functions**, **Numeric Functions**, **String Functions**, **Date Time Functions**, and **Null Functions**.

Table 7-1 lists common operators and **Table 7-2** lists common functions.

Table 7-1. Common operators

Logical Operators	Comparison Operators (Symbols)	Comparison Operators (Keywords)	Math Operators
AND	=	BETWEEN	+
OR	!=, <>	EXISTS	-
NOT	<	IN	*
	<=	IS NULL	/
	>	LIKE	%
	v=		

Table 7-2. Common functions

Aggregate Functions	Numeric Functions	String Functions	Datetime Functions	Null Functions
COUNT()	ABS()	LENGTH()	CURRENT_DATE	COALESCE()
SUM()	SQRT()	TRIM()	CURRENT_DATE	LESC()
AVG()	LOG()	CONCAT()	CURRENT_TIME	
MIN()	ROUND()	SUBSTR()	TIME	
MAX()	CAST()	REGEXP()	DATEDIFF() EXTRACT() CONVERT()	

Operators

Operators can be symbols or keywords. They can perform calculations (+) or comparisons (BETWEEN). This section describes the available operators in SQL.

Logical Operators

Logical operators are used to modify conditions, which result in TRUE, FALSE, or NULL. The logical operators in the code block (NOT, AND, OR) are bolded:

```
SELECT *
FROM employees
WHERE start_date IS NOT NULL
      AND (title = 'analyst' OR pay_rate < 25);
```

TIP

When using AND and OR to combine multiple conditional statements, it's a good idea to clearly state the order of operations with parentheses: () .

Table 7-3 lists the logical operators in SQL.

Table 7-3. Logical operators

Operator	Description
AND	Returns TRUE if both conditions are TRUE. Returns FALSE if either is FALSE. Returns NULL otherwise.
OR	Returns TRUE if either condition is TRUE. Returns FALSE if both are FALSE. Returns NULL otherwise.
NOT	Returns TRUE if the condition is FALSE. Returns FALSE if it is TRUE. Returns NULL otherwise.

Imagine there is a column called name. **Table 7-4** shows how values in the column would be evaluated in a conditional statement without a NOT and with a NOT.

Table 7-4. NOT example

name	name IN ('Henry', 'Harper')	name NOT IN ('Henry', 'Harper')
Henry	TRUE	FALSE
Lily	FALSE	TRUE
NULL	NULL	NULL

Imagine there are two columns called `name` and `age`. Table 7-5 shows how values in the columns would be evaluated in a conditional statement with an AND and with an OR.

Table 7-5. AND and OR example

name	age	name = 'Henry'	age > 3	name = 'Henry' AND age > 3	name = 'Henry' OR age > 3
Henry	5	TRUE	TRUE	TRUE	TRUE
Henry	1	TRUE	FALSE	FALSE	TRUE
Lily	2	FALSE	FALSE	FALSE	FALSE
Henry	NULL	TRUE	NULL	NULL	TRUE
Lily	NULL	FALSE	NULL	FALSE	NULL

Comparison Operators

Comparison operators are used in predicates.

Operators Versus Predicates

Predicates are comparisons that include an *operator*:

- The predicate `age = 35` includes the `=` operator.
- The predicate `COUNT(id) < 20` includes the `<` operator.

Predicates are also known as conditional statements. These comparisons are evaluated on each row in a table, and result in a value of TRUE, FALSE, or NULL.

The comparison operators in the code block (IS NULL, =, BETWEEN) are bolded:

```
SELECT *
  FROM employees
 WHERE start_date IS NOT NULL
       AND (title = 'analyst'
          OR pay_rate BETWEEN 15 AND 25);
```

Table 7-6 lists comparison operators that are symbols and **Table 7-7** lists comparison operators that are keywords.

Table 7-6. Comparison operators (symbols)

Operator	Description
=	Tests for equality
!=, <>	Tests for inequality
<	Tests for less than
<=	Tests for less than or equal to
>	Tests for greater than
>=	Tests for greater than or equal to

NOTE

MySQL also allows for `<=>`, which is a null-safe test for equality.

When using `=`, if two values are compared and one of them is `NULL`, the resulting value is `NULL`.

When using `<=>`, if two values are compared and one of them is `NULL`, the resulting value is `0`. If they are both `NULL`, the resulting value is `1`.

Table 7-7. Comparison operators (keywords)

Operator	Description
BETWEEN	Tests whether a value lies within a given range
EXISTS	Tests whether rows exist in a subquery
IN	Tests whether a value is contained in a list of values
IS NULL	Tests whether a value is null or not
LIKE	Tests whether a value matches a simple pattern

NOTE

The `LIKE` operator is used to match simple patterns, like finding text that starts with the letter A. More details can be found in the [LIKE](#) section.

Regular expressions are used to match more complex patterns, like extracting any text located between two punctuation marks. More details can be found in the [regular expressions](#) section.

Each keyword comparison operator is explained in detail in the following sections.

BETWEEN

Use `BETWEEN` to test if a value falls within a range. `BETWEEN` is a combination of `>=` and `<=`. The smaller of the two values should always be written first, with the `AND` operator separating the two.

To find all rows where the ages are greater than or equal to 35 and less than or equal to 44:

```
SELECT *
FROM my_table
WHERE age BETWEEN 35 AND 44;
```

To find all rows where the ages are less than 35 or greater than 44:

```
SELECT *
FROM my_table
WHERE age NOT BETWEEN 35 AND 44;
```

EXISTS

Use EXISTS to test if a subquery returns results or not. Typically, the subquery references another table.

The following query returns employees who also happen to be customers:

```
SELECT e.id, e.name
FROM employees e
WHERE EXISTS (SELECT *
               FROM customers c
               WHERE c.email = e.email);
```

EXISTS Versus JOIN

The EXISTS query could also be written with a JOIN:

```
SELECT *
FROM employees e INNER JOIN customers c
ON e.email = c.email;
```

A JOIN is preferred when you want values from both tables to be returned (SELECT *).

An EXISTS is preferred when you want values from a single table to be returned (SELECT e.id, e.name). This type of query is sometimes referred to as a *semi-join*. EXISTS is also useful when the second table has duplicate rows and you're only interested in whether a row exists or not.

The following query returns customers who have never made a purchase:

```
SELECT c.id, c.name
FROM customers c
WHERE NOT EXISTS (SELECT *
                   FROM orders o
                   WHERE o.email = c.email);
```

IN

Use IN to test whether a value falls within a list of values.

The following query returns values for a few employees:

```
SELECT *
FROM employees
WHERE e.id IN (10001, 10032, 10057);
```

The following query returns employees who have not taken a vacation day:

```
SELECT e.id
FROM employees e
WHERE e.id NOT IN (SELECT v.emp_id
                    FROM vacations v);
```

WARNING

When using NOT IN, if there is even a single NULL value in the column in the subquery (v.emp_id in this case), the subquery will never be TRUE, meaning no rows will be returned.

If there are potentially NULL values in the column in the subquery, it is better to use NOT EXISTS:

```
SELECT e.id
FROM employees e
WHERE NOT EXISTS (SELECT *
                   FROM vacations v
                   WHERE v.emp_id = e.id);
```

IS NULL

Use IS NULL or IS NOT NULL to test whether a value is null or not.

The following query returns employees who don't have a manager:

```
SELECT *
FROM employees
WHERE manager IS NULL;
```

The following query returns employees who have a manager:

```
SELECT *
FROM employees
WHERE manager IS NOT NULL;
```

LIKE

Use LIKE to match a simple pattern. The percent sign (%) is a wildcard that means one or more characters.

Here is a sample table:

```
SELECT * FROM my_table;
```

id	txt
1	You are great.
2	Thank you!
3	Thinking of you.
4	I'm 100% positive.

Find all rows that *contain* the term you:

```
SELECT *
FROM my_table
WHERE txt LIKE '%you%';
```

```
-- MySQL, SQL Server and SQLite results
+-----+-----+
| id   | txt      |
+-----+-----+
| 1    | You are great. |
| 2    | Thank you! |
| 3    | Thinking of you. |
+-----+-----+
-- Oracle and PostgreSQL results
+-----+-----+
| id   | txt      |
+-----+-----+
| 2    | Thank you! |
| 3    | Thinking of you. |
+-----+-----+
```

In *MySQL*, *SQL Server*, and *SQLite*, the pattern is case insensitive. Both *You* and *you* are captured by '%*you*%'.

In *Oracle* and *PostgreSQL*, the pattern is case sensitive. Only *you* is captured by '%*you*%'.

Find all rows that *start with* the term *You*:

```
SELECT *
FROM my_table
WHERE txt LIKE 'You%';
```

```
+-----+-----+
| id   | txt      |
+-----+-----+
| 1    | You are great. |
+-----+-----+
```

Use NOT LIKE to return rows that don't contain the characters.

Instead of the percent sign (%) to match one or more characters, you can use the underscore (_) to match exactly one character.

WARNING

Because % and _ have special meaning when used with LIKE, if you want to search for those actual characters, you'll need to add the ESCAPE keyword.

The following code finds all rows that contain the % symbol:

```
SELECT *
FROM my_table
WHERE txt LIKE '%!%%' ESCAPE '!';
```

id	txt
4	I'm 100% positive.

After the ESCAPE keyword, we've declared ! as the escape character, so when the ! is put in front of the middle % in %!%%, !% is interpreted as %.

LIKE is useful when searching for a particular string of characters. For more advanced pattern searches, you can use regular expressions, which are covered in the [regular expressions](#) section later in this chapter.

Math Operators

Math operators are math symbols that can be used in SQL. The math operator in the code block (/) is bolded:

```
SELECT salary / 52 AS weekly_pay
FROM my_table;
```

[Table 7-8](#) lists the math operators in SQL.

Table 7-8. Math operators

Operator	Description
+	Addition
-	Subtraction
*	Multiplication
/	Division
%	Modulo (remainder)
(MOD in <i>Oracle</i>)	

NOTE

In *PostgreSQL*, *SQL Server*, and *SQLite*, dividing an integer by an integer results in an integer:

```
SELECT 15/2;  
7
```

If you want the result to include decimals, you can either divide by a decimal or use the **CAST** function:

```
SELECT 15/2.0;  
7.5  
  
-- PostgreSQL and SQL Server  
SELECT CAST(15 AS DECIMAL) /  
       CAST(2 AS DECIMAL);  
7.5  
  
-- SQLite  
SELECT CAST(15 AS REAL) /  
       CAST(2 AS REAL);  
7.5
```

Other math operators include:

- *Bitwise operators* such as & (AND), | (OR), and ^ (XOR) for working with bits (0 and 1 values).
- *Assignment operators* such as += (add equals) and -= (subtract equals) for updating values in a table.

Aggregate Functions

An *aggregate function* performs a calculation on many rows of data and results in a single value. **Table 7-9** lists the five basic aggregate functions in SQL.

Table 7-9. Basic aggregate functions

Function	Description
COUNT()	Counts the number of values
SUM()	Calculates the sum of a column
AVG()	Calculates the average of a column
MIN()	Finds the minimum of a column
MAX()	Finds the maximum of a column

Aggregate functions apply calculations to non-null values in a column. The only exception is COUNT(*), which counts *all* rows, including null values.

You can also aggregate multiple rows into a single list using functions like ARRAY_AGG, GROUP_CONCAT, LISTAGG, and STRING_AGG. More details can be found in the [Aggregate rows into a single value or list](#) section in Chapter 8.

NOTE

Oracle supports additional aggregate functions like median (MEDIAN), mode (STATS_MODE), and standard deviation (STDDEV).

Aggregate functions (bolded in the example) are located in the SELECT and HAVING clauses of a query:

```
SELECT COUNT(*) AS total_rows,  
       AVG(age) AS average_age  
  FROM my_table;  
  
SELECT region, MIN(age), MAX(age)  
  FROM my_table  
 GROUP BY region  
 HAVING MIN(age) < 18;
```

WARNING

If you choose to have both aggregate and nonaggregate columns in the SELECT statement, you *must* include all nonaggregate columns in the **GROUP BY** clause (*region* in the preceding example).

Some RDBMSs will throw an error if you do not do this. Other RDBMSs (such as *SQLite*), will not throw an error and allow the statement to run, even though the results returned will be *inaccurate*. It is good practice to double-check your results to make sure they make sense.

MIN/MAX Versus LEAST/GREATEST

The `MIN` and `MAX` functions find the smallest and largest values within a column.

The `LEAST` and `GREATEST` functions find the smallest and largest values within a row. Inputs can be numeric, string, or datetime values. If one value is `NULL`, the function returns `NULL`.

The following table shows the total miles run each quarter, and the query finds the miles run in the best quarter:

```
SELECT * FROM goat;
```

name	q1	q2	q3	q4
Ali	100	200	150	NULL
Bolt	350	400	380	300
Jordan	200	250	300	320

```
SELECT name, GREATEST(q1, q2, q3, q4)
      AS most_miles
FROM goat;
```

name	most_miles
Ali	NULL
Bolt	400
Jordan	320

Numeric Functions

Numeric functions can be applied to columns with **numeric data types**. This section covers common numeric functions in SQL.

Apply Math Functions

There are multiple types of math calculations in SQL:

Math Operators

Calculations using symbols such as +, -, *, /, and %

Aggregate Functions

Calculations that summarize an entire column of data into a single value such as COUNT, SUM, AVG, MIN, and MAX

Math Functions

Calculations using keywords that apply to each row of data such as SQRT, LOG, and more that are listed in **Table 7-10**

NOTE

SQLite only supports the ABS function. Other math functions need to be manually enabled. More details can be found on the math functions page on the [SQLite website](#).

Table 7-10. Math functions

Category	Function	Description	Code	Result
Positive and Negative Values	ABS	Absolute value	SELECT ABS(-5);	5
	SIGN	Returns -1, 0, or 1 depending on if a number is negative, zero, or positive	SELECT SIGN(-5);	-1

Category	Function	Description	Code	Result
Exponents and Logarithms	POWER	x raised to the power of y	SELECT POWER(5,2);	25
	SQRT	Square root	SELECT SQRT(25);	5
	EXP	$e (=2.71828)$ raised to the power of x	SELECT EXP(2);	7.389
	LOG (LOG(y,x) in SQL Server)	Log of y base x	SELECT LOG(2,10); SELECT LOG(10,2);	3.322
	LN (LOG in SQL Server)	Natural log (base e)	SELECT LN(10); SELECT LOG(10);	2.303
Other	LOG10 (LOG(10,x) in Oracle)	Log base 10	SELECT LOG10(100); SELECT LOG(10,100) FROM dual;	2
	MOD (x%y in SQL Server)	Remainder of x / y	SELECT MOD(12,5); SELECT 12%5;	2
	PI (not available in Oracle)	Value of pi	SELECT PI();	3.14159
	COS, SIN, etc.	Cosine, sine, and other trig functions (input is in radians)	SELECT COS(.78);	0.711

Generate Random Numbers

Table 7-11 shows how to generate a random number in each RDBMS. In some cases, you can input a *seed* so that the random number generated is the same each time.

Table 7-11. Random number generator

RDBMS	Code	Range of Results
MySQL,	SELECT RAND();	0 to 1
SQL Server	-- Optional seed SELECT RAND(22);	
Oracle	SELECT DBMS_RANDOM.VALUE FROM dual; SELECT DBMS_RANDOM.RANDOM FROM dual;	0 to 1 -2E31 to +2E31
PostgreSQL	SELECT RANDOM();	0 to 1
SQLite	SELECT RANDOM();	-9E18 to +9E18

The random number function is sometimes used to return a few random rows of a table. While not the most efficient query (since the table has to be sorted), it is a quick hack:

```
-- Return 5 random rows
SELECT *
FROM my_table
ORDER BY RANDOM()
LIMIT 5;
```

Oracle and *SQL Server* allow you to randomly sample a table:

```
-- Return random 20% of rows in Oracle
SELECT *
FROM my_table
SAMPLE(20);

-- Return random 100 rows in SQL Server
SELECT *
```

```
FROM my_table  
TABLESAMPLE(100 ROWS);
```

Round and Truncate Numbers

Table 7-12 shows the various ways to round numbers in each RDBMS.

Table 7-12. Rounding options

Function	Description	Code	Output
CEIL (CEILING in SQL Server)	Rounds up to the nearest integer	SELECT CEIL(98.7654); SELECT CEILING(98.7654);	99
FLOOR	Rounds down to the nearest integer	SELECT FLOOR(98.7654);	98
ROUND	Rounds to a specific number of decimal places, defaults to 0 decimals	SELECT ROUND(98.7654,2);	98.77
TRUNC (TRUNCATE in MySQL; ROUND(x,y,1) in SQL Server)	Cuts off number at a specific number of decimal places, default to 0 decimals	SELECT TRUNC(98.7654,2); SELECT TRUNCATE(98.7654,2); SELECT ROUND(98.7654,2,1);	98.76

NOTE

SQLite only supports the ROUND function. Other rounding options need to be manually enabled. More details can be found on the math functions page on the [SQLite website](#).

Convert Data to a Numeric Data Type

The `CAST` function is used to convert between various data types, and is often used for numeric data.

In the following example, we want to compare a string column with a numeric column

Here is a table with a string column:

id	str_col
1	1.33
2	5.5
3	7.8

Try to compare the string column with numeric value:

```
SELECT *
FROM my_table
WHERE str_col > 3;

-- MySQL, Oracle, and SQLite results
+-----+
| id    | str_col |
+-----+
|      2 | 5.5     |
|      3 | 7.8     |
+-----+

-- PostgreSQL and SQL Server results
Error
```

NOTE

In *MySQL*, *Oracle*, and *SQLite*, the query returns the correct results because the string column is recognized as a numeric column when the `>` operator is introduced.

In *PostgreSQL* and *SQL Server*, you must explicitly `CAST` the string column into a numeric column.

Cast the string column to a decimal column to compare it with a number:

```
SELECT *
FROM my_table
WHERE CAST(str_col AS DECIMAL) > 3;

id | str_col
---+-----
2  | 5.5
3  | 7.8
```

NOTE

Using `CAST` does not permanently change the data type of the column; it is only for the duration of the query. To permanently change the data type of a column, you can [alter the table](#).

String Functions

String functions can be applied to columns with [string data types](#). This section covers common string operations in SQL.

Find the Length of a String

Use the `LENGTH` function.

In the SELECT clause:

```
SELECT LENGTH(name)  
FROM my_table;
```

In the WHERE clause:

```
SELECT *  
FROM my_table  
WHERE LENGTH(name) < 10;
```

In SQL Server, use LEN instead of LENGTH.

NOTE

Most RDBMSs exclude trailing spaces when calculating the length of a string, while *Oracle* includes them.

Example string: 'Al '

Length: 2

Length in Oracle: 5

To exclude trailing spaces in Oracle, use the **TRIM** function:

```
SELECT LENGTH(TRIM(name))  
FROM my_table;
```

Change the Case of a String

Use the UPPER or LOWER function.

UPPER:

```
SELECT UPPER(type)  
FROM my_table;
```

LOWER:

```
SELECT *  
FROM my_table  
WHERE LOWER(type) = 'public';
```

Oracle and *PostgreSQL* also have `INITCAP(string)` to uppercase the first letter of each word in a string and lowercase the other letters.

Trim Unwanted Characters Around a String

Use the `TRIM` function to remove both leading and trailing characters around a string value. The following table has several characters that we'd like to remove:

```
SELECT * FROM my_table;
```

color
!!red
.orange!
..yellow..

Remove spaces around a string

By default, `TRIM` removes spaces from both the left and right sides of a string:

```
SELECT TRIM(color) AS color_clean  
FROM my_table;
```

color_clean
!!red
.orange!
..yellow..

Remove other characters around a string

You can specify other characters to remove besides a single space. The following code removes exclamation marks around a string:

```
SELECT TRIM('!' FROM color) AS color_clean  
FROM my_table;
```

```
+-----+  
| color_clean |  
+-----+  
| red |  
| .orange |  
| ..yellow.. |  
+-----+
```

In *SQLite*, use `TRIM(color, '!')` instead.

Remove characters from the left or right side of a string

There are two options for removing characters on either side of a string.

Option 1: TRIM(LEADING ...) and TRIM(TRAILING ...)

In *MySQL*, *Oracle* and *PostgreSQL*, you can remove characters from either the left or right side of a string with `TRIM(LEADING ...)` and `TRIM(TRAILING ...)`, respectively. The following code removes exclamation marks from the beginning of a string:

```
SELECT TRIM(LEADING '!' FROM color) AS color_clean  
FROM my_table;
```

```
+-----+  
| color_clean |  
+-----+  
| red |  
| .orange! |  
| ..yellow.. |  
+-----+
```

Option 2: LTRIM and RTRIM

Use the keywords `LTRIM` and `RTRIM` to remove characters from either the left or right side of a string, respectively.

In *Oracle*, *PostgreSQL*, and *SQLite*, all unwanted characters can be listed within a single string. The following code

removes periods, exclamation marks, and spaces from the beginning of a string:

```
SELECT LTRIM(color, '.! ') AS color_clean  
FROM my_table;
```

```
+-----+  
| color_clean |  
+-----+  
| red |  
| orange! |  
| yellow.. |  
+-----+
```

In MySQL and SQL Server, only whitespace characters can be removed using `LTRIM(color)` or `RTRIM(color)`.

Concatenate Strings

Use the `CONCAT` function or the concatenation operator (`||`).

```
-- MySQL, PostgreSQL, and SQL Server  
SELECT CONCAT(id, '_', name) AS id_name  
FROM my_table;
```

```
-- Oracle, PostgreSQL, and SQLite  
SELECT id || '_' || name AS id_name  
FROM my_table;
```

```
+-----+  
| id_name |  
+-----+  
| 1_Boots |  
| 2_Pumpkin |  
| 3_Tiger |  
+-----+
```

Search for Text in a String

There are two approaches to search for text in a string.

Approach 1: Does the text appear in the string or not?

Use the `LIKE` operator to determine whether text appears in a string or not. With the following query, only rows that contain the text `some` will be returned:

```
SELECT *
FROM my_table
WHERE my_text LIKE '%some%';
```

More details can be found in the `LIKE` section earlier in this chapter.

Approach 2: Where does the text appear in the string?

Use the `INSTR/POSITION/CHARINDEX` function to determine the location of text in a string.

Table 7-13 lists the parameters required by the location functions in each RDBMS.

Table 7-13. Functions to find the location of text in a string

RDBMS	Code Format
MySQL	<code>INSTR(string, substring)</code> <code>LOCATE(substring, string, position)</code>
Oracle	<code>INSTR(string, substring, position, occurrence)</code>
PostgreSQL	<code>POSITION(substring IN string)</code> <code>STRPOS(string, substring)</code>
SQL Server	<code>CHARINDEX(substring, string, position)</code>
SQLite	<code>INSTR(string, substring)</code>

The inputs are:

- *string (required)*: the string you are searching in (i.e., the name of a `VARCHAR` column)
- *substring (required)*: the string you are searching for (i.e., a character, a word, etc.)

- *position (optional)*: the starting position for the search. The default is to start at the first character (1). If *position* is negative, the search begins at the end of the string.
- *occurrence (optional)*: the first/second/third, etc. time the substring appears in the string. The default is the first occurrence (1).

Here is a sample table:

my_text
Here is some text.
And some numbers - 1 2 3 4 5
And some punctuation! :)

Find the location of the substring `some` within the string `my_text`:

```
SELECT INSTR(my_text, 'some') AS some_location  
FROM my_table;
```

some_location
9
5
5

Counting in SQL Starts at 1

Unlike other programming languages that are zero-indexed (the count starts from 0), the count starts from 1 in SQL.

The 9 in the preceding output means the ninth character.

NOTE

In *Oracle*, regular expressions can also be used to search for a substring using `REGEXP_INSTR`. More details are in the [regular expressions in Oracle](#) section.

Extract a Portion of a String

Use the `SUBSTR` or `SUBSTRING` function. The function name and inputs differ for each RDBMS:

```
-- MySQL, Oracle, PostgreSQL, and SQLite
SUBSTR(string, start, length)
```

```
-- MySQL
SUBSTR(string FROM start FOR length)
```

```
-- MySQL, PostgreSQL, and SQL Server
SUBSTRING(string, start, length)
```

```
-- MySQL and PostgreSQL
SUBSTRING(string FROM start FOR length)
```

The inputs are:

- *string (required)*: the string you are searching in (i.e., the name of a `VARCHAR` column)
- *start (required)*: the starting location of the search. If *start* is set to 1, the search will start from the first character, 2 is the second character, and so on. If *start* is set to 0, it will be treated like a 1. If *start* is negative, the search will start from the last character.
- *length (optional)*: the length of the string returned. If *length* is omitted, then all characters from the *start* to the end of the string will be returned. In *SQL Server*, *length* is required. Here is a sample table:

```
+-----+  
| my_text |  
+-----+  
| Here is some text. |  
| And some numbers - 1 2 3 4 5 |  
| And some punctuation! :) |  
+-----+
```

Extract a substring:

```
SELECT SUBSTR(my_text, 14, 8) AS sub_str  
FROM my_table;
```

```
+-----+  
| sub_str |  
+-----+  
| text. |  
| ers - 1 |  
| tuation! |  
+-----+
```

NOTE

In *Oracle*, regular expressions can also be used to extract a substring using REGEXP_SUBSTR. More details are in the [regular expressions in Oracle](#) section.

Replace Text in a String

Use the REPLACE function. Note the order of the inputs for the function:

```
REPLACE(string, old_string, new_string)
```

Here is a sample table:

```
+-----+  
| my_text |  
+-----+  
| Here is some text. |
```

```
| And some numbers - 1 2 3 4 5 |
| And some punctuation! :)      |
+-----+
```

Replace the word `some` with the word `the`:

```
SELECT REPLACE(my_text, 'some', 'the')
       AS new_text
  FROM my_table;
```

```
+-----+
| new_text          |
+-----+
| Here is the text. |
| And the numbers - 1 2 3 4 5 |
| And the punctuation! :)        |
+-----+
```

NOTE

In *Oracle* and *PostgreSQL*, regular expressions can also be used to replace a string using `REGEXP_REPLACE`. More details are in the [regular expressions in Oracle](#) and [regular expressions in PostgreSQL](#) sections.

Delete Text from a String

You can use the `REPLACE` function, but specify an empty string as the replace value.

Replace the word `some` with an empty string:

```
SELECT REPLACE(my_text, 'some ', '')
       AS new_text
  FROM my_table;
```

```
+-----+
| new_text          |
+-----+
| Here is text.    |
+-----+
```

```
| And numbers - 1 2 3 4 5 |
| And punctuation! :)      |
+-----+
```

Use Regular Expressions

Regular expressions allow you to match complex patterns. For example, finding all words that have exactly five letters or finding all words that start with a capital letter.

Imagine you have the following recipe for taco seasoning:

- 1 tablespoon chili powder
- .5 tablespoon ground cumin
- .5 teaspoon paprika
- .25 teaspoon garlic powder
- .25 teaspoon onion powder
- .25 teaspoon crushed red pepper flakes
- .25 teaspoon dried oregano

You want to exclude the amounts and just have a list of ingredients. To do so, you can write a regular expression to extract all of the text that follows the term spoon.

The regular expression would look like:

```
(?=<spoon ).*$
```

and the results would look like:

```
chili powder
ground cumin
paprika
garlic powder
onion powder
crushed red pepper flakes
dried oregano
```

The regular expression went through all of the text and extracted any text that was sandwiched between the term spoon and the end of the line.

A couple things to note about regular expressions:

- Regular expression syntax is not intuitive. It is helpful to break down the meaning of each part of a regular expression using an online tool, such as [Regex101](#).
- Regular expressions are not SQL-specific. They can be used within many programming languages and text editors.
- [RegexOne](#) provides a quick introductory tutorial. You can also reference Thomas Nield's O'Reilly post, “[An Introduction to Regular Expressions](#).”

TIP

Instead of memorizing regular expression syntax, I recommend finding existing regular expressions and modifying them to fit your needs.

For the previous regular expression, I searched for “regular expression text after string.”

The second Google search result got me to `(?=<=WORD).*$`. I used [Regex101](#) to understand each part of the regular expression, and finally replaced WORD with spoon.

Regular expression functions vary widely by RDBMS, so there is a separate section for each one. SQLite does not support regular expressions by default, but they can be implemented. More details can be found in the [SQLite documentation](#).

Regular Expressions in MySQL

Use `REGEXP` to look for a regular expression pattern anywhere in a string.

Here is a sample table:

title	city
spoon	spoon

```
+-----+-----+
| 10 Things I Hate About You | Seattle      |
| 22 Jump Street          | New Orleans |
| The Blues Brothers       | Chicago     |
| Ferris Bueller's Day Off | Chi         |
+-----+-----+
```

Find all variant spellings of Chicago:

```
SELECT *
FROM movies
WHERE city REGEXP '(Chicago|CHI|Chitown)';
```

```
+-----+-----+
| title           | city    |
+-----+-----+
| The Blues Brothers | Chicago |
| Ferris Bueller's Day Off | Chi |
+-----+-----+
```

MySQL's regular expressions are case-insensitive for character strings; CHI and Chi are seen as equivalent.

Find all movies with numbers in the title:

```
SELECT *
FROM movies
WHERE title REGEXP '\\d';
```

```
+-----+-----+
| title           | city    |
+-----+-----+
| 10 Things I Hate About You | Seattle      |
| 22 Jump Street          | New Orleans |
+-----+-----+
```

In MySQL, any single backslash in a regular expression (\d = any digit) needs to be changed to a double backslash.

Regular expressions in Oracle

Oracle supports many regular expression functions including:

- REGEXP_LIKE matches a regular expression pattern within the text.
- REGEXP_COUNT counts the number of times a pattern appears in the text.
- REGEXP_INSTR locates the positions that a pattern appears in the text.
- REGEXP_SUBSTR returns the substrings in the text that match a pattern.
- REGEXP_REPLACE replaces substrings that match a pattern with other text.

Here is a sample table:

TITLE	CITY
10 Things I Hate About You	Seattle
22 Jump Street	New Orleans
The Blues Brothers	Chicago
Ferris Bueller's Day Off	Chi

Find all movies with numbers in the title:

```
SELECT *
FROM movies
WHERE REGEXP_LIKE(title, '\d');
```

TITLE	CITY
10 Things I Hate About You	Seattle
22 Jump Street	New Orleans

NOTE

The following expressions are equivalent:

```
REGEXP_LIKE(title, '\d')
REGEXP_LIKE(title, '[0-9]')
REGEXP_LIKE(title, '[:digit:]')
```

The third option uses **POSIX** regular expression syntax, which is supported by Oracle.

Count the number of capital letters in the title:

```
SELECT title, REGEXP_COUNT(title, '[A-Z]')
      AS num_caps
  FROM movies;
```

TITLE	NUM_CAPS
10 Things I Hate About You	5
22 Jump Street	2
The Blues Brothers	3
Ferris Bueller's Day Off	4

Find the location of the first vowel in the title:

```
SELECT title, REGEXP_INSTR(title, '[aeiou]')
      AS first_vowel
  FROM movies;
```

TITLE	FIRST_VOWEL
10 Things I Hate About You	6
22 Jump Street	5
The Blues Brothers	3
Ferris Bueller's Day Off	2

Return all numbers in the title:

```
SELECT title, REGEXP_SUBSTR(title, '[0-9]+')
      AS nums
  FROM movies
```

```
WHERE REGEXP_SUBSTR(title, '[0-9]+') IS NOT NULL;
```

TITLE	NUMS
10 Things I Hate About You	10
22 Jump Street	22

Replace all numbers in the title with the number 100:

```
SELECT REGEXP_REPLACE(title, '[0-9]+', '100')
      AS one_hundred_title
FROM movies;
```

```
ONE_HUNDRED_TITLE
```

100 Things I Hate About You
100 Jump Street

NOTE

More details and examples on regular expressions in Oracle can be found in the *Oracle Regular Expressions Pocket Reference* by Jonathan Gennick and Peter Linsley (O'Reilly).

Regular expressions in PostgreSQL

Use SIMILAR TO or ~ to look for a regular expression pattern anywhere in a string.

Here is a sample table:

title		city
10 Things I Hate About You		Seattle
22 Jump Street		New Orleans
The Blues Brothers		Chicago
Ferris Bueller's Day Off		Chi

Find all variant spellings of Chicago:

```
SELECT *
FROM movies
WHERE city SIMILAR TO '(Chicago|CHI|Chi|Chitown)';

      title      |   city
-----+-----
The Blues Brothers    | Chicago
Ferris Bueller's Day Off | Chi
```

PostgreSQL's regular expressions are case-sensitive for character strings; CHI and Chi are seen as different values.

SIMILAR TO Versus ~

SIMILAR TO offers limited regular expression capabilities, and is most often used to offer multiple alternatives (Chicago|CHI|Chi). Other common regex symbols to use with SIMILAR TO are * (0 or more), + (1 or more), and {} (exact number of times).

The tilde (~) should be used for more advanced regular expressions along with **POSIX** syntax, which is another flavor of regular expression that PostgreSQL supports.

The full list of supported symbols can be found in the [PostgreSQL documentation](#).

The following example uses ~ instead of SIMILAR TO.

Find all movies with numbers in the title:

```
SELECT *
FROM movies
WHERE title ~ '\d';
```

```
+-----+-----+
| title          | city      |
+-----+-----+
| 10 Things I Hate About You | Seattle    |
| 22 Jump Street       | New Orleans|
+-----+-----+
```

PostgreSQL also supports REGEXP_REPLACE, which allows you to replace characters in a string that match a particular pattern.

Replace all numbers in the title with the number 100:

```
SELECT REGEXP_REPLACE(title, '\d+', '100')
FROM movies;
```

```
regexp_replace
```

```
-----  
100 Things I Hate About You  
100 Jump Street  
The Blues Brothers  
Ferris Bueller's Day Off
```

The regular expression \d is equivalent to [0-9] and [[:digit::]].

Regular expressions in SQL Server

SQL Server supports a very limited amount of regular expressions through its LIKE keyword.

Here is a sample table:

title	city
10 Things I Hate About You	Seattle
22 Jump Street	New Orleans
The Blues Brothers	Chicago
Ferris Bueller's Day Off	Chi

SQL Server uses a slightly different flavor of regular expression syntax, which is detailed in the [Microsoft documentation](#).

Find all movies with numbers in the title:

```
SELECT *
FROM movies
WHERE title LIKE '%[0-9]%' ;
```

title	city
10 Things I Hate About You	Seattle
22 Jump Street	New Orleans

Convert Data to a String Data Type

When string functions are applied to nonstring data types, it is typically not an issue even though there is a data type mismatch.

The following table has a numeric column called `numbers`:

numbers
1.33
2.5
3.777

When the string function `LENGTH` (or `LEN` in *SQL Server*) is applied on the numeric column, the statement executes without error in most RDBMSs:

```
SELECT LENGTH(numbers) AS len_num
FROM my_table;

-- MySQL, Oracle, SQL Server, and SQLite results
+-----+
| len_num |
+-----+
|      4 |
|      3 |
|      5 |
+-----+

-- PostgreSQL results
Error
```

In *PostgreSQL*, you must explicitly `CAST` the numeric column into a string column:

```
SELECT LENGTH(CAST(numbers AS CHAR(5))) AS len_num
FROM my_table;

len_num
-----
4
3
5
```

NOTE

Using `CAST` does not permanently change the data type of the column—it is only for the duration of the query. To permanently change the data type of a column, you can [alter the table](#).

Datetime Functions

Datetime functions can be applied to columns with [datetime data types](#). This section covers common datetime functions in SQL.

Return the Current Date or Time

The following statements return the current date, current time, and current date and time:

```
-- MySQL, PostgreSQL, and SQLite
SELECT CURRENT_DATE;
SELECT CURRENT_TIME;
SELECT CURRENT_TIMESTAMP;

-- Oracle
SELECT CURRENT_DATE FROM dual;
SELECT CAST(CURRENT_TIMESTAMP AS TIME) FROM dual;
SELECT CURRENT_TIMESTAMP FROM dual;

-- SQL Server
```

```
SELECT CAST(CURRENT_TIMESTAMP AS DATE);
SELECT CAST(CURRENT_TIMESTAMP AS TIME);
SELECT CURRENT_TIMESTAMP;
```

There are many other functions equivalent to these including CURDATE() in MySQL, GETDATE() in SQL Server, etc.

The following three situations show how these functions are used in practice.

Display the current time:

```
SELECT CURRENT_TIME;
```

-----+ <td>current_time </td>	current_time
+-----+ <td> 20:53:35 </td>	20:53:35
+-----+	

Create a table that marks the date and time of creation:

```
CREATE TABLE my_table
    (id INT,
     creation_datetime TIMESTAMP DEFAULT
          CURRENT_TIMESTAMP);
```

```
INSERT INTO my_table (id)
    VALUES (1), (2), (3);
```

-----+ <td>id creation_datetime </td>	id creation_datetime
+-----+ <td> 1 2021-02-15 20:57:12 </td>	1 2021-02-15 20:57:12
2 2021-02-15 20:57:12	
3 2021-02-15 20:57:12	
-----+	

Find all rows of data before a certain date:

```
SELECT *
FROM   my_table
WHERE  creation_datetime < CURRENT_DATE;
```

id	creation_datetime
1	2021-01-15 10:47:02
2	2021-01-15 10:47:02
3	2021-01-15 10:47:02

Add or Subtract a Date or Time Interval

You can add or subtract various time intervals (years, months, days, hours, minutes, seconds, etc.) from date and time values.

Table 7-14 lists the ways to subtract a day.

Table 7-14. Return yesterday's date

RDBMS	Code
MySQL	<code>SELECT CURRENT_DATE - INTERVAL 1 DAY;</code> <code>SELECT SUBDATE(CURRENT_DATE, 1);</code> <code>SELECT DATE_SUB(CURRENT_DATE, INTERVAL 1 DAY);</code>
Oracle	<code>SELECT CURRENT_DATE - INTERVAL '1' DAY FROM dual;</code>
PostgreSQL	<code>SELECT CAST(CURRENT_DATE - INTERVAL '1 day' AS DATE);</code>
SQL Server	<code>SELECT CAST(CURRENT_TIMESTAMP - 1 AS DATE);</code> <code>SELECT DATEADD(DAY, -1, CAST(CURRENT_TIMESTAMP AS DATE));</code>
SQLite	<code>SELECT DATE(CURRENT_DATE, '-1 day');</code>

Table 7-15 lists the ways to add three hours.

Table 7-15. Return the date and time three hours from now

RDBMS	Code
MySQL	<pre>SELECT CURRENT_TIMESTAMP + INTERVAL 3 HOUR; SELECT ADDDATE(CURRENT_TIMESTAMP, INTERVAL 3 HOUR); SELECT DATE_ADD(CURRENT_TIMESTAMP, INTERVAL 3 HOUR);</pre>
Oracle	<pre>SELECT CURRENT_TIMESTAMP + INTERVAL '3' HOUR FROM dual;</pre>
PostgreSQL	<pre>SELECT CURRENT_TIMESTAMP + INTERVAL '3 hours';</pre>
SQL Server	<pre>SELECT DATEADD(HOUR, 3, CURRENT_TIMESTAMP);</pre>
SQLite	<pre>SELECT DATETIME(CURRENT_TIMESTAMP, '+3 hours');</pre>

Find the Difference Between Two Dates or Times

You can find the difference between two dates, times, or datetimes in terms of various time intervals (years, months, days, hours, minutes, seconds, etc.).

Finding a date difference

Given a start and end date, [Table 7-16](#) lists the ways to find the days between the two dates.

Here is a sample table:

start_date	end_date
2016-10-10	2020-11-11
2019-03-03	2021-04-04

Table 7-16. Days between two dates

RDBMS	Code
MySQL	<pre>SELECT DATEDIFF(end_date, start_date) AS day_diff FROM my_table;</pre>
Oracle	<pre>SELECT (end_date - start_date) AS day_diff FROM my_table;</pre>
PostgreSQL	<pre>SELECT AGE(end_date, start_date) AS day_diff FROM my_table;</pre>
SQL Server	<pre>SELECT DATEDIFF(day, start_date, end_date) AS day_diff FROM my_table;</pre>
SQLite	<pre>SELECT (julianday(end_date) - julianday(start_date)) AS day_diff FROM my_table;</pre>

After running the code in the table, these are the results:

```
-- MySQL, Oracle, SQL Server, and SQLite
+-----+
| day_diff |
+-----+
|     1493 |
|      763 |
+-----+
-- PostgreSQL

      day_diff
-----
4 years 1 mon 1 day
2 years 1 mon 1 day
```

Finding a time difference

Given a start and end time, Table 7-17 lists the ways to find the seconds between the two times.

Here is a sample table:

start_time	end_time
10:30:00	11:30:00
14:50:32	15:22:45

Table 7-17. Seconds between two times

RDBMS	Code
MySQL	<pre>SELECT TIMEDIFF(end_time, start_time) AS time_diff FROM my_table;</pre>
Oracle	No time data type
PostgreSQL	<pre>SELECT EXTRACT(epoch from end_time - start_time) AS time_diff FROM my_table;</pre>
SQL Server	<pre>SELECT DATEDIFF(second, start_time, end_time) AS time_diff FROM my_table;</pre>
SQLite	<pre>SELECT (strftime('%s',end_time) - strftime('%s',start_time)) AS time_diff FROM my_table;</pre>

After running the code in the table, these are the results:

-- MySQL
+-----+
time_diff
+-----+
01:00:00
00:32:13
+-----+

```
-- PostgreSQL, SQL Server, and SQLite  
  
time_diff  
-----  
3600  
1933
```

Finding a datetime difference

Given a start and end datetime, [Table 7-18](#) lists the ways to find the number of hours between the two datetimes.

Here is a sample table:

start_dt	end_dt
2016-10-10 10:30:00	2020-11-11 11:30:00
2019-03-03 14:50:32	2021-04-04 15:22:45

Table 7-18. Hours between two datetimes

RDBMS	Code
MySQL	<pre>SELECT TIMESTAMPDIFF(hour, start_dt, end_dt) AS hour_diff FROM my_table;</pre>
Oracle	<pre>SELECT (end_dt - start_dt) AS hour_diff FROM my_table;</pre>
PostgreSQL	<pre>SELECT AGE(end_dt, start_dt) AS hour_diff FROM my_table;</pre>
SQL Server	<pre>SELECT DATEDIFF(hour, start_dt, end_dt) AS hour_diff FROM my_table;</pre>
SQLite	<pre>SELECT ((julianday(end_dt) - julianday(start_dt))*24) AS hour_diff FROM my_table;</pre>

After running the code in the table, these are the results:

```
-- MySQL, SQL Server, and SQLite
+-----+
| hour_diff |
+-----+
|      35833 |
|     18312 |
+-----+

-- Oracle

HOUR_DIFF
-----
+000001493 01:00:00.000000
+000000763 00:32:13.000000

-- PostgreSQL

hour_diff
-----
4 years 1 mon 1 day 01:00:00
2 years 1 mon 1 day 00:32:13
```

NOTE

The *PostgreSQL* result is lengthy:

```
SELECT AGE(end_dt, start_dt)
FROM my_table;
```

age

```
-----  
4 years 1 mon 1 day 01:00:00  
2 years 1 mon 1 day 00:32:13
```

Use the **EXTRACT** function to pull out only the year field.

```
SELECT EXTRACT(year FROM
                AGE(end_dt, start_dt))
FROM my_table;
```

date_part

```
-----  
4  
2
```

Extract a Part of a Date or Time

There are multiple ways to extract a time unit (month, hour, etc.) from a date or time value. **Table 7-19** shows how to do so, specifically for the month time unit.

Table 7-19. Extract the month from a date

RDBMS	Code
MySQL	<pre>SELECT EXTRACT(month FROM CURRENT_DATE); SELECT MONTH(CURRENT_DATE);</pre>
Oracle	<pre>SELECT EXTRACT(month FROM CURRENT_DATE) FROM dual;</pre>
PostgreSQL	<pre>SELECT EXTRACT(month FROM CURRENT_DATE); SELECT DATE_PART('month', CURRENT_DATE);</pre>

RDBMS	Code
SQL Server	<code>SELECT DATEPART(month, CURRENT_TIMESTAMP); SELECT MONTH(CURRENT_TIMESTAMP);</code>
SQLite	<code>SELECT strftime('%m', CURRENT_DATE);</code>

Both *MySQL* and *SQL Server* support time unit specific functions like `MONTH()`, as seen in [Table 7-19](#).

- *MySQL* supports `YEAR()`, `QUARTER()`, `MONTH()`, `WEEK()`, `DAY()`, `HOUR()`, `MINUTE()`, and `SECOND()`.
- *SQL Server* supports `YEAR()`, `MONTH()`, and `DAY()`.

You can replace the `month` or `%m` values in [Table 7-19](#) with other time units. [Table 7-20](#) lists the time units accepted by each RDBMS.

Table 7-20. Time unit options

MySQL	Oracle	PostgreSQL	SQL Server	SQLite
microsecond	second	microsecond	nanosecond	%f (fractional second)
second	minute	millisecond	microsecond	%S (second)
minute	hour	second	millisecond	%s (seconds since 1970-01-01)
hour	day	minute	second	1970-01-01)
day	month	hour	minute	%M (minute)
week	year	day	hour	%H (hour)
month		dow	week	%J (Julian day number)
quarter		week	weekday	%w (day of week)
year		month	day	%d (day of month)
		quarter	dayofyear	%j (day of year)
		year	month	%W (week of year)
		decade	quarter	%m (month)
		century	year	%Y (year)

NOTE

You can also extract a time unit from a string value. The code can be found in [Table 7-28: Extract year from a string](#).

Determine the Day of the Week of a Date

Given a date, determine the day of the week:

- Date: 2020-03-16
- Numeric day of the week: 2 (Sunday is the first day)
- Day of the week: Monday

[Table 7-21](#) returns the numeric day of the week of a given date. Sunday is the first day, Monday the second day, and so on.

Table 7-21. Return the numeric day of the week

RDBMS	Code	Range of Values
MySQL	<code>SELECT DAYOFWEEK('2020-03-16');</code>	1 to 7
Oracle	<code>SELECT TO_CHAR(date '2020-03-16', 'd') FROM dual;</code>	1 to 7
PostgreSQL	<code>SELECT DATE_PART('dow', date '2020-03-16');</code>	0 to 6
SQL Server	<code>SELECT DATEPART(weekday, '2020-03-16');</code>	1 to 7
SQLite	<code>SELECT strftime('%w', '2020-03-16');</code>	0 to 6

[Table 7-22](#) returns the day of the week of a given date.

Table 7-22. Return the day of the week

RDBMS	Code
MySQL	SELECT DAYNAME('2020-03-16');
Oracle	SELECT TO_CHAR(date '2020-03-16', 'day') FROM dual;
PostgreSQL	SELECT TO_CHAR(date '2020-03-16', 'day');
SQL Server	SELECT DATENAME(weekday, '2020-03-16');
SQLite	Not available

Round a Date to the Nearest Time Unit

Oracle and *PostgreSQL* support rounding and truncating (also known as rounding down).

Rounding in Oracle

Oracle supports rounding and truncating a date to the nearest year, month, or day (first day of the week).

To round down to the first of the month:

```
SELECT TRUNC(date '2020-02-25', 'month')  
FROM dual;
```

01-FEB-20

To round to the nearest month:

```
SELECT ROUND(date '2020-02-25', 'month')  
FROM dual;
```

01-MAR-20

Rounding in PostgreSQL

PostgreSQL supports truncating a date to the nearest year, quarter, month, week (first day of the week), day, hour, minute, or second. Additional time units can be found in the [PostgreSQL documentation](#).

To round down to the first of the month:

```
SELECT DATE_TRUNC('month', DATE '2020-02-25');
```

```
2020-02-01 00:00:00-06
```

To round down to the minute:

```
SELECT DATE_TRUNC('minute', TIME '10:30:59.12345');
```

```
10:30:00
```

Convert a String to a Datetime Data Type

There are two ways to convert a string to a datetime data type:

- Use the `CAST` function for a simple case.
- Use `STR_TO_DATE/TO_DATE/CONVERT` for a custom case.

The `CAST` function

If a string column contains dates in a standard format, you can use the `CAST` function to turn it into a date data type.

Table 7-23 shows the code for converting to a date data type.

Table 7-23. Convert a string to a date

RDBMS	Required Date Format	Code
MySQL, PostgreSQL, SQL Server	YYYY-MM-DD	<pre>SELECT CAST('2020-10-15' AS DATE);</pre>
Oracle	DD-MON-YYYY	<pre>SELECT CAST('15-OCT-2020' AS DATE) FROM dual;</pre>
SQLite	YYYY-MM-DD	<pre>SELECT DATE('2020-10-15');</pre>

Table 7-24 shows the code for converting to a time data type.

Table 7-24. Convert a string to a time

RDBMS	Required Time Format	Code
MySQL, PostgreSQL, SQL Server	hh:mm:ss	SELECT CAST('14:30' AS TIME);
Oracle	hh:mm:ss hh:mm:ss AM/PM	SELECT CAST('02:30:00 PM' AS TIME) FROM dual;
SQLite	hh:mm:ss	SELECT TIME('14:30');

Table 7-25 shows the code for converting to a datetime data type.

Table 7-25. Convert a string to a datetime

RDBMS	Required Datetime Format	Code
MySQL,	YYYY-MM-DD	SELECT CAST('2020-10-15'
SQL Server	hh:mm:ss	14:30' AS DATETIME);
Oracle	DD-MON-YYYY hh:mm:ss DD-MON-YYYY hh:mm:ss AM/PM	SELECT CAST('15-OCT-20 02:30:00 PM' AS TIMESTAMP) FROM dual;
PostgreSQL	YYYY-MM-DD hh:mm:ss	SELECT CAST('2020-10-15 14:30' AS TIMESTAMP);
SQLite	YYYY-MM-DD hh:mm:ss	SELECT DATETIME('2020-10-15 14:30');

The `CAST` function can also be used to convert dates to **numeric** and **string** data types.

The `STR_TO_DATE`, `TO_DATE`, and `CONVERT` functions

For dates and times not in the standard `YYYY-MM-DD/DD-MON-YYYY/hh:mm:ss` formats, use a string to date or a string to time function instead.

Table 7-26 lists the string to date and string to time functions for each RDBMS. The example strings in the code are in non-standard formats MM-DD-YY and hhmm.

Table 7-26. String to date and string to time functions

RDBMS	String to date	String to time
MySQL	<pre>SELECT STR_TO_DATE('10-15-22', '%m-%d-%y');</pre>	<pre>SELECT STR_TO_DATE('1030', '%H%i');</pre>
Oracle	<pre>SELECT TO_DATE('10-15-22', 'MM-DD-YY') FROM dual;</pre>	<pre>SELECT TO_TIMESTAMP('1030', 'HH24MI') FROM dual;</pre>
PostgreSQL	<pre>SELECT TO_DATE('10-15-22', 'MM-DD-YY');</pre>	<pre>SELECT TO_TIMESTAMP('1030', 'HH24MI');</pre>
SQL Server	<pre>SELECT CONVERT(VARCHAR, '10-15-22', 105);</pre>	<pre>SELECT CAST(CONCAT(10,':',30) AS TIME);</pre>
SQLite	No nonstandard date function	No non-standard time function

NOTE

SQL Server uses the `CONVERT` function to change a string to a datetime data type. `VARCHAR` is the original data type, `10-15-22` is the date, and `105` stands for the format `MM-DD-YYYY`.

Other date formats are `MM/DD/YYYY` (101), `YYYY.MM.DD` (102), `DD/MM/YYYY` (103), and `DD.MM.YYYY` (104). More formats are listed in the [Microsoft documentation](#).

The time formats are `hh:mi:ss` (108) and `hh:mi:ss:mmm` (114), neither which match the format in **Table 7-26**, which is why the time can't be read in by SQL Server using `CONVERT`.

You can replace the %H%i or HH24MI values in [Table 7-26](#) with other time units. [Table 7-27](#) lists common format specifiers for *MySQL*, *Oracle*, and *PostgreSQL*.

Table 7-27. Datetime format specifiers

MySQL	Oracle and PostgreSQL	Description
%Y	YYYY	4-digit year
%y	YY	2-digit year
%m	MM	Numeric month (1–12)
%b	MON	Abbreviated month (Jan–Dec)
%M	MONTH	Name of month (January–December)
%d	DD	Day (1–31)
%h	HH or HH12	12 hours (1–12)
%H	HH24	24 hours (0–23)
%i	MI	Minutes (0–59)
%s	SS	Seconds (0–59)

Apply a date function to a string column

Imagine you have the following string column:

```
str_column
10/15/2022
10/16/2023
10/17/2024
```

You want to extract the year from each date:

```
year_column
2022
2023
2024
```

Problem

You cannot use a datetime function (`EXTRACT`) on a string column (`str_column`).

Solution

First convert the string column into a date column. Then apply the datetime function. **Table 7-28** lists how to do so in each RDBMS.

Table 7-28. Extract year from a string

RDBMS	Code
MySQL	<pre>SELECT YEAR(STR_TO_DATE(str_column, '%m/%d/%Y')) FROM my_table;</pre>
Oracle	<pre>SELECT EXTRACT(YEAR FROM TO_DATE(str_column, 'MM/DD/YYYY')) FROM my_table;</pre>
PostgreSQL	<pre>SELECT EXTRACT(YEAR FROM TO_DATE(str_column, 'MM/DD/YYYY')) FROM my_table;</pre>
SQL Server	<pre>SELECT YEAR(CONVERT(CHAR, str_column, 101)) FROM my_table;</pre>
SQLite	<pre>SELECT SUBSTR(str_column, 7) FROM my_table;</pre>

NOTE

SQLite does not have datetime functions, but a work-around is to use the **SUBSTR** (substring) function to extract the last four digits.

Null Functions

Null functions can be applied to any type of column and are triggered when a null value is encountered.

Return an Alternative Value if There Is a Null Value

Use the COALESCE function.

Here is a sample table:

id	greeting
1	hi there
2	hello!
3	NULL

When there is no greeting, return hi:

```
SELECT COALESCE(greeting, 'hi') AS greeting  
FROM my_table;
```

greeting
hi there
hello!
hi

MySQL and *SQLite* also accept IFNULL(greeting, 'hi').

Oracle also accepts NVL(greeting, 'hi').

SQL Server also accepts ISNULL(greeting, 'hi').

Advanced Querying Concepts

This chapter covers a few advanced ways of wrangling data using SQL queries, beyond the six main clauses covered in [Chapter 4, *Querying Basics*](#), and the common keywords covered in [Chapter 7, *Operators and Functions*](#).

Table 8-1 includes descriptions and code examples of the four concepts covered in this chapter.

Table 8-1. Advanced querying concepts

Concept	Description	Code Example
Case Statements	If a condition is met, return a particular value. Otherwise, return another value.	<pre>SELECT house_id, CASE WHEN flg = 1 THEN 'for sale' ELSE 'sold' END FROM houses;</pre>
Grouping and Summarizing	Split data into groups, aggregate the data within each group, and return a value for each group.	<pre>SELECT zip, AVG(ft) FROM houses GROUP BY zip;</pre>

Concept	Description	Code Example
Window Functions	Split data into groups, aggregate or order the data within each group, and return a value for each row.	<pre>SELECT zip, ROW_NUMBER() OVER (PARTITION BY zip ORDER BY price) FROM houses;</pre>
Pivoting and Unpivoting	Turn values in a column into multiple columns or consolidate multiple columns into a single column. Supported by Oracle and SQL Server.	<pre>-- Oracle syntax SELECT * FROM listing_info PIVOT (COUNT(*) FOR room IN ('bd','br'));</pre>

This chapter describes each of the concepts in [Table 8-1](#) in detail, along with common use cases.

Case Statements

A CASE statement is used to apply if-else logic within a query. For example, you could use a CASE statement to spell out values. If a 1 is seen, display `vip`. Otherwise, display `general admission`.

ticket	ticket
1	vip
0	general admission
1	vip

In *Oracle*, you may also see the `DECODE` function, which is an older function that operates similarly to the CASE statement.

NOTE

Using a CASE statement temporarily updates values for the duration of a query. To save the updated values, you can do so with an UPDATE statement.

The following two sections go over two types of CASE statements:

- *Simple* CASE statement for a *single column* of data
- *Searched* CASE statement for *multiple columns* of data

Display Values Based on If-Then Logic for a Single Column

To check for equality within a single column of data, use the *simple* CASE statement syntax.

Our goal:

Instead of displaying the values 1/0/NULL, display the values *vip/reserved seating/general admission*:

- If `flag = 1`, then `ticket = vip`
- If `flag = 0`, then `ticket = reserved seating`
- Else, `ticket = general admission`

Here is a sample table:

```
SELECT * FROM concert;
```

name	flag
anton	1
julia	0

```
+-----+-----+
| maren |    1 |
| sarah | NULL |
+-----+-----+
```

Implement the if-else logic with a simple CASE statement:

```
SELECT name, flag,
       CASE flag WHEN 1 THEN 'vip'
                  WHEN 0 THEN 'reserved seating'
                  ELSE 'general admission' END AS ticket
  FROM concert;
```

```
+-----+-----+-----+
| name  | flag  | ticket          |
+-----+-----+-----+
| anton |    1 | vip            |
| julia |    0 | reserved seating |
| maren |    1 | vip            |
| sarah | NULL | general admission |
+-----+-----+-----+
```

If no WHEN clause is a match and no ELSE value is specified, a NULL will be returned.

Display Values Based on If-Then Logic for Multiple Columns

To check for any **condition** (=, <, IN, IS NULL, etc.) within potentially multiple columns of data, use the *searched* CASE statement syntax.

Our goal:

Instead of displaying the values 1/0/NULL, display the values vip/reserved seating/general admission:

- If name = anton, then ticket = vip
- If flag = 0 or flag = 1, then ticket = reserved seating
- Else, ticket = general admission

Here is a sample table:

```
SELECT * FROM concert;
```

name	flag
anton	1
julia	0
maren	1
sarah	NULL

Implement the if-else logic with a searched CASE statement:

```
SELECT name, flag,  
      CASE WHEN name = 'anton' THEN 'vip'  
            WHEN flag IN (0,1) THEN 'reserved seating'  
            ELSE 'general admission' END AS ticket  
FROM concert;
```

name	flag	ticket
anton	1	vip
julia	0	reserved seating
maren	1	reserved seating
sarah	NULL	general admission

If multiple conditions are met, the first listed condition takes precedence.

NOTE

To replace all NULL values in a column with another value, you could use a CASE statement, but it is more common to use the NULL function **COALESCE** instead.

Grouping and Summarizing

SQL allows you to separate rows into groups and summarize the rows within each group in some way, ultimately returning just one row per group.

Table 8-2 lists the concepts associated with grouping and summarizing data.

Table 8-2. Grouping and summarizing concepts

Category	Keyword	Description
The main concept	GROUP BY	Use the GROUP BY clause to separate rows of data into groups.
Ways to summarize rows within each group	COUNT SUM MIN MAX AVG	These aggregate functions summarize multiple rows of data into a <i>single value</i> .
	ARRAY_AGG GROUP_CONCAT LISTAGG STRING_AGG	These functions combine multiple rows of data into a <i>single list</i> .
Extensions of the GROUP BY clause	ROLLUP CUBE GROUPING SETS	<p>Includes rows for subtotals and the grand total as well.</p> <p>Includes aggregations for all possible combinations of the grouped by columns.</p> <p>Allows you to specify particular groupings to display.</p>

GROUP BY Basics

The following table shows the number of calories burned by two people:

```
SELECT * FROM workouts;
```

name	calories
ally	80
ally	75
ally	90
jess	100
jess	92

To create a summary table, you need to decide how to:

1. Group the data: separate all the `name` values into two groups—ally and jess.
2. Aggregate the data within the groups: find the total `calories` within each group.

Use the `GROUP BY` clause to create a summary table:

```
SELECT name,
       SUM(calories) AS total_calories
  FROM workouts
 GROUP BY name;
```

name	total_calories
ally	245
jess	192

More details on how `GROUP BY` works behind the scenes can be found in [The GROUP BY Clause](#) section in [Chapter 4](#).

Grouping by multiple columns

The following table shows the number of calories burned by two people during their daily workouts:

```
SELECT * FROM daily_workouts;
```

id	name	date	calories
1	ally	2021-03-03	80
1	ally	2021-03-04	75
1	ally	2021-03-05	90
2	jess	2021-03-03	100
2	jess	2021-03-05	92

When writing a query with a `GROUP BY` clause that groups by multiple columns and/or includes multiple aggregations:

- The `SELECT` clause should include all *column names* and *aggregations* that you want to appear in the output.
- The `GROUP BY` clause should include the same *column names* that are in the `SELECT` clause.

Use the `GROUP BY` clause to summarize the stats for each person, returning both the `id` and `name` along with two aggregations:

```
SELECT id, name,  
       COUNT(date) AS workouts,  
       SUM(calories) AS calories  
FROM daily_workouts  
GROUP BY id, name;
```

id	name	workouts	calories
1	ally	3	245
2	jess	2	192

Reduce the GROUP BY List for Efficiency

If you know that each `id` is linked to a single `name`, you can exclude the `name` column from the `GROUP BY` clause and get the same results as the previous query:

```
SELECT id,  
       MAX(name) AS name,  
       COUNT(date) AS workouts,  
       SUM(calories) AS calories  
  FROM daily_workouts  
 GROUP BY id;
```

This runs more efficiently behind the scenes since the GROUP BY only has to occur on one column.

To compensate for dropping the name from the GROUP BY clause, you'll notice that an arbitrary aggregate function (MAX) was applied to the name column within the SELECT clause. Because there is only one name value within each id group, MAX(name) will simply return the name associated with each id.

Aggregate Rows into a Single Value or List

With the GROUP BY clause, you must specify how the rows of data within each group should be summarized using either:

- An *aggregate function* to summarize rows into a single value: COUNT, SUM, MIN, MAX, and AVG
- A *function to summarize rows into a list* (shown in the sample table): GROUP_CONCAT and others listed in [Table 8-3](#)

Here is a sample table:

```
SELECT * FROM workouts;
```

name	calories
ally	80
ally	75
ally	90
jess	100
jess	92

Use GROUP_CONCAT in MySQL to create a list of calories:

```

SELECT name,
       GROUP_CONCAT(calories) AS calories_list
  FROM workouts
 GROUP BY name;

+-----+-----+
| name | calories_list |
+-----+-----+
| ally | 80,75,90      |
| jess | 100,92         |
+-----+-----+

```

The GROUP_CONCAT function differs for each RDBMS. **Table 8-3** shows the syntax supported by each RDBMS:

Table 8-3. Aggregate rows into a list in each RDBMS

RDBMS	Code	Default Separator
MySQL	GROUP_CONCAT(calories) GROUP_CONCAT(calories SEPARATOR ',')	Comma
Oracle	LISTAGG(calories) LISTAGG(calories, ',')	No value
PostgreSQL	ARRAY_AGG(calories)	Comma
SQL Server	STRING_AGG(calories, ',')	Separator required
SQLite	GROUP_CONCAT(calories) GROUP_CONCAT(calories, ',')	Comma

In *MySQL*, *Oracle*, and *SQLite*, the separator portion (' ,') is optional. *PostgreSQL* doesn't accept a separator, and *SQL Server* requires one.

You can also return a sorted list or a unique list of values. **Table 8-4** shows the syntax supported by each RDBMS.

Table 8-4. Return a sorted or unique list of values in each RDBMS

RDBMS	Sorted List	Unique List
MySQL	GROUP_CONCAT(calories ORDER BY calories)	GROUP_CONCAT(DISTINCT calories)
Oracle	LISTAGG(calories) WITHIN GROUP(ORDER BY calories)	LISTAGG(DISTINCT calories)
PostgreSQL	ARRAY_AGG(calories ORDER BY calories)	ARRAY_AGG(DISTINCT calories)
SQL Server	STRING_AGG(calories, ',') WITHIN GROUP(ORDER BY calo ries)	Not supported
SQLite	Not supported	GROUP_CONCAT(DISTINCT calories)

ROLLUP, CUBE, and GROUPING SETS

In addition to GROUP BY, you can also add on the ROLLUP, CUBE, or GROUPING SETS keywords to include additional summary information.

The following table lists five purchases over the course of three months:

```
SELECT * FROM spendings;
```

YEAR	MONTH	AMOUNT
2019	1	20
2019	1	30
2020	1	42
2020	2	37
2020	2	100

The examples in this section build on the following GROUP BY example, which returns the total monthly spendings:

```
SELECT year, month,  
SUM(amount) AS total
```

```
FROM spendings  
GROUP BY year, month  
ORDER BY year, month;
```

YEAR	MONTH	TOTAL
2019	1	50
2020	1	42
2020	2	137

ROLLUP

MySQL, Oracle, PostgreSQL, and SQL Server support ROLLUP, which extends the GROUP BY by including additional rows for subtotals and the grand total.

Use ROLLUP to display the yearly and total spendings as well. The 2019, 2020, and total spending rows are added with the addition of ROLLUP:

```
SELECT year, month,  
       SUM(amount) AS total  
  FROM spendings  
 GROUP BY ROLLUP(year, month)  
 ORDER BY year, month;
```

YEAR	MONTH	TOTAL
2019	1	50
2019		50 -- 2019 spendings
2020	1	42
2020	2	137
2020		179 -- 2020 spendings
		229 -- Total spendings

The preceding syntax works in *Oracle, PostgreSQL, and SQL Server*. The *MySQL* syntax is GROUP BY year, month WITH ROLLUP, which also works in *SQL Server*.

CUBE

Oracle, *PostgreSQL*, and *SQL Server* support CUBE, which extends the ROLLUP by including additional rows for all possible combinations of the columns that you are grouping by, as well as the grand total.

Use CUBE to display monthly spendings (single month across multiple years) as well. The January and Feburary spending rows are added with the addition of CUBE:

```
SELECT year, month,
       SUM(amount) AS total
  FROM spendings
 GROUP BY CUBE(year, month)
 ORDER BY year, month;
```

YEAR	MONTH	TOTAL
2019	1	50
2019		50
2020	1	42
2020	2	137
2020		179
	1	92 -- January spendings
	2	137 -- February spendings
		229

The preceding syntax works in *Oracle*, *PostgreSQL*, and *SQL Server*. *SQL Server* also supports the syntax GROUP BY year, month WITH CUBE.

GROUPING SETS

Oracle, *PostgreSQL*, and *SQL Server* support GROUPING SETS, which lets you specify particular groupings that you want to display.

This data is a subset of the results generated by CUBE, only including groupings of one column at a time. In this case, only the total yearly and total monthly spendings are returned:

```
SELECT year, month,
       SUM(amount) AS total
  FROM spendings
 GROUP BY GROUPING SETS(year, month)
 ORDER BY year, month;
```

YEAR	MONTH	TOTAL
2019		50
2020		179
	1	92
	2	137

Window Functions

A *window function* (or *analytic function* in *Oracle*) is similar to an **aggregate function** in that they both perform a calculation on rows of data. The difference is that an aggregate function returns a single value while a window function returns a value for each row of data.

The following table lists employees along with their monthly sales. The following queries use this table to show the difference between an aggregate function and a window function.

```
SELECT * FROM sales;
```

name	month	sales
David	3	2
David	4	11
Laura	3	3
Laura	4	14
Laura	5	7
Laura	6	1

Aggregate Function

SUM() is an aggregate function. The following query sums up the sales for each person and returns each name along with its total_sales value.

```
SELECT name,  
       SUM(sales) AS total_sales  
FROM sales  
GROUP BY name;
```

name	total_sales
David	13
Laura	25

Window Function

ROW_NUMBER() OVER (PARTITION BY name ORDER BY month) is a window function. In the bolded portion of the following query, for each person, a row number is generated that represents the first month, second month, etc., that they sold something. The query returns each row along with its sale_month value.

```
SELECT name,  
       ROW_NUMBER() OVER (PARTITION BY name  
                           ORDER BY month) AS sale_month  
FROM sales;
```

name	sale_month	name	month	sales
David	1	David	3	2
David	2	David	4	11
Laura	1	Laura	3	3
Laura	2	Laura	4	14
Laura	3	Laura	5	7
Laura	4	Laura	6	1

Breaking Down the Window Function

`ROW_NUMBER() OVER (PARTITION BY name ORDER BY month)`

A *window* is a group of rows. In the preceding example, there were two windows. The name `David` had a window of two rows and the name `Laura` had a window of four rows:

`ROW_NUMBER()`

The function you want to apply to each window. Other common functions include `RANK()`, `FIRST_VALUE()`, `LAG()`, etc. This is required.

`OVER`

This states that you are specifying a window function. This is required.

`PARTITION BY name`

This states how you want to split your data into windows. It can be split according to one or more columns. This is optional. If excluded, the window is the entire table.

`ORDER BY month`

This states how each window should be sorted before the function is applied. This is optional in *MySQL*, *PostgreSQL*, and *SQLite*. It is required in *Oracle* and *SQL Server*.

The following sections include examples of how window functions are used in practice.

Rank the Rows in a Table

Use the `ROW_NUMBER()`, `RANK()`, or `DENSE_RANK()` function to add a row number to each row of a table.

The following table shows the number of babies given popular names:

```
SELECT * FROM baby_names;
```

gender	name	babies
F	Emma	92
F	Mia	88
F	Olivia	100
M	Liam	105
M	Mateo	95
M	Noah	110

The two following queries:

- Rank the names by popularity
- Rank the names by popularity for each gender

Rank the names by popularity:

```
SELECT gender, name,
       ROW_NUMBER() OVER (
           ORDER BY babies DESC) AS popularity
FROM baby_names;
```

gender	name	popularity
M	Noah	1
M	Liam	2
F	Olivia	3
M	Mateo	4
F	Emma	5
F	Mia	6

Rank the names by popularity for each gender:

```
SELECT gender, name,
       ROW_NUMBER() OVER (PARTITION BY gender
           ORDER BY babies DESC) AS popularity
FROM baby_names;
```

gender	name	popularity
F	Olivia	1
F	Emma	2
F	Mia	3
M	Noah	1
M	Liam	2
M	Mateo	3

ROW_NUMBER Versus RANK Versus DENSE_RANK

There are three approaches to adding row numbers. Each one has a different way of handling ties.

ROW_NUMBER breaks the tie:

NAME	BABIES	POPULARITY
Olivia	99	1
Emma	80	2
Sophia	80	3
Mia	75	4

RANK keeps the tie:

NAME	BABIES	POPULARITY
Olivia	99	1
Emma	80	2
Sophia	80	2
Mia	75	4

DENSE_RANK keeps the tie and doesn't skip numbers:

NAME	BABIES	POPULARITY
Olivia	99	1
Emma	80	2
Sophia	80	2
Mia	75	3

Return the First Value in Each Group

Use FIRST_VALUE and LAST_VALUE to return the first and last rows of a window, respectively.

The following queries break down the two-step process to return the most popular name for each gender.

Step 1: Display the most popular name for each gender.

```
SELECT gender, name, babies,
       FIRST_VALUE(name) OVER (PARTITION BY gender
                                ORDER BY babies DESC) AS top_name
  FROM baby_names;
```

gender	name	babies	top_name
F	Olivia	100	Olivia
F	Emma	92	Olivia
F	Mia	88	Olivia
M	Noah	110	Noah
M	Liam	105	Noah
M	Mateo	95	Noah

Use the output as a subquery for the next step, which filters on the subquery.

Step 2: Return only the two rows containing the most popular names.

```
SELECT * FROM

(SELECT gender, name, babies,
       FIRST_VALUE(name) OVER (PARTITION BY gender
                                ORDER BY babies DESC) AS top_name
  FROM baby_names) AS top_name_table

WHERE name = top_name;
```

gender	name	babies	top_name
F	Olivia	100	Olivia

M	Noah	110	Noah

In Oracle, exclude the AS `top_name_table` portion.

Return the Second Value in Each Group

Use `NTH_VALUE` to return a specific rank number within each window. *SQL Server* does not support `NTH_VALUE`. Instead, refer to the code in the next section, [Return the first two values in each group](#), but only return the second value.

The following queries break down the two-step process to return the second most popular name for each gender.

Step 1: Display the second most popular name for each gender.

```
SELECT gender, name, babies,
       NTH_VALUE(name, 2) OVER (PARTITION BY gender
                                ORDER BY babies DESC) AS second_name
  FROM baby_names;
```

gender	name	babies	second_name
F	Olivia	100	NULL
F	Emma	92	Emma
F	Mia	88	Emma
M	Noah	110	NULL
M	Liam	105	Liam
M	Mateo	95	Liam

The second parameter in `NTH_VALUE(name, 2)` is what specifies the second value in the window. This can be any positive integer.

Use the output as a subquery for the next step, which filters on the subquery.

Step 2: Return only the two rows containing the second most popular names.

```
SELECT * FROM
  (SELECT gender, name, babies,
```

```

        NTH_VALUE(name, 2) OVER (PARTITION BY gender
          ORDER BY babies DESC) AS second_name
FROM baby_names) AS second_name_table

WHERE name = second_name;

```

gender	name	babies	second_name
F	Emma	92	Emma
M	Liam	105	Liam

In *Oracle*, exclude the AS second_name_table portion.

Return the First Two Values in Each Group

Use **ROW_NUMBER** within a subquery to return multiple rank numbers within each group.

The following queries break down the two-step process to return the first and second most popular names for each gender.

Step 1: Display the popularity rank for each gender.

```

SELECT gender, name, babies,
       ROW_NUMBER() OVER (PARTITION BY gender
                           ORDER BY babies DESC) AS popularity
FROM baby_names;

```

gender	name	babies	popularity
F	Olivia	100	1
F	Emma	92	2
F	Mia	88	3
M	Noah	110	1
M	Liam	105	2
M	Mateo	95	3

Use the output as a subquery for the next step, which filters on the subquery.

Step 2: Filter on the rows that contain ranks 1 and 2.

```
SELECT * FROM  
  
(SELECT gender, name, babies,  
       ROW_NUMBER() OVER (PARTITION BY gender  
                           ORDER BY babies DESC) AS popularity  
    FROM baby_names) AS popularity_table  
  
WHERE popularity IN (1,2);
```

gender	name	babies	popularity
F	Olivia	100	1
F	Emma	92	2
M	Noah	110	1
M	Liam	105	2

In Oracle, exclude the AS popularity_table portion.

Return the Prior Row Value

Use LAG and LEAD to look a certain number of rows behind and ahead, respectively.

Use LAG to return the previous row:

```
SELECT gender, name, babies,  
       LAG(name) OVER (PARTITION BY gender  
                           ORDER BY babies DESC) AS prior_name  
    FROM baby_names;
```

gender	name	babies	prior_name
F	Olivia	100	NULL
F	Emma	92	Olivia
F	Mia	88	Emma
M	Noah	110	NULL
M	Liam	105	Noah
M	Mateo	95	Liam

Use `LAG(name, 2, 'No name')` to return the names from two rows prior and replace NULL values with No name:

```
SELECT gender, name, babies,
       LAG(name, 2, 'No name')
    OVER (PARTITION BY gender
          ORDER BY babies DESC) AS prior_name_2
  FROM baby_names;
```

gender	name	babies	prior_name_2
F	Olivia	100	No name
F	Emma	92	No name
F	Mia	88	Olivia
M	Noah	110	No name
M	Liam	105	No name
M	Mateo	95	Noah

The `LAG` and `LEAD` functions each take three arguments:
`LAG(name, 2, 'None')`

- `name` is the value you want to return. It is required.
- `2` is the row offset. It is optional and defaults to `1`.
- `'No name'` is the value that will be returned when there is no value. It is optional and defaults to `NULL`.

Calculate the Moving Average

Use a combination of the `AVG` function and the `ROWS BETWEEN` clause to calculate the moving average.

Here is a sample table:

```
SELECT * FROM sales;
```

name	month	sales
David	1	2
David	2	11
David	3	6
David	4	8
Laura	1	3
Laura	2	14
Laura	3	7
Laura	4	1
Laura	5	20

For each person, find the three-month moving average of sales, from two months prior to the current month:

```
SELECT name, month, sales,
       AVG(sales) OVER (PARTITION BY name
                         ORDER BY month
                         ROWS BETWEEN 2 PRECEDING AND
                         CURRENT ROW) three_month_ma
FROM sales;
```

name	month	sales	three_month_ma
David	1	2	2.0000
David	2	11	6.5000
David	3	6	6.3333
David	4	8	8.3333
Laura	1	3	3.0000
Laura	2	14	8.5000
Laura	3	7	8.0000
Laura	4	1	7.3333
Laura	5	20	9.3333

NOTE

The preceding example looks at the two rows prior through the current row:

ROWS BETWEEN 2 PRECEDING AND CURRENT ROW

You can also look at the next rows using the FOLLOWING keyword:

ROWS BETWEEN 2 PRECEDING AND 3 FOLLOWING

These ranges are sometimes referred to as *sliding windows*.

Calculate the Running Total

Use a combination of the SUM function and the ROWS BETWEEN UNBOUNDED clause to calculate the running total.

For each person, find the running total of sales, up to the current month:

```
SELECT name, month, sales,
       SUM(sales) OVER (PARTITION BY name
                        ORDER BY month
                        ROWS BETWEEN UNBOUNDED PRECEDING AND
                        CURRENT ROW) running_total
  FROM sales;
```

name	month	sales	running_total
David	1	2	2
David	2	11	13
David	3	6	19
David	4	8	27
Laura	1	3	3
Laura	2	14	17
Laura	3	7	24
Laura	4	1	25
Laura	5	20	45

NOTE

Here, we calculated the running total for each person. To calculate the running total for the entire table, you can remove the PARTITION BY name portion of the code.

ROWS Versus RANGE

An alternative to ROWS BETWEEN is RANGE BETWEEN. The following query calculates the running total of sales made by all employees, using both the ROWS and RANGE keywords:

```
SELECT month, name,
    SUM(sales) OVER (ORDER BY month ROWS BETWEEN
        UNBOUNDED PRECEDING AND CURRENT ROW) rt_rows,
    SUM(sales) OVER (ORDER BY month RANGE BETWEEN
        UNBOUNDED PRECEDING AND CURRENT ROW) rt_range
FROM sales;
```

month	name	rt_rows	rt_range
1	David	2	5
1	Laura	5	5
2	David	16	30
2	Laura	30	30
3	David	36	43
3	Laura	43	43
4	David	51	52
4	Laura	52	52
5	Laura	72	72

The difference between the two is that RANGE will return the same running total value for each month (since the data was ordered by month), while ROWS will have a different running total value for each row.

Pivoting and Unpivoting

Oracle and *SQL Server* support the PIVOT and UNPIVOT operations. PIVOT takes a single column and splits it out into multiple columns. UNPIVOT takes multiple columns and consolidates them into a single column.

Break Up the Values of a Column into Multiple Columns

Imagine you have a table where each row is a person followed by a fruit that they ate that day. You want to take the fruit column and create a separate column for each fruit.

Here is a sample table:

```
SELECT * FROM fruits;
```

id	name	fruit
1 Henry strawberries		
2 Henry grapefruit		
3 Henry watermelon		
4 Lily strawberries		
5 Lily watermelon		
6 Lily strawberries		
7 Lily watermelon		

Expected output:

name	strawberries	grapefruit	watermelon
Henry	1	1	1
Lily	2	0	2

Use the PIVOT operation in *Oracle* and *SQL Server*:

```
-- Oracle
SELECT *
FROM fruits
PIVOT
(COUNT(id) FOR fruit IN ('strawberries',
                           'grapefruit', 'watermelon'));

-- SQL Server
SELECT *
FROM fruits
PIVOT
(COUNT(id) FOR fruit IN ([strawberries],
                           [grapefruit], [watermelon])
 ) AS fruits_pivot;
```

Within the PIVOT section, the id and fruit columns are referenced, but the name column is not. Therefore, the name column will stay as its own column in the final result and each fruit will be turned into a new column.

The values of the table are the count of the number of rows in the original table that contained each particular name/fruit combination.

PIVOT Alternative: CASE

A more manual way of doing a PIVOT is to use a **CASE statement** instead in MySQL, PostgreSQL, and SQLite since they do not support **PIVOT**.

```
SELECT name,
       SUM(CASE WHEN fruit = 'strawberries'
                  THEN 1 ELSE 0 END) AS strawberries,
       SUM(CASE WHEN fruit = 'grapefruit'
                  THEN 1 ELSE 0 END) AS grapefruit,
       SUM(CASE WHEN fruit = 'watermelon'
                  THEN 1 ELSE 0 END) AS watermelon
  FROM fruits
 GROUP BY name
 ORDER BY name;
```

List the Values of Multiple Columns in a Single Column

Imagine you have a table where each row is a person followed by multiple columns that contain their favorite fruits. You want to rearrange the data so that all of the fruits are in one column.

Here is a sample table:

```
SELECT * FROM favorite_fruits;
```

id	name	fruit_one	fruit_two	fruit_thr
1	Anna	apple	banana	
2	Barry	raspberry		
3	Liz	lemon	lime	orange
4	Tom	peach	pear	plum

Expected output:

id	name	fruit	rank
1	Anna	apple	1
1	Anna	banana	2
2	Barry	raspberry	1
3	Liz	lemon	1
3	Liz	lime	2
3	Liz	orange	3
4	Tom	peach	1
4	Tom	pear	2
4	Tom	plum	3

Use the UNPIVOT operation in *Oracle* and *SQL Server*:

```
-- Oracle
SELECT *
FROM favorite_fruits
UNPIVOT
(fruit FOR rank IN (fruit_one AS 1,
```

```
        fruit_two AS 2,  
        fruit_thr AS 3));  
  
-- SQL Server  
SELECT *  
FROM favorite_fruits  
UNPIVOT  
(fruit FOR rank IN (fruit_one,  
                      fruit_two,  
                      fruit_thr)  
) AS fruit_unpivot  
WHERE fruit <> '';
```

The UNPIVOT section takes the columns `fruit_one`, `fruit_two`, and `fruit_thr` and consolidates them into a single column called `fruit`.

Once that's done, you can go ahead and use a typical SELECT statement to pull the original `id` and `name` columns along with the newly created `fruit` column.

UNPIVOT Alternative: UNION ALL

A more manual way of doing an UNPIVOT is to use UNION ALL instead in *MySQL*, *PostgreSQL*, and *SQLite* since they do not support UNPIVOT.

```
WITH all_fruits AS  
(SELECT id, name,  
       fruit_one as fruit,  
       1 AS rank  
  FROM favorite_fruits  
UNION ALL  
  SELECT id, name,  
         fruit_two as fruit,  
         2 AS rank  
    FROM favorite_fruits  
UNION ALL  
  SELECT id, name,  
         fruit_three as fruit,  
         3 AS rank  
    FROM favorite_fruits)
```

```
SELECT *
FROM all_fruits
WHERE fruit <> ''
ORDER BY id, name, fruit;
```

MySQL does not support inserting a constant into a column within a query (1 AS rank, 2 AS rank, and 3 AS rank). Remove those lines for the code to run.

Working with Multiple Tables and Queries

This chapter covers how to bring together multiple tables by either joining them or using union operators, and also how to work with multiple queries using common table expressions.

Table 9-1 includes descriptions and code examples of the three concepts covered in this chapter.

Table 9-1. Working with multiple tables and queries

Concept	Description	Code Example
Joining Tables	Combine the columns of two tables based on matching rows.	<pre>SELECT c.id, l.city FROM customers c INNER JOIN loc l ON c.lid = l.id;</pre>
Union Operators	Combine the rows of two tables based on matching columns.	<pre>SELECT name, city FROM employees; UNION SELECT name, city FROM customers;</pre>

Concept	Description	Code Example
Common Table Expressions	Temporarily save the output of a query, for another query to reference it. Also includes recursive and hierarchical queries.	<pre>WITH my_cte AS (SELECT name, SUM(order_id) AS num_orders FROM customers GROUP BY name) SELECT MAX(num_orders) FROM my_cte;</pre>

Joining Tables

In SQL, *joining* means combining data from multiple tables together within a single query. The following two tables list the state a person lives in and the pets they own:

-- states			-- pets		
name	state		name	pet	
Ada	AZ		Deb	dog	
Deb	DE		Deb	duck	

Use the `JOIN` clause to join the two tables into one table:

```
SELECT *
  FROM states s INNER JOIN pets p
    ON s.name = p.name;
```

name	state	name	pet
Deb	DE	Deb	dog
Deb	DE	Deb	duck

The resulting table only includes rows for Deb since she is present in both tables.

The left two columns are from the states table and the right two are from the pets table. The columns in the output can be referenced using the aliases s.name, s.state, p.name, and p.pet.

Breaking Down the JOIN Clause

```
states s INNER JOIN pets p ON s.name = p.name
```

Tables (states, pets)

The tables we would like to combine.

Aliases (s, p)

These are nicknames for the tables. This is optional, but recommended for simplicity. Without aliases, the ON clause could be written as states.name = pets.name.

Join Type (INNER JOIN)

The INNER portion specifies that only matching rows should be returned. If only JOIN is written, then it defaults to an INNER JOIN. Other join types can be found in [Table 9-2](#).

Join Condition (ON s.name = p.name)

The condition that must be true in order for two rows to be considered matching. Equal (=) is the most common operator, but others can be used as well including not equal (!= or <>), >, <, BETWEEN, etc.

In addition to the INNER JOIN, [Table 9-2](#) lists the various types of joins in SQL. The following query shows the general format for joining tables together:

```
SELECT *
FROM states s [JOIN_TYPE] pets p
ON s.name = p.name;
```

Replace the bolded [JOIN_TYPE] portion with the keywords in the Keyword column to get the results shown in the Resulting Rows column. For the CROSS JOIN join type, exclude the ON clause to get the results shown in the table.

Table 9-2. Ways to join together tables

Keyword	Description	Resulting Rows																												
JOIN	Defaults to an INNER JOIN.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> </tbody> </table>	nm	st	nm	pt	Deb	DE	Deb	dog	Deb	DE	Deb	duck																
nm	st	nm	pt																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
INNER JOIN	Returns the rows in common.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> </tbody> </table>	nm	st	nm	pt	Deb	DE	Deb	dog	Deb	DE	Deb	duck																
nm	st	nm	pt																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
LEFT JOIN	Returns the rows in the left table and the matching rows in the other table.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Ada</td><td>AZ</td><td>NULL</td><td>NULL</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> </tbody> </table>	nm	st	nm	pt	Ada	AZ	NULL	NULL	Deb	DE	Deb	dog	Deb	DE	Deb	duck												
nm	st	nm	pt																											
Ada	AZ	NULL	NULL																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
RIGHT JOIN	Returns the rows in the right table and the matching rows in the other table.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> <tr> <td>NULL</td><td>NULL</td><td>Pat</td><td>pig</td></tr> </tbody> </table>	nm	st	nm	pt	Deb	DE	Deb	dog	Deb	DE	Deb	duck	NULL	NULL	Pat	pig												
nm	st	nm	pt																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
NULL	NULL	Pat	pig																											
FULL OUTER JOIN	Returns the rows in both tables.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Ada</td><td>AZ</td><td>NULL</td><td>NULL</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> <tr> <td>NULL</td><td>NULL</td><td>Pat</td><td>pig</td></tr> </tbody> </table>	nm	st	nm	pt	Ada	AZ	NULL	NULL	Deb	DE	Deb	dog	Deb	DE	Deb	duck	NULL	NULL	Pat	pig								
nm	st	nm	pt																											
Ada	AZ	NULL	NULL																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
NULL	NULL	Pat	pig																											
CROSS JOIN	Returns all combinations of rows in the two tables.	<table border="1"> <thead> <tr> <th>nm</th><th>st</th><th>nm</th><th>pt</th></tr> </thead> <tbody> <tr> <td>Ada</td><td>AZ</td><td>Deb</td><td>dog</td></tr> <tr> <td>Ada</td><td>AZ</td><td>Deb</td><td>duck</td></tr> <tr> <td>Ada</td><td>AZ</td><td>Pat</td><td>pig</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>dog</td></tr> <tr> <td>Deb</td><td>DE</td><td>Deb</td><td>duck</td></tr> <tr> <td>Deb</td><td>DE</td><td>Pat</td><td>pig</td></tr> </tbody> </table>	nm	st	nm	pt	Ada	AZ	Deb	dog	Ada	AZ	Deb	duck	Ada	AZ	Pat	pig	Deb	DE	Deb	dog	Deb	DE	Deb	duck	Deb	DE	Pat	pig
nm	st	nm	pt																											
Ada	AZ	Deb	dog																											
Ada	AZ	Deb	duck																											
Ada	AZ	Pat	pig																											
Deb	DE	Deb	dog																											
Deb	DE	Deb	duck																											
Deb	DE	Pat	pig																											

In addition to joining tables using the standard `JOIN ... ON ...` syntax, **Table 9-3** lists others ways to join tables in SQL.

Table 9-3. Syntax to join together tables

Type	Description	Code
JOIN ... ON ... Syntax	Most common join syntax that works with INNER JOIN, LEFT JOIN, RIGHT JOIN, FULL OUTER JOIN, and CROSS JOIN.	<pre>SELECT * FROM states s INNER JOIN pets p ON s.name = p.name;</pre>
USING Shortcut	Use USING instead of the ON clause if the names of the columns that you are joining on match.	<pre>SELECT * FROM states INNER JOIN pets USING (name);</pre>
NATURAL JOIN Shortcut	Use NATURAL JOIN instead of INNER JOIN if the names of all of the columns that you are joining on match.	<pre>SELECT * FROM states NATURAL JOIN pets;</pre>
Old Join Syntax	Return all the combinations of the rows in two tables. Equivalent to a CROSS JOIN.	<pre>SELECT * FROM states s, pets p WHERE s.name = p.name;</pre>

Type	Description	Code
Self Join	Use either the old join or new join syntax to return all the combinations of the rows in a table with itself.	>SELECT * FROM states s1, states s2 WHERE s1.region = s2.region; SELECT * FROM states s1 INNER JOIN states s2 WHERE s1.region = s2.region;

The following sections describe the concepts in Tables 9-2 and 9-3 in detail.

Join Basics and INNER JOIN

This section walks through how a join works conceptually, as well as the basic join syntax using an INNER JOIN.

Join basics

You can think of joining tables in two steps:

1. Display all combinations of rows in the tables.
2. Filter on the rows that have matching values.

Here are two tables we'd like to join:

-- states			-- pets		
+	-----+	-----+	+	-----+	-----+
	name	state		name	pet
+	-----+	-----+	+	-----+	-----+
	Ada	AZ		Deb	dog
	Deb	DE		Deb	duck
+	-----+	-----+	+	Pat	pig
			+	-----+	-----+

Step 1: Display all combinations of rows.

By listing the table names in the `FROM` clause, all possible combinations of rows from the two tables are returned.

```
SELECT *
FROM states, pets;
```

name	state	name	pet
Ada	AZ	Deb	dog
Deb	DE	Deb	dog
Ada	AZ	Deb	duck
Deb	DE	Deb	duck
Ada	AZ	Pat	pig
Deb	DE	Pat	pig

The `FROM states, pets` syntax is an older way of doing a join in SQL. A more modern way of doing the same thing is using a **CROSS JOIN**.

Step 2: Filter on the rows that have matching names.

You likely don't want to display all combinations of rows in the two tables, but rather only situations where the `name` column of both tables match.

```
SELECT *
FROM states s, pets p
WHERE s.name = p.name;
```

name	state	name	pet
Deb	DE	Deb	dog
Deb	DE	Deb	duck

The row `Deb/DE` is listed twice because it matched two `Deb` values in the `pets` table.

The preceding code is equivalent to an **INNER JOIN**.

NOTE

The two-step process described previously is purely conceptual. Databases will rarely do a cross join when executing a join, but instead do it in a more optimized way.

However, thinking in these conceptual terms will help you correctly write join queries and understand their results.

INNER JOIN

The most common way to join together two tables is using an **INNER JOIN**, which returns rows that are in both tables.

Use INNER JOIN to only return people in both tables

```
SELECT *
FROM states s INNER JOIN pets p
ON s.name = p.name;
```

+-----+	-----+	+-----+	+-----+	+-----+
name	state	name	pet	
Deb	DE	Deb	dog	
Deb	DE	Deb	duck	
+-----+	-----+	+-----+	+-----+	+-----+

Join together more than two tables

This can be done by including additional sets of the **JOIN .. ON ..** keywords:

```
SELECT *
FROM states s
INNER JOIN pets p
ON s.name = p.name
INNER JOIN lunch l
ON s.name = l.name;
```

Join on more than one column

This can be done by including additional conditions within the ON clause. Imagine you want to join the following tables on both name and age:

```
-- states_ages          -- pets_ages
+-----+-----+-----+ +-----+-----+-----+
| name | state | age | | name | pet | age |
+-----+-----+-----+ +-----+-----+-----+
| Ada  | AK   | 25 | | Ada  | ant | 30 |
| Ada  | AZ   | 30 | | Pat  | pig | 45 |
+-----+-----+-----+ +-----+-----+-----+
SELECT *
FROM states_ages s INNER JOIN pets_ages p
  ON s.name = p.name
  AND s.age = p.age;

+-----+-----+-----+-----+-----+
| name | state | age | name | pet | age |
+-----+-----+-----+-----+-----+
| Ada  | AZ   | 30 | Ada  | ant | 30 |
+-----+-----+-----+-----+-----+
```

LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN

Use LEFT JOIN, RIGHT JOIN, and FULL OUTER JOIN to bring together rows from two tables, including ones that don't appear in both tables.

LEFT JOIN

Use LEFT JOIN to return all people in the states table. People in the states table that are not in the pets table get returned with NULL values.

```
SELECT *
FROM states s LEFT JOIN pets p
  ON s.name = p.name;

+-----+-----+-----+
| name | state | name | pet |
+-----+-----+-----+
```

Ada	AZ	NULL	NULL	
Deb	DE	Deb	dog	
Deb	DE	Deb	duck	

A `LEFT JOIN` is equivalent to a `LEFT OUTER JOIN`.

RIGHT JOIN

Use `RIGHT JOIN` to return all people in the `pets` table. People in the `pets` table that are not in the `states` table get returned with `NULL` values.

```
SELECT *
FROM states s RIGHT JOIN pets p
    ON s.name = p.name;
```

+-----+	-----+	+-----+	+-----+	+-----+
name	state	name	pet	
+-----+	-----+	+-----+	+-----+	+-----+
Deb	DE	Deb	dog	
Deb	DE	Deb	duck	
NULL	NULL	Pat	pig	
+-----+	-----+	+-----+	+-----+	+-----+

A `RIGHT JOIN` is equivalent to a `RIGHT OUTER JOIN`.

SQLite does not support `RIGHT JOIN`.

TIP

The `LEFT JOIN` is much more common than the `RIGHT JOIN`. If a `RIGHT JOIN` is needed, swap the two tables within the `FROM` clause and do a `LEFT JOIN` instead.

FULL OUTER JOIN

Use `FULL OUTER JOIN` to return all people in both the `states` and `pets` tables. Missing values from both tables are returned with `NULL` values.

```
SELECT *
FROM states s FULL OUTER JOIN pets p
    ON s.name = p.name;
```

name	state	name	pet
Ada	AZ	NULL	NULL
Deb	DE	Deb	dog
Deb	DE	Deb	duck
NULL	NULL	Pat	pig

A FULL OUTER JOIN is equivalent to a FULL JOIN.

MySQL and *SQLite* do not support FULL OUTER JOIN.

USING and NATURAL JOIN

When joining tables together, to save on typing, you can use the USING or NATURAL JOIN shortcuts instead of the standard JOIN .. ON .. syntax.

USING

MySQL, *Oracle*, *PostgreSQL*, and *SQLite* support the USING clause.

You can use the USING shortcut in place of the ON clause to join on two columns of the exact same name. The join must be an equi-join (= in the ON clause) to use USING.

```
-- ON clause
SELECT *
FROM states s INNER JOIN pets p
    ON s.name = p.name;
```

name	state	name	pet
Deb	DE	Deb	dog
Deb	DE	Deb	duck

```
+-----+-----+-----+-----+
-- Equivalent USING shortcut
SELECT *
FROM states INNER JOIN pets
    USING (name);

+-----+-----+-----+
| name | state | pet   |
+-----+-----+-----+
| Deb  | DE    | dog   |
| Deb  | DE    | duck  |
+-----+-----+-----+
```

The difference between the two queries is that the first query returns four columns including `s.name` and `p.name`, while the second query returns three columns because the two `name` columns get merged together as one and is simply called `name`.

NATURAL JOIN

MySQL, Oracle, PostgreSQL, and SQLite support a NATURAL JOIN.

You can use the NATURAL JOIN shortcut in place of the INNER JOIN .. ON .. syntax to join two tables based on all columns of the exact same name. The join must be an equi-join (= in the ON clause) to use a NATURAL JOIN.

```
-- INNER JOIN ... ON ... AND ...
SELECT *
FROM states_ages s NATURAL JOIN pets_ages p
    ON s.name = p.name
    AND s.age = p.age;
```

```
+-----+-----+-----+-----+
| name | state | age  | name | pet   | age  |
+-----+-----+-----+-----+
| Ada  | AZ   | 30   | Ada  | ant   | 30   |
+-----+-----+-----+-----+
```

```
-- Equivalent NATURAL JOIN shortcut
SELECT *
FROM states_ages NATURAL JOIN pets_ages;

+-----+-----+-----+
| name | age  | state | pet  |
+-----+-----+-----+
| Ada  | 30   | AZ    | ant  |
+-----+-----+-----+
```

The difference between the two queries is that the first query returns six columns including `s.name`, `s.age`, `p.name`, and `p.age`, while the second query returns four columns because the duplicate `name` and `age` columns get merged together and are simply called `name` and `age`.

WARNING

Be careful when using a `NATURAL JOIN`. It saves quite a bit of typing, but can do an unexpected join if a column of a matching name is added or removed from a table. It is better to use for quick queries versus production code.

CROSS JOIN and Self Join

Another way of joining tables together is by displaying all combinations of the rows in two tables. This can be done with a `CROSS JOIN`. If this is done on a table with itself, it is called a *self join*. A self join is useful when you want to compare rows within the same table.

CROSS JOIN

Use `CROSS JOIN` to return all combinations of the rows in two tables. It is equivalent to listing out the tables in the `FROM` clause (which is sometimes referred to as “old join syntax”).

```
-- CROSS JOIN
SELECT *
```

```
FROM states CROSS JOIN pets;  
  
-- Equivalent table list  
SELECT *  
FROM states, pets;
```

name	state	name	pet
Ada	AZ	Deb	dog
Deb	DE	Deb	dog
Ada	AZ	Deb	duck
Deb	DE	Deb	duck
Ada	AZ	Pat	pig
Deb	DE	Pat	pig

Once all combinations are listed out, you can choose to filter on the results by adding a WHERE clause to return fewer rows based on what you're looking for.

Self join

You can join a table with itself using a self join. There are typically two steps to a self join:

1. Display all combinations of the rows in a table with itself.
2. Filter on the resulting rows based on some criteria.

The following are two examples of self joins in practice.

Here is a table of employees and their managers:

```
SELECT * FROM employee;
```

dept	emp_id	emp_name	mgr_id
tech	201	lisa	101
tech	202	monica	101
data	203	nancy	201

data	204	olivia	201
data	205	penny	202

Example 1: Return a list of employees and their managers.

```
SELECT e1.emp_name, e2.emp_name as mgr_name
FROM employee e1, employee e2
WHERE e1.mgr_id = e2.emp_id;
```

emp_name	mgr_name
nancy	lisa
olivia	lisa
penny	monica

Example 2: Match each employee with another employee in their department.

```
SELECT e.dept, e.emp_name, matching_emp.emp_name
FROM employee e, employee matching_emp
WHERE e.dept = matching_emp.dept
      AND e.emp_name <> matching_emp.emp_name;
```

dept	emp_name	emp_name
tech	monica	lisa
tech	lisa	monica
data	penny	nancy
data	olivia	nancy
data	penny	olivia
data	nancy	olivia
data	olivia	penny
data	nancy	penny

NOTE

The preceding query has duplicate rows (`monica/lisa` and `lisa/monica`). To remove the duplicates and return just four rows instead of eight, you can add the line:

```
AND e.emp_name < matching_emp.emp_name
```

to the `WHERE` clause to only return rows where the first name is before the second name alphabetically. Here is the output without duplicates:

dept	emp_name	emp_name
tech	lisa	monica
data	nancy	olivia
data	nancy	penny
data	olivia	penny

Union Operators

Use the `UNION` keyword to combine the results of two or more `SELECT` statements. The difference between a `JOIN` and a `UNION` is that `JOIN` links together multiple tables within a single query, whereas `UNION` stacks the results of multiple queries:

```
-- JOIN example
SELECT *
FROM birthdays b JOIN candles c
    ON b.name = c.name;
```

```
-- UNION example
SELECT * FROM writers
UNION
SELECT * FROM artists;
```

Figure 9-1 shows the difference between the results of a `JOIN` and a `UNION`, based on the preceding code.

JOIN

birthdays		candles		birthdays_and_candles	
Name	Birthday	Name	Candles	Name	Birthday
Molly	1/1/84	Molly	37	Molly	1/1/84
Polly	2/14/05	Polly	16	Polly	2/14/05

UNION

writers		artists		writers_and_artists	
Name	Birthday	Name	Birthday	Name	Birthday
Molly	1/1/84	Cindy	7/4/90	Molly	1/1/84
Polly	2/14/05	Mindy	8/9/80	Polly	2/14/05
+		=			

Figure 9-1. JOIN versus UNION

There are three ways to combine the rows of two tables together. These are also known as *union operators*:

UNION

Combines the results of multiple statements.

EXCEPT (MINUS in Oracle)

Returns the results minus another set of results.

INTERSECT

Returns overlapping results.

UNION

The UNION keyword combines the results of two or more SELECT statements into one output.

Here are two tables we'd like to combine:

-- staff	
name	origin
michael	NULL
janet	NULL
tahani	england

-- residents		
name	country	occupation
eleanor	usa	temp
chidi	nigeria	professor
tahani	england	model
jason	usa	dj

Use `UNION` to combine the two tables and eliminate any duplicate rows:

```
SELECT name, origin FROM staff
UNION
SELECT name, country FROM residents;
```

name	origin
michael	NULL
janet	NULL
tahani	england
eleanor	usa
chidi	nigeria
jason	usa

Note that `tahani/england` appears in both the `staff` and `residents` tables. However, it only shows up as one row in the

result set because UNION removes duplicate rows from the output.

Which Queries Can You Union Together?

When doing a UNION on two queries, some characteristics of the queries must match and others do not have to match.

Number of Columns: MUST MATCH

When you union together two queries, you must specify the same number of columns in both queries.

Column Names: DO NOT HAVE TO MATCH

The column names of the two queries do not need to match to do a UNION. The column names used in the first SELECT statement in a UNION query become the names of the output columns.

Data Types: MUST MATCH

The data types of the two queries need to match to do a UNION. If they do not match, you can use the **CAST** function to cast them into the same data type before doing a UNION.

UNION ALL

Use UNION ALL to combine the two tables and preserve duplicate rows:

```
SELECT name, origin FROM staff  
UNION ALL  
SELECT name, country FROM residents;
```

name	origin
michael	NULL
janet	NULL
tahani	england
leanor	usa
chidi	nigeria

tahani	england	
jason	usa	

TIP

If you know with certainty that no duplicate rows are possible, use UNION ALL to improve performance. UNION does an additional sort behind the scenes to identify the duplicates.

UNION with other clauses

You can also include other clauses when using a UNION, such as WHERE, JOIN, etc. However, only one ORDER BY clause is allowed for the whole query, and it should be at the very end.

Filter out null values and sort the results of a UNION query:

```
SELECT name, origin  
FROM staff  
WHERE origin IS NOT NULL
```

UNION

```
SELECT name, country  
FROM residents
```

```
ORDER BY name;
```

name	origin	
chidi	nigeria	
eleanor	usa	
jason	usa	
tahani	england	

UNION with more than two tables

You can union together more than two tables by including additional UNION clauses.

Combine the rows of more than two tables:

```
SELECT name, origin  
FROM staff
```

UNION

```
SELECT name, country  
FROM residents
```

UNION

```
SELECT name, country  
FROM pets;
```

TIP

UNION is typically used to combine results from multiple tables. If you are combining results from a single table, it is better to write a single query instead and use the appropriate WHERE clause, CASE statement, etc.

EXCEPT and INTERSECT

In addition to using a UNION to combine the rows of multiple tables, you can use EXCEPT and INTERSECT to combine the rows in different ways.

EXCEPT

Use EXCEPT to “subtract” the results of one query from another query.

Return the staff members that are not residents:

```
SELECT name FROM staff  
EXCEPT  
SELECT name FROM residents;
```

```
+-----+  
| name |  
+-----+  
| michael |  
| janet |  
+-----+
```

MySQL does not support EXCEPT. Instead, you can use the **NOT IN** keywords as a workaround:

```
SELECT name  
FROM staff  
WHERE name NOT IN (SELECT name FROM residents);
```

Oracle uses MINUS instead of EXCEPT.

PostgreSQL also supports EXCEPT ALL, which does not remove duplicates. EXCEPT removes all occurrences of a value, while EXCEPT ALL removes specific instances.

INTERSECT

Use INTERSECT to find the rows in common between two queries.

Return the staff members that are residents as well:

```
SELECT name, origin FROM staff  
INTERSECT  
SELECT name, country FROM residents;
```

```
+-----+-----+  
| name | origin |  
+-----+-----+  
| tahani | england |  
+-----+-----+
```

MySQL does not support INTERSECT. Instead, you can use an **INNER JOIN** as a workaround:

```
SELECT s.name, s.origin  
FROM staff s INNER JOIN residents r  
ON s.name = r.name;
```

PostgreSQL also supports `INTERSECT ALL`, which preserves duplicate values.

Union Operators: Order of Evaluation

When writing a statement with multiple union operators (`UNION`, `EXCEPT`, `INTERSECT`), use parentheses to specify the order in which the operations should occur.

```
SELECT * FROM staff  
EXCEPT  
(SELECT * FROM residents  
UNION  
SELECT * FROM pets);
```

Unless otherwise specified, union operators are performed in top-down order, except that `INTERSECT` takes precedence over `UNION` and `EXCEPT`.

Common Table Expressions

A *common table expression (CTE)* is a temporary result set. In other words, it temporarily saves the output of a query for you to write other queries that reference it.

You can spot a CTE when you see the `WITH` keyword. There are two types of CTEs:

Nonrecursive CTE

A query for other queries to reference (see “[CTEs Versus Subqueries](#)” on page 293).

Recursive CTE

A query that references itself (see “[Recursive CTEs](#)” on page 295).

NOTE

Nonrecursive CTEs are a lot more common than recursive CTEs. Most of the time, if someone mentions a CTE, they are referring to a nonrecursive CTE.

Here is an example of a nonrecursive CTE in practice:

```
-- Query the results of my_cte
WITH my_cte AS (
    SELECT name, AVG(grade) AS avg_grade
    FROM my_table
    GROUP BY name)

SELECT *
FROM my_cte
WHERE avg_grade < 70;
```

Here is an example of a recursive CTE in practice:

```
-- Generate the numbers 1 through 10
WITH RECURSIVE my_cte(n) AS
(
    SELECT 1 -- Include FROM dual in Oracle
    UNION ALL
    SELECT n + 1 FROM my_cte WHERE n < 10
)

SELECT * FROM my_cte;
```

In *MySQL* and *PostgreSQL*, the `RECURSIVE` keyword is required. In *Oracle* and *SQL Server*, the `RECURSIVE` keyword must be left out. *SQLite* works with either syntax.

In *Oracle*, you may see older code that uses the `CONNECT BY` syntax for recursive queries, but CTEs are much more common these days.

CTEs Versus Subqueries

Both CTEs and subqueries allow you to write a query, and then write another query that references the first query. This section describes the difference between the two approaches.

Imagine your goal is to find the department that has the largest average salary. This can be done in two steps: write a query that returns the average salary for each department; use a CTE or subquery to write a query around the first query to return the department with the largest average salary.

Step 1. Query that finds the average salary for each department

```
SELECT dept, AVG(salary) AS avg_salary
FROM employees
GROUP BY dept;
```

dept	avg_salary
mktg	78000
sales	61000
tech	83000

Step 2. CTE and subquery that find the department with the largest average salary using the preceding query

```
-- CTE approach
WITH avg_dept_salary AS (
    SELECT dept, AVG(salary) AS avg_salary
    FROM employees
    GROUP BY dept)

SELECT *
FROM avg_dept_salary
ORDER BY avg_salary DESC
LIMIT 1;

-- Equivalent subquery approach
SELECT *
FROM
```

```
(SELECT dept, AVG(salary) AS avg_salary
FROM employees
GROUP BY dept) avg_dept_salary

ORDER BY avg_salary DESC
LIMIT 1;

+-----+-----+
| dept | avg_salary |
+-----+-----+
| tech |      83000 |
+-----+-----+
```

The `LIMIT` clause syntax differs by software. Replace `LIMIT 1` with `ROWNUM = 1` in *Oracle* and `TOP 1` in *SQL Server*. More details can be found in [The LIMIT Clause](#) section in [Chapter 4](#).

Advantages of a CTE Versus a Subquery

There are a few advantages to using a CTE instead of a subquery.

Multiple References

Once a CTE is defined, you can reference it by name multiple times within the `SELECT` queries that follow:

```
WITH my_cte AS (...)

    SELECT * FROM my_cte WHERE id > 10
    UNION
    SELECT * FROM my_cte WHERE score > 90;
```

With a subquery, you would need to write out the full subquery each time.

Multiple Tables

CTE syntax is more readable when working with multiple tables because you can list all the CTEs up front:

```
WITH my_cte1 AS (...),
     my_cte2 AS (...)

    SELECT *
```

```
FROM my_cte1 m1
    INNER JOIN my_cte2 m2
        ON m1.id = m2.id;
```

With subqueries, the subqueries would be scattered throughout the overall query.

CTEs are *not supported* in older SQL software, which is why subqueries are still commonly used.

Recursive CTEs

This section walks through two practical situations where a recursive CTE would be useful.

Fill in missing rows in a sequence of data

The following table includes dates and prices. Note that the date column is missing data for the second and fifth of the month.

```
SELECT * FROM stock_prices;
```

date	price
2021-03-01	668.27
2021-03-03	678.83
2021-03-04	635.40
2021-03-06	591.01

Fill in the dates with a two-step process:

1. Use a recursive CTE to generate a sequence of dates.
2. Join the sequence of dates with the original table.

NOTE

The following code runs in MySQL. **Table 9-4** has the syntax for each RDBMS.

Step 1: Use a recursive CTE to generate a sequence of dates called my_dates.

The `my_dates` table starts with the date `2021-03-01`, and adds on the next date again and again, up until the date `2021-03-06`:

```
-- MySQL syntax
WITH RECURSIVE my_dates(dt) AS (
    SELECT '2021-03-01'
    UNION ALL
    SELECT dt + INTERVAL 1 DAY
    FROM my_dates
    WHERE dt < '2021-03-06')

SELECT * FROM my_dates;
```

dt
2021-03-01
2021-03-02
2021-03-03
2021-03-04
2021-03-05
2021-03-06

Step 2: Left join the recursive CTE with the original table.

```
-- MySQL syntax
WITH RECURSIVE my_dates(dt) AS (
    SELECT '2021-03-01'
    UNION ALL
    SELECT dt + INTERVAL 1 DAY
    FROM my_dates
    WHERE dt < '2021-03-06')
```

```
SELECT d.dt, s.price
FROM my_dates d
    LEFT JOIN stock_prices s
        ON d.dt = s.date;
```

dt	price
2021-03-01	668.27
2021-03-02	NULL
2021-03-03	678.83
2021-03-04	635.40
2021-03-05	NULL
2021-03-06	591.01

Step 3 (Optional): Fill in the null values with the previous day's price.

Replace the SELECT clause (SELECT d.dt, s.price) with:

```
SELECT d.dt, COALESCE(s.price,
    LAG(s.price) OVER
        (ORDER BY d.dt)) AS price
...
```

dt	price
2021-03-01	668.27
2021-03-02	668.27
2021-03-03	678.83
2021-03-04	635.40
2021-03-05	635.40
2021-03-06	591.01

There are syntax differences for each RDBMS.

Here is the general syntax for generating a date column. The bolded portions differ by RDBMS, and the software-specific code is listed in [Table 9-4](#).

```
[WITH] my_dates(dt) AS (
    SELECT [DATE]
    UNION ALL
```

```

SELECT [DATE PLUS ONE]
FROM my_dates
WHERE dt < [LAST DATE])

SELECT * FROM my_dates;

```

Table 9-4. Generating a date column in each RDBMS

RDBMS	WITH	DATE	DATE PLUS ONE	LAST DATE
MySQL	WITH RECUR SIVE	'2021-03-01' INTERVAL 1 DAY	dt +	'2021-03-06'
Oracle	WITH	DATE '2021-03-01'	dt + INTERVAL '1' DAY	DATE '2021-03-06'
PostgreSQL	WITH RECUR SIVE	CAST('2021-03-01' AS DATE)	CAST(dt + INTERVAL '1 day' AS DATE)	'2021-03-06'
SQL Server	WITH	CAST('2021-03-01' AS DATE)	DATEADD(DAY, 1, CAST(dt AS DATE))	'2021-03-06'
SQLite	WITH RECUR SIVE	DATE('2021-03-01') INTERVAL 1 DAY	DATE(dt, '1 day')	'2021-03-06'

Return all the parents of a child row

The following table includes the roles of various family members. The rightmost column includes the `id` of a person's parent.

```
SELECT * FROM family_tree;
```

id name role parent_id			
1 Lao Ye Grandpa NULL			
2 Lao Lao Grandma NULL			

	3	Ollie	Dad	NULL	
	4	Alice	Mom	1	
	4	Alice	Mom	2	
	5	Henry	Son	3	
	5	Henry	Son	4	
	6	Lily	Daughter	3	
	6	Lily	Daughter	4	

NOTE

The following code runs in MySQL. [Table 9-5](#) has the syntax for each RDBMS.

You can list each person's parents and grandparents with a recursive CTE:

```
-- MySQL syntax
WITH RECURSIVE my_cte (id, name, lineage) AS (
    SELECT id, name, name AS lineage
    FROM family_tree
    WHERE parent_id IS NULL
    UNION ALL
    SELECT ft.id, ft.name,
           CONCAT(mc.lineage, ' > ', ft.name)
    FROM family_tree ft
    INNER JOIN my_cte mc
    ON ft.parent_id = mc.id)

SELECT * FROM my_cte ORDER BY id;
```

	id	name	lineage	
	1	Lao Ye	Lao Ye	
	2	Lao Lao	Lao Lao	
	3	Ollie	Ollie	
	4	Alice	Lao Ye > Alice	
	4	Alice	Lao Lao > Alice	

5	Henry	Ollie > Henry
5	Henry	Lao Ye > Alice > Henry
5	Henry	Lao Lao > Alice > Henry
6	Lily	Ollie > Lily
6	Lily	Lao Ye > Alice > Lily
6	Lily	Lao Lao > Alice > Lily

In the preceding code (also known as a *hierarchical query*), `my_cte` contains two statements that are unioned together:

- The first `SELECT` statement is the starting point. The rows where the `parent_id` is `NULL` are treated as the tree roots.
- The second `SELECT` statement defines the recursive link between the parent and child rows. The children of each tree root are returned and tacked on to the `lineage` column until the full lineage is spelled out.

There are syntax differences for each RDBMS.

Here is the general syntax for listing all the parents. The bolded portions differ by RDBMS, and the software-specific code is listed in [Table 9-5](#).

```
[WITH] my_cte (id, name, lineage) AS (
    SELECT id, name, [NAME] AS lineage
    FROM family_tree
    WHERE parent_id IS NULL
    UNION ALL
    SELECT ft.id, ft.name, [LINEAGE]
    FROM family_tree ft
        INNER JOIN my_cte mc
        ON ft.parent_id = mc.id)

SELECT * FROM my_cte ORDER BY id;
```

Table 9-5. Listing all the parents in each RDBMS

RDBMS	WITH	NAME	LINEAGE
MySQL	WITH RECURSIVE	name	CONCAT(mc.lineage, ' > ', ft.name)

RDBMS	WITH	NAME	LINEAGE
Oracle	WITH	name	mc.lineage ' > ' ft.name
PostgreSQL	WITH RECURSIVE	CAST(name AS VARCHAR(30))	CAST(CONCAT(mc.lineage, ' > ', ft.name) AS VARCHAR(30))
SQL Server	WITH	CAST(name AS VARCHAR(30))	CAST(CONCAT(mc.lineage, ' > ', ft.name) AS VARCHAR(30))
SQLite	WITH RECURSIVE	name	mc.lineage ' > ' ft.name

How Do I...?

This chapter is intended to be a quick reference for frequently asked SQL questions that combine multiple concepts:

- Find the rows containing duplicate values
- Select rows with the max value for another column
- Concatenate text from multiple fields into a single field
- Find all tables containing a specific column name
- Update a table where the ID matches another table

Find the Rows Containing Duplicate Values

The following table lists seven types of teas and the temperatures they should be steeped at. Note that there are two sets of duplicate tea/temperature values, which are in bold.

```
SELECT * FROM teas;
```

id	tea	temperature
1	green	170
2	black	200

	3		black		200	
	4		herbal		212	
	5		herbal		212	
	6		herbal		210	
	7		oolong		185	

This section covers two different scenarios:

- Return all unique tea/temperature combinations
- Return only the rows with duplicate tea/temperature values

Return All Unique Combinations

To exclude duplicate values and return only the unique rows of a table, use the **DISTINCT** keyword.

```
SELECT DISTINCT tea, temperature  
FROM teas;
```

tea	temperature
green	170
black	200
herbal	212
herbal	210
oolong	185

Potential extensions

To return the number of unique rows in a table, use the **COUNT** and **DISTINCT** keywords together. More details can be found in the DISTINCT section in [Chapter 4](#).

Return Only the Rows with Duplicate Values

The following query identifies the rows in the table with duplicate values.

```
WITH dup_rows AS (
    SELECT tea, temperature,
           COUNT(*) as num_rows
    FROM teas
   GROUP BY tea, temperature
  HAVING COUNT(*) > 1)

SELECT t.id, d.tea, d.temperature
  FROM teas t INNER JOIN dup_rows d
    ON t.tea = d.tea
   AND t.temperature = d.temperature;
```

id tea temperature
2 black 200
3 black 200
4 herbal 212
5 herbal 212

Explanation

The bulk of the work happens in the `dup_rows` query. All of the `tea/temperature` combinations are counted, and then only the combinations that occur more than once are kept with the `HAVING` clause. This is what `dup_rows` looks like:

tea temperature num_rows
black 200 2
herbal 212 2

The purpose of the `JOIN` in the second half of the query is to pull the `id` column back into the final output.

Keywords in the query

- `WITH dup_rows` is the start of a **common table expression**, which allows you to work with multiple `SELECT` statements within a single query.
- `HAVING COUNT(*) > 1` uses the **HAVING** clause, which allows you to filter on an **aggregation** like `COUNT()`.
- `teas t INNER JOIN dup_rows d` uses an **INNER JOIN**, which allows you to bring together the `teas` table and the `dup_rows` query.

Potential extensions

To delete particular duplicate rows from a table, use a **DELETE** statement. More details can be found in [Chapter 5](#)

Select Rows with the Max Value for Another Column

The following table lists employees and the number of sales they've made. You want to return each employee's most recent number of sales, which are in bold.

```
SELECT * FROM sales;
```

+-----+	+-----+	+-----+	+-----+	+-----+
id	employee	date	sales	
+-----+	+-----+	+-----+	+-----+	+-----+
1	Emma	2021-08-01	6	
2	Emma	2021-08-02	17	
3	Jack	2021-08-02	14	
4	Emma	2021-08-04	20	
5	Jack	2021-08-05	5	
6	Emma	2021-08-07	1	
+-----+	+-----+	+-----+	+-----+	+-----+

Solution

The following query returns the number of sales that each employee made on their most recent sale date (aka each employee's largest date value).

```
SELECT s.id, r.employee, r.recent_date, s.sales
FROM (SELECT employee, MAX(date) AS recent_date
      FROM sales
     GROUP BY employee) r
INNER JOIN sales s
  ON r.employee = s.employee
  AND r.recent_date = s.date;
```

id	employee	recent_date	sales
5	Jack	2021-08-05	5
6	Emma	2021-08-07	1

Explanation

The key to this problem is to break it down into two parts. The first goal is to identify the most recent sale date for each employee. This is what the output of the subquery `r` looks like:

employee	recent_date
Emma	2021-08-07
Jack	2021-08-05

The second goal is to pull the `id` and `sales` columns back into the final output, which is done using the `JOIN` in the second half of the query.

Keywords in the query

- `GROUP BY employee` uses the **GROUP BY** clause, which splits up the table by employee and finds the **MAX(date)** for each employee.
- `r INNER JOIN sales s` uses an **INNER JOIN**, which allows you to bring together the `r` subquery and the `sales` table.

Potential extensions

An alternative to the `GROUP BY` solution is to use a **window function** (`OVER ... PARTITION BY ...`) with a **FIRST_VALUE** function, which would return the same results. More details can be found in the “**Window Functions**” on page 250 section in Chapter 8.

Concatenate Text from Multiple Fields into a Single Field

This section covers two different scenarios:

- Concatenate text from fields *in a single row* into a single value
- Concatenate text from fields *in multiple rows* into a single value

Concatenate Text from Fields in a Single Row

The following table has two columns, and you want to concatenate them into one column.

id	name		id_name
1	Boots	--->	1_Boots
2	Pumpkin		2_Pumpkin
3	Tiger		3_Tiger

Use the `CONCAT` function or the concatenation operator (`||`) to bring together the values:

```
-- MySQL, PostgreSQL, and SQL Server
SELECT CONCAT(id, '_', name) AS id_name
FROM my_table;
```

```
-- Oracle, PostgreSQL, and SQLite
SELECT id || '_' || name AS id_name
FROM my_table;
```

id_name	
1_Boots	
2_Pumpkin	
3_Tiger	

Potential extensions

Chapter 7, *Operators and Functions*, covers other ways to work with string values in addition to `CONCAT`, including:

- Finding the length of a string
- Finding words in a string
- Extracting text from a string

Concatenate Text from Fields in Multiple Rows

The following table lists the calories burned by each person. You want to concatenate the calories for each person into a single row.

name	calories
ally	80
ally	75
ally	90
jess	100
jess	92

--->

name	calories
ally	80,75,90
jess	100,92

Use a function like GROUP_CONCAT, LISTAGG, ARRAY_AGG, or STRING_AGG to create the list.

```
SELECT name,
       GROUP_CONCAT(calories) AS calories_list
  FROM workouts
 GROUP BY name;
```

name	calories_list
ally	80,75,90
jess	100,92

This code works in *MySQL* and *SQLite*. Replace GROUP_CONCAT(calories) with the following in other RDBMSs:

Oracle

```
LISTAGG(calories, ',')
```

PostgreSQL

```
ARRAY_AGG(calories)
```

SQL Server

```
STRING_AGG(calories, ',')
```

Potential extensions

The [aggregate rows into a single value or list](#) section in [Chapter 8](#) includes details on how to use other separators besides the comma (,), how to sort the values, and how to return unique values.

Find All Tables Containing a Specific Column Name

Imagine you have a database with many tables. You want to quickly find all tables that contain a column name with the word city in it.

Solution

In most RDBMSs, there is a special table that contains all table names and column names. [Table 10-1](#) shows how to query that table in each RDBMS.

The last line of each code block is optional. You can include it if you want to narrow down the results for a particular database or user. If excluded, all tables will be returned.

Table 10-1. Find all tables containing a specific column name

RDBMS	Code
MySQL	<pre>SELECT table_name, column_name FROM information_schema.columns WHERE column_name LIKE '%city%' AND table_schema = 'my_db_name';</pre>
Oracle	<pre>SELECT table_name, column_name FROM all_tab_columns WHERE column_name LIKE '%CITY%' AND owner = 'MY_USER_NAME';</pre>
PostgreSQL, SQL Server	<pre>SELECT table_name, column_name FROM information_schema.columns WHERE column_name LIKE '%city%' AND table_catalog = 'my_db_name';</pre>

The output will display all column names that contain the term `city` along with the tables they are in:

TABLE_NAME	COLUMN_NAME
customers	city
employees	city
locations	metro_city

NOTE

SQLite does not have a table that contains all column names. Instead, you can manually show all tables and then view the column names within each table:

```
.tables  
pragma table_info(my_table);
```

Potential extensions

[Chapter 5, *Creating, Updating, and Deleting*](#), covers more ways to interact with databases and tables, including:

- Viewing existing databases
- Viewing existing tables
- Viewing the columns of a table

[Chapter 7, *Operators and Functions*](#), covers more ways to search for text in addition to `LIKE`, including:

- `=` to search for an exact match
- `IN` to search for multiple terms
- `Regular expressions` to search for a pattern

Update a Table Where the ID Matches Another Table

Imagine you have two tables: products and deals. You'd like to update the names in the deals table with the names of items in the products table that have a matching id.

```
SELECT * FROM products;
```

```
+-----+-----+
| id   | name        |
+-----+-----+
| 101  | Mac and cheese mix |
| 102  | MIDI keyboard    |
| 103  | Mother's day card |
+-----+-----+
```

```
SELECT * FROM deals;
```

```
+-----+-----+
| id   | name        |
+-----+-----+
| 102  | Tech gift    | --> MIDI keyboard
| 103  | Holiday card | --> Mother's day card
+-----+-----+
```

Solution

Use an **UPDATE** statement to modify values in a table using the **UPDATE ... SET ...** syntax. **Table 10-2** shows how to do this in each RDBMS.

Table 10-2. Update a table where the ID matches another table

RDBMS	Code
MySQL	<pre>UPDATE deals d, products p SET d.name = p.name WHERE d.id = p.id;</pre>

RDBMS	Code
Oracle	UPDATE deals d SET name = (SELECT p.name FROM products p WHERE d.id = p.id);
PostgreSQL, SQLite	UPDATE deals SET name = p.name FROM deals d INNER JOIN products p ON d.id = p.id WHERE deals.id = p.id;
SQL Server	UPDATE d SET d.name = p.name FROM deals d INNER JOIN products p ON d.id = p.id;

The `deals` table is now updated with the names from the `products` table:

```
SELECT * FROM deals;
```

```
+-----+-----+  
| id   | name           |  
+-----+-----+  
| 102  | MIDI keyboard |  
| 103  | Mother's day card |  
+-----+
```

WARNING

Once the `UPDATE` statement is executed, the results cannot be undone. The exception is if you start a `transaction` before executing the `UPDATE` statement.

Potential extensions

Chapter 5, *Creating, Updating, and Deleting*, covers more ways to modify tables, including:

- Updating a column of data
- Updating rows of data
- Updating rows of data with the results of a query
- Adding a column to a table

Final Words

This book covers the most popular concepts and keywords in SQL, but we've only scratched the surface. SQL can be used to perform many tasks, using a variety of different approaches. I encourage you to keep on learning and exploring.

You may have noticed that SQL syntax varies widely by RDBMS. Writing SQL code requires a lot of practice, patience, and looking up syntax. I hope you've found this pocket guide to be helpful for doing so.

Symbols

- != (inequality) operator, 183
- "" (double quotes)
 - column aliases, 59
 - enclosing identifiers, 44, 49, 155
 - enclosing text values in files, 113
- \$\$ (dollar signs), enclosing strings in PostgreSQL, 155
- % (percent sign)
 - modulo operator, 189
 - use with LIKE, 187, 188
- " (single quotes)
 - embedding single quote in string value, 154
 - enclosing string values, 49, 154
 - escape sequences in, 156
- () (parentheses)
 - enclosing subqueries, 61
 - indicating order of operations, 181
- * (asterisk)
 - multiplication operator, 189
 - selecting all columns, 55
- + (plus sign), addition operator, 179, 189
- , (comma) separator, 55, 112, 246, 311
- (minus sign) subtraction operator, 189
- comment for single line of code, 8, 48
- . (dot notation), 59
- / (forward slash) division operator, 189
- /* */ comment for multiple lines of code, 48
- ; (semicolon) ending SQL statements, 45
- < (less than) operator, 183
- <= (less than or equal to) operator, 183
- <=> (null-safe equality) operator in MySQL, 183
- <> (inequality) operator, 183
- = (equal sign)
 - equals operator, 179, 183
 - equi-joins, 280
 - most popular comparison operator, 47
- > (greater than) operator, 183

`>=` (greater than or equal to) operator, 183
`\` (backslash) escaping regular expressions in MySQL, 211
`\d` (digits) in regular expressions, 211, 216
`_` (underscore), use with LIKE, 188
`||` (concatenation) operator, 203, 309
`~` (tilde) limited regular expression support in PostgreSQL, 215

using when renaming columns, 44
using with table aliases, 61
ASCENDING (ASC) keyword, 86
ASCII versus Unicode encoding, 159
assignment operators, 191
atomicity, 138
attributes, 11
AUTOINCREMENT, 110
AVG function, 191
ROWS BETWEEN clause, using with, 259

A

absolute value, 194
aggregate functions, 191-193
aggregating rows into single value or list, 245, 311
`COUNT`, 80
multiple, in GROUP BY, 244
versus window functions, 250
aliases, 44
aliasing columns, 57
aliasing subqueries, 70
column versus table, 60
`ALL` keyword, 63
`ALTER` privileges, 115
`ALTER TABLE` statements, 115, 118
ambiguous column name (error), 60
`AND` operator, 75, 179, 180, 181
ANSI standards, 39-41
deciding whether to use in writing SQL code, 40
deciding which to use in SQL code, 41
`ARRAY_AGG` function, 191, 242, 310
`AS` keyword
column aliases, 58, 60

B

`BETWEEN` operator, 179, 184
binary data types, 175
binary values, 174
storing external files as, 174
bits, 154
bitwise operators, 191
`BLOB` data type, 176
Boolean data types, 173
bytea data type (PostgreSQL), 176
bytes, 154

C

`CASCADE` keyword, 128
caution with, 129
case insensitivity
keywords in SQL, 42
LIKE patterns in MySQL, SQL Server, and SQLite, 188
regular expressions in MySQL, 211
in SQL, 7
case sensitivity
aliases in double quotes, 59
LIKE patterns in Oracle and PostgreSQL, 188

regular expressions in PostgreSQL, 215

CASE statements, 237, 238-242

- alternative to PIVOT operation, 264
- using instead of Oracle's DECODE function, 40

case, changing for a string, 200

CAST function

- converting string to datetime data type, 230
- converting to a numeric data type, 198-199
- converting to a string data type, 217
- referencing date value, 161
- referencing datetime value, 164
- referencing time value, 163

CHAR data type, 157

character data types, 156

- ASCII versus Unicode encoding, 159
- Unicode encodings, 159
- VARCHAR, CHAR, and TEXT, 157

characters, trimming from around strings, 201-203

CHARINDEX function, 204

CHECK constraint, 103

- DROP CHECK and DROP CONSTRAINT, 123

clauses, 7, 45, 51, 53

- operators and functions in, 179
- order of clauses in SELECT statement, 8
- order of execution in SELECT statement, 9
- sample query using six main clauses, 54

cloud, databases hosted on, 22

cloud-based storage solutions

requiring SQL queries on, xi

COALESCE function, 86, 235, 297

code examples from this book, xv

column aliases, 68, 271

- with case sensitivity and punctuation, 59
- creating, 57
- versus table aliases, 60
- use with subqueries in SELECT clause, 62

column names, searching for, 311

columns, 11

- adding to a table, 117
- deleting from a table, 118
- displaying column names in SQLite output, 16
- displaying for a table, 117
- filtering by column in WHERE clause, 74

multiple, counting unique combinations, 65

multiple, in GROUP BY, 244

qualifying column names, 59

renaming, 116

renaming with aliases, 44

updating data in, 124

command prompt, 15

- for MySQL, 17
- for Oracle, 18
- for PostgreSQL, 18
- for SQL Server, 19
- for SQLite, 16

comments, 48

COMMIT command, 138

- confirming changes with, 140
- no ROLLBACK after, 141

common table expressions (CTEs), 291-301

- nonrecursive, 291-295
- recursive, 295-301

versus subqueries, 293-295

- advantages of CTEs, 294

comparison operators, 182-189
 BETWEEN, 184
 EXISTS, 185
 IN, 186
 IS NULL, 187
keywords, 183
 LIKE, 187
 symbols, 183
composite indexes, 131
composite key, 105
CONCAT function, 309
concatenating strings
 CONCAT function, 203
 text from multiple fields into
 single field, 308-311
 concatenating text from
 fields in multiple
 rows, 310-311
 concatenating text from
 fields in single row,
 309
|| (concatenation) operator,
 203
conditional statements, 47, 75,
 183
 (see also predicates)
connections (database)
 connecting database tool to
 database, 23
 setting up for Python, 26-29
 setting up for R, 31-34
constants, 144
 (see also literals)
CONSTRAINT keyword, 102
constraints, 101, 106, 120-124
 adding to a table, 122
 deleting before deleting a col-
 umn, 118
 deleting from a table, 123
 displaying for a table, 120
 foreign key, deleting table
 with, 128
 modifying in a table, 122
not allowing NULL values in
 a column, 102
requiring unique values in
 column, 104
restricting values in column,
 103
setting default values in col-
 umn, 103
specifying foreign key, 106
specifying primary key, 105
CONVERT function, 231
correlated subqueries, 62
 performance issues with, 62
COUNT function, 80, 191
HAVING clause referring to,
 84
HAVING COUNT in query,
 306
using with DISTINCT, 64,
 304
counting starting at 1 in SQL, 205
CREATE statements
 for databases, 96
 for indexes, 131
 privileges to execute, 95, 100,
 132, 135
 for tables, 98, 100
 for views, 134
Create, Read, Update, and Delete
 (see CRUD operations)
CROSS JOIN, 271, 272, 275, 281
 syntax, 273
CRUD operations, 6, 91-142
 for databases, 91-97
 for indexes, 129-132
 for tables, 97-129
 using with transaction man-
 agement, 138-142
 for views, 133-137
CSV files, inserting data from into
 a table, 112-115
CTEs (see common table expres-
 sions)

- CUBE keyword, 249
current date or time, getting, 218-220
CURRENT_DATE function, 47, 218
CURRENT_TIME function, 218
CURRENT_TIMESTAMP function, 218
- D**
- data analysis workflow, 24
Data Control Language (DCL), 51
Data Definition Language (DDL), 50
Data Manipulation Language (DML), 50
data models, 10-12, 91
 versus schemas, 93
Data Query Language (DQL), 50
data types, 143-178
 choosing for a column, 145
 of columns, 97
 datetime data, 161-172
 datetime data types, 165-172
 numeric data, 147-154
 decimal data types, 150
 floating point data types, 151
 integer data types, 148-150
 other data, 172-178
 Boolean data types, 173
 external files, 173-178
 string data, 154-161
 character data types, 156-159
 Unicode data types, 159
database drivers, 20
 installing driver for Python, 25
 installing driver for R, 31
database files, 22
- Database Management Systems (DBMS), 3
database objects, 91
database tools, 13, 20-24
 comparisons of, 21
 connecting to a database, 22
databases, 10, 91-97
 about, 1
 creating, 22, 95
 data model, 10
 data model versus schema, 93
 deleting, 96
 displaying name of current, 94
 displaying names of existing, 93
NoSQL, 2
 qualifying tables with database name, 60
querying, 54
schemas, 92
 showing in MySQL, 17
 showing in Oracle, 18
 showing in PostgreSQL, 19
 showing in SQL Server, 20
 showing in SQLite, 16
SQL, 1
 switching to another, 95
DATE data type, 143
DATE keyword, 161
DATEDIFF function, 222
dates, 161
 converting string to date data type, 232
data types (see datetime data types)
date formats, 232
date units in different RDBMSs, 226
date values, 161
datetime format specifiers, 233

filling in missing dates with recursive CTE, 295-298

datetime data, 161-172

DATETIME data type, 166

datetime functions, 43, 161, 218-234

converting string to datetime data type, 230-234

applying date function to string column, 233

using CAST function, 230

using STRING_TO_DATE, TO_DATE, and CONVERT, 231

determining day of week for a date, 228

extracting part of date or time, 226-228

finding difference between two dates, 221

finding difference between two datetimes, 224

finding difference between two times, 222

returning current date or time, 218-220

rounding date to nearest time unit, 229

in SQLite, 171

subtracting date or time interval, 220

DATETIME keyword, 164

datetime values, 161-165

date and time values, 164

date values, 161

datetime data types, 165-172

time values, 162

DB files (see database files)

DCL (Data Control Language), 51

DDL (Data Definition Language), 50

DECIMAL data type, 153

decimals

decimal data types, 150

decimal values, 147

declarative programming, 38

DEFAULT constraint, 103

DELETE statements, 101, 120, 140, 142

operators and functions in, 180

DENSE_RANK function, 252

versus ROW_NUMBER and RANK, 254

derived tables, 69

DESCENDING (DESC) keyword, 86

digits, 154

DISTINCT keyword, 63, 304

using with COUNT, 80

using with COUNT in SELECT clause, 64, 304

distributed data processing

frameworks, SQL-like interfaces, xi

DML (Data Manipulation Language), 50

documents, storing for use in SQL, 173

DOUBLE data type, 152

double precision (floating point data types), 152

DQL (Data Query Language), 50

DROP statements

DROP CHECK and DROP CONSTRAINT, 123

for databases, 96

for indexes, 132

for tables, 101, 128

for views, 137

duplicates

excluding duplicate rows with UNION, 287

finding rows containing duplicate values, 303-306

preserving duplicate rows
with UNION ALL, 287

removing duplicate rows
from result with DISTINCT, 64, 304

E

ESCAPE keyword, 189

escape sequences
only for strings enclosed in single quotes, not dollar signs (\$\$), 156

escape sequences in a string, 155

EXCEPT operator, 289
order of execution, 291

EXISTS operator, 185
JOIN versus, 185
NOT EXISTS, 186, 186

expressions, 47
(see also common table expressions)

extensions for SQL, 38

external files (images, documents, etc.), 173

EXTRACT function, 233

F

FALSE and TRUE values, 172

filepaths to desktop (example), 114

filtering data
alternatives to WHERE clause, 78
using HAVING clause, 83
using WHERE clause, 73-77
filtering on columns, 74
filtering on subqueries, 75-77

FIRST_VALUE function, 255

fixed point numbers, 150

FLOAT data type, 152

floating point data types, 151

floating point values, 147

foreign keys, 11
deleting table with foreign key reference, 128
specifying, 106

frequently asked questions (FAQs), 303-315

concatenating text from multiple fields into single field, 308

from fields in a single row, 308

from fields in multiple rows, 310

finding all tables containing specific column name, 311-313

finding rows containing duplicate values, 303-306
returning all unique combinations, 304
returning only rows with duplicate values, 305-306

selecting rows with max value for another column, 306-308

FROM clause, 8, 46, 66-73, 78
data from multiple tables, using JOIN, 66
table aliases, 67

order of execution in SELECT statement, 9

subqueries within, 69-73
benefits of, 72-73

table aliases defined within, 61

FULL OUTER JOIN, 272, 278

functions, 42, 179, 191-235
aggregate, 191, 193
datetime, 218-234
most common, 180
null, 234

- numeric, 193-199
operators versus, 179
string, 199-218
- G**
- go command (SQL Server), 20
GREATEST function, 193
GROUP BY clause, 8, 46, 78-82, 242-247
aggregating rows into single value or list, 245
collecting of rows, 79
grouping by multiple columns, 243
HAVING clause following, 83
nonaggregate columns in, 192
order of execution in SELECT statement, 9
in query selecting rows with max value, 308
steps to take when using, 82
summarizing of rows in groups, 80
using to create summary table, 243
grouping and summarizing, 237, 242-250
GROUP BY clause, 242
using ROLLUP, CUBE, and GROUPING SETS, 247
CUBE keyword, 249
GROUPING SETS keywords, 249
ROLLUP keyword, 248
GROUPING SETS keywords, 249
GROUP_CONCAT function, 191, 245, 310
GUIs (graphical user interfaces)
in database tools, 13
using database tool GUI to modify table, 119
- H**
- HAVING clause, 8, 83-85
filtering results of GROUP BY, 83
HAVING COUNT in query, 306
operators and functions in, 179
order of execution in SELECT statement, 9
in query finding duplicate values, 305
requirement to follow GROUP BY, 83
used with SELECT, order of execution, 84
using to filter data, 78
WHERE versus, 84
hexadecimal values, 174
Homebrew package manager (Linux and macOS), 15
How do I... (see frequently asked questions)
- I**
- identifiers, 43
double quotes ("") enclosing, 49
naming, 43
IF EXISTS keywords, 128
IF NOT EXISTS keywords, 100
images, 144
in external files, 173
immutability, primary keys, 106
imperative programming, 37
IN operator, 186
NOT IN, 181, 186
indexes, 129-132
book index versus SQL index, 129
creating to speed up queries, 131
deleting, 132

- limiting number of, 131
INNER JOIN, 69, 271, 272, 276-277
in query finding duplicate values, 306
in query selecting rows with max value, 308
INSERT statements, 98, 110, 120 operators and functions in, 180
installation, RDBMSs, 15
INSTR function, 204
INTEGER data type, 98, 143, 146, 171
integers
 dividing by an integer, 190
 integer data types, 148
 integer values, 147
 sample of integer data types, 145
INTERSECT operator, 290
 order of execution, 291
intervals
 INTERVAL data type, 172
 subtracting date or time interval, 220
IS NOT NULL operator, 187
IS NULL operator, 48, 187
- J**
joins, 270-274
 basics of, 274
 CROSS JOIN, 281
 in duplicate rows query, 305
 EXISTS versus JOIN, 185
 FULL OUTER JOIN, 278
 INNER JOIN, 274, 276-277
 JOIN clause, 270
 breakdown of, 271
 join condition, 271
 JOIN defaulting to INNER JOIN, 69, 271
 join types, 271
- JOIN versus UNION, 284
JOIN ... ON ... syntax, 67, 276
LEFT JOIN, 277
in query selecting rows with max value, 307
rewriting subqueries with, 63, 72
RIGHT JOIN, 278
self join, 282-284
syntax to join together tables, 272
USING and NATURAL JOIN
 shortcuts, 279-281
using instead of correlated subquery, 62
using JOIN clause in FROM clause, 66
- K**
keywords, 7, 42, 51
 case insensitive in SQL, 42
operators as, 180
- L**
LAG function, 258, 297
LAST_VALUE function, 255
LEAD function, 259
LEAST function, 193
LEFT JOIN, 272, 277, 296
LEFT OUTER JOIN, 278
LENGTH function, 199, 217
 LEN function in SQL Server, 217
LIKE operator, 184, 187
 limited regular expressions
 support in SQL Server, 216
NOT LIKE, 188
searching for text in strings, 204
LIMIT clause, 78, 88

LISTAGG function, 191, 242, 310
literals, 144
localhost, 28, 33
logical operators, 181
LOWER function, 200
LTRIM and RTRIM functions,
202

M

math functions, 194
math operators, 189
MAX function, 191, 193, 245
Microsoft SQL Server (see SQL Server)
MIN function, 191, 193
missing data
 from CSV file, interpretations by RDBMS, 114
 replacing missing values with another value, 103
MongoDB, 3
MySQL
 MySQL Workbench, 21
 schemas and databases, 93
 writing SQL code with, 17

N

naming conventions
 column aliases, 59
 identifiers, 43
NATURAL JOIN, 273
 caution with, 281
 using to replace INNER JOIN, 280
noncorrelated subqueries, 62
nonrecursive CTEs, 291
NoSQL, 2
NOT NULL constraint, 102, 106
NOT operator, 181
NTH_VALUE function, 256
NULL values, 48
 allowing in a column, 102

from CSV file, interpretations by RDBMS, 114
IS NULL and IS NOT NULL operators, 187
NOT IN versus NOT EXISTS operator, 186
NULL and NOT NULL column constraints, 102
null functions, 234
NULL literal, 145
 replaced by COALESCE function, 86
numeric data, 147-154
 comparing numeric column to string column, 198
 decimal data types, 150
 floating point data types, 151
 integer data types, 148
 numeric values, 147
numeric functions, 43, 147, 193-199
CAST, converting to a numeric data type, 198-199
generating random numbers, 196
math, 194
rounding and truncating numbers, 197
NVARCHAR data type, 160
VARCHAR versus, 160

O

object relational mappers (ORMs), 30
ON clause, 67
 JOIN ... ON ... syntax, multiple conditions in, 277
operators in, 179
 replacing with USING in joins, 279
operators, 179-191
 comparison, 182-189

versus functions, 179
in join conditions, 271
logical, 181
math, 189
most common, 180
versus predicates, 182
union, 285-291
using to combine predicates, 75

OR operator, 75, 179, 181, 182

Oracle
compared to other RDBMSs, 4
Oracle SQL Developer, 21
procedural language SQL (PL/SQL), 38
schemas and users, 93
writing SQL code with, 17

Oracle Database (see Oracle)

ORDER BY clause, 8, 85-88, 252
alphabetic sorting in ascending order, 86
inability to use in subqueries, 88
order of execution in SELECT statement, 9
sorting by columns and expressions not in SELECT, 87
sorting by descending order, 87
in UNION queries, 288

order of execution
SELECT and HAVING clauses, 84
SELECT statement, 9
union operators, 291

OVER keyword, 252
OVER ... PARTITION BY ... syntax, 252, 308

P
pandas dataframes, 29

PARTITION BY keywords, 252, 262
passwords
to connect to database tool, 23
for Python connection to database, 28
for R connection to database, 33

pattern matching
LIKE operator, 187
(see also regular expressions)

performance
correlated subqueries, issues with, 62
optimizing code, 77
speeding up queries with indexes, 131

PIVOT operation, 263-265
CASE statement as alternative to, 264

pivoting and unpivoting, 237, 263-267
breaking up values of column into multiple columns, 263-265
listing values of multiple columns in single column, 265-267

PL/SQL (procedural language SQL), 38

POSITION function, 204

POSIX syntax, regular expressions, 213, 215

PostgreSQL
compared to other RDBMSs, 4
default database, postgres, 97
pgAdmin, 21
writing SQL code with, 18

precision, 151
double precision floating point data types, 152

- single precision floating point
 data types, 152
- predicates, 47, 75
 combining multiple, using
 operators, 75
- operators versus, 182
- primary keys, 11, 105-106
 best practices, 106
- foreign keys referring to, 106
- programming languages
 comparison of SQL to other,
 37
- writing SQL code within, 13,
 24-35
- Python, 25-31
- R, 31-35
- punctuation in column aliases, 59
- Python, 37
 connecting to a database,
 25-29
- connecting to RDBMS and
 writing SQL code within,
 14
- SQLAlchemy package, 30
- use of SQL with, xiii
- writing SQL code within,
 29-31
- Q**
- qualifying column names, 59, 68
- qualifying table names, 60
- queries, 6-10, 53-89
 basics, 53
- combining with UNION, 287
- creating view to save results
 of, 135
- inserting results of into a
 table, 110
- speeding up by creating an
 index, 131
- updating rows with results of,
 126
- querying concepts, advanced,
 237-267
- CASE statements, 238-242
- grouping and summarizing,
 242-250
- pivoting and unpivoting,
 263-267
- window functions, 250-262
- quotes in SQL, 48
- R**
- R language
 connecting to a database,
 31-34
- connecting to RDBMS and
 writing SQL code within,
 14
- use of SQL with, xiii
- writing SQL code within, 34
- random number generator, 196
- ranges
 RANGE BETWEEN versus
 ROWS BETWEEN, 262
- testing if value falls in with
 BETWEEN, 184
- RANK function, 252
 versus ROW_NUMBER and
 DENSE_RANK, 254
- RDBMS software, 13
- RDBMSs
 ANSI SQL versus RDBMS-specific SQL, 39
- comparisons of, 4
- deciding which RDBMS to
 use, 14
- MySQL, 17
- Oracle, 17
- PostgreSQL, 18
- SQL Server, 19
- SQLite, 15
- defined, 3
- using SQLAlchemy with, 30

variations in SQL syntax, 3, 315
writing SQL code with, 14
REAL data type, 171
recursive CTEs, 291, 295-301
 filling in missing rows in data sequence, 295-298
 returning all parents of child rows, 298-301
RECURSIVE keyword, 292
REGEXP function (MySQL), 210
REGEXP_REPLACE function, 208
 in Oracle, 212
 in PostgreSQL, 216
regular expressions, 184, 209-217
 important notes about, 210
 in MySQL, 210
 in Oracle, 211
 in PostgreSQL, 214
 in SQL Server, 216
 using in Oracle to extract substrings, 207
 using in Oracle to search for substrings, 206
Relational Database Management Systems (see RDBMSs)
relational databases, 1
relationship, defined, 11
remote servers, databases on, 22
renaming
 of columns, 116
 of tables, 116
REPLACE function, 207
result sets, 54
RIGHT JOIN, 272, 278
RIGHT OUTER JOIN, 278
ROLLBACK command, 138
 undoing changes with, 141
ROLLUP keyword, 248
rounding numbers, 197
rows
 aggregating into single value or list, 245
 collecting in GROUP BY, 79
 deleting from a table, 120
 displaying for a table, 119
 inserting into a table, 120
 ranking in a table, 252
 ROWS BETWEEN versus RANGE BETWEEN, 262
 updating values in, 125
 updating with query results, 126
 ROWS BETWEEN clause, 259
 ROWS BETWEEN UNBOUNDED clause, 261
ROW_NUMBER function, 252, 252
 versus RANK and DENSE_RANK, 254
 using in subquery to return multiple ranks in each group, 257
 within a window function, 251
RTRIM and LTRIM functions, 202

S

scale, 151
schemas, 91
 data models versus, 93
 qualifying tables with schema name, 60
in SQL databases, 2
varying definitions in RDBMSs, 93
SELECT clause, 8, 46, 55-66
aliasing columns, 57
 aliases with case sensitivity and punctuation, 59
COUNT and DISTINCT in, 64

DISTINCT in, 63
operators and functions in,
179
order of execution in SELECT
statement, 9
qualifying columns, 59
qualifying tables, 60
selecting subqueries, 61-63
SELECT statements, 7, 45, 140
clauses in, 7, 45
combining results of two or
more with UNION, 285
operators and functions in,
180
self join, 274, 282-284
semi-join, 185
separators, 246, 311
SIMILAR TO, limited regular
expression support in Post-
greSQL, 215
single precision (floating point
data types), 152
SMALLINT data type, 146
sorting
 ascending and descending
 order, 86
 by columns and expressions
 not in SELECT list, 87
SQL
 about, 1
 comparison to other pro-
 gramming languages, 37
 integral role in data land-
 scape, xi
 language summary, 51
 searching for syntax online, 4
 sublanguages, 50
SQL Server, 5
 compared to other RDBMSs,
 4
 default database, master, 97
 SQL Server Management Stu-
 dio, 22
 writing SQL code with, 19
SQL statements, 6
SQLAlchemy package (Python),
30
SQLite
 compared to other RDBMSs,
 4
 database stored in outside
 files, 94
DB Browser for SQLite, 21
fastest RDBMS to set up, 14
star schema, 92
Start Fridays with grandma's
homemade oatmeal
 (mnemonic for clauses), 8
START TRANSACTION, 138
statements, 6, 45
 clauses in, 45
string functions, 43, 154, 199-218
 changing case of a string, 200
 concatenating strings, 203
 converting to a string data
 type, 217
 deleting text from a string,
 208
 extracting portion of a string,
 206
 finding length of a string, 199
 replacing text in a string, 207
 searching for text in strings,
 203
 trimming unwanted charac-
 ters around strings,
 201-203
 using regular expressions,
 209-217
strings, 154-161
 casting to decimal to compare
 with a number, 199
 character data types, 156
 comparing string column to
 numeric column, 198

- converting to datetime data type, 230-234
- single quotes ("") enclosing, 49
- string to date functions, 162
- string to datetime functions, 164
- string to time functions, 162
- string values, 154, 156
- alternative to using single quotes, 155
- escape sequences for, 155
- Unicode data types, 159
- STRING_AGG function, 191, 242, 310
- STR_TO_DATE function, 231
- sublanguages (SQL), 50
- subqueries
- common table expressions versus, 293-295
 - correlated versus noncorrelated, 62
 - correlated, performance issues with, 62
 - filtering data on, 255, 256
 - within FROM clause, 69-73
 - benefits of using, 72-73
 - order of execution for sample query, 69-71
 - naming temporarily with aliases, 44
 - no ORDER BY clauses in, 88
 - within SELECT clause, 61
 - using instead of a view, 134
 - versus views, 135
 - within WHERE clause, 75
 - benefits of, 76
 - versus WITH clause, 71
 - wrapping DISTINCT in and using COUNT on, 65
- SUBSTR or SUBSTRING function, 206
- substrings, finding in strings, 205
- SUM function, 191, 251
- using with ROWS BETWEEN UNBOUNDED clause, 261
- Sweaty feet will give horrible odors (mnemonic for clauses), 8
- T**
- table aliases, 60
- column aliases versus, 60
- using in FROM clause, 67
- table names, displaying, 311
- tables, 1, 10
- combining with UNION and eliminating duplicate rows, 286
- creating simple table, 98
- creating table not already existing, 100
- creating table with automatically generated field, 108
- creating table with constraints, 101-105
- not allowing NULL values in column, 102
- requiring unique values in column, 104
- restricting values in column with CHECK, 103
- setting default values in column, 103
- creating table with primary and foreign keys, 105-108
- specifying foreign key, 106
- specifying primary key, 105
- creating to populate new database, 22
- derived tables, 69
- displaying names of existing tables, 100

inserting query results into, 110-112
inserting text file data into, 112-115
modifying, 115-129
 adding a column, 117
 adding a constraint, 122
 adding rows, 120
 changing data type for a column, 146
 deleting a table, 128
 deleting column from a table, 118
 deleting constraints, 123
 deleting rows, 120
 deleting table with foreign key reference, 128
 displaying columns, 117
 displaying constraints, 120
 displaying rows, 119
 modifying a constraint, 122
 renaming a column or table, 115-117
 updating column of data, 124
 updating rows of data, 125
 updating rows with query results, 126
qualifying table names, 60
renaming with aliases, 44
showing in MySQL, 17
showing in Oracle, 18
showing in PostgreSQL, 19
showing in SQL Server, 20
showing in SQLite, 16
SQL requirements for, 97
views versus, 54, 133
working with multiple, 269-301

TCL (Transaction Control Language), 51

terminal window, writing SQL code in, 13, 15
(see also command prompt)

text
 concatenating from multiple fields into single field, 308-311
 finding all tables containing specific column name, 311-313
 inserting data from text file into a table, 112-115
 replacing in a string, 207
 searching for in a string, 204

TEXT data type, 157, 171

time, 161
(see also datetime data)
 converting string to time data type, 230
data types (see datetime data types)
time formats, 232
time units in different RDBMSs, 227
time values, 162

TIME keyword, 163

TIMESTAMP data type, 166

TIMESTAMP keyword, 164

TINYINT data type, 146

TO_DATE function, 231

trailing spaces
 excluding in string length, 200
 removing with RTRIM, 202

Transaction Control Language (TCL), 51

transaction management, 138-142
 benefits of, 138
 double-checking changes before COMMIT, 139
no ROLLBACK after COMMIT, 141

- undoing changes with ROLLBACK, 141
- transaction, defined, 138
- TRIM function, 200, 201
- removing leading or trailing characters, 202
- TRUE and FALSE values, 172
- truncating existing table, 101, 120
- truncating numbers, 197
- type affinities (SQLite), 98
- U**
- Unicode data types, 159
- ASCII versus Unicode encoding, 159
- UNION ALL, alternative to UNPIVOT, 266
- union operators, 284-291
- EXCEPT, 289
- INTERSECT, 290
- JOIN versus UNION, 284
- order of execution, 291
- requirement for matching data types, 287
- UNION, 285
- with more than two tables, 289
- with other clauses, 288
- UNION ALL, 287, 296
- UNIQUE constraint, 104, 106
- unique rows (see DISTINCT keyword)
- UNPIVOT operation, 265-267
- UNION ALL as alternative to, 266
- UPDATE statements, 125
- inability to undo results, 314
- operators and functions in, 180
- updating a column of data, 124
- updating rows of data, 125
- updating table where ID matches another table, 313
- updating values based on a query, 127
- UPPER function, 179, 200
- usernames
- to connect to database tool, 23
- for Python connection to database, 28
- for R connection to database, 33
- USING clause, replacing ON in joins, 273, 279
- UTF (Unicode Transformation Format), 159
- V**
- VARBINARY data type (SQL Server), 176
- VARCHAR data type, 98, 143, 157
- versus NVARCHAR, 160
- VARCHAR2 in Oracle, 110
- vi text editor, 19
- views, 133-137
- creating to restrict access to tables, 133
- creating to save query results, 135
- deleting, 137
- displaying existing views, 136
- querying, 54
- subqueries versus, 135
- updating, 137
- W**
- web page for this book, xvi
- WHERE clause, 8, 46, 73-78
- in CROSS JOINs, 282
- DELETE statement using, 140
- filtering on columns, 74, 89
- filtering on subqueries, 75

versus HAVING, 84
omitting in DELETE statement, 120
operators and functions in, 179
order of execution in SQL code, 9
in self join, 284
in UPDATE statement, 126
whitespace
removing spaces around a string, 201
spaces in column alias names, 59
in SQL, 49, 51
window functions, 237, 250-262
aggregate functions versus, 250
breakdown of, 252
calculating moving average, 259
calculating running total, 261
ranking rows in a table, 252
returning first two values in each group, 257
returning first value in each group, 255
returning prior row value, 258
returning second value in each group, 256
window, defined, 252
WITH clause
in common table expressions, 291
RDBMSs supporting, 71
subqueries versus, 71
working code, writing then optimizing, 77

Y

YEAR function, 179, 227

About the Author

Alice Zhao is a data scientist who is passionate about teaching and making complex things easy to understand. She has taught numerous courses in SQL, Python, and R as a senior data scientist at Metis and as a cofounder of Best Fit Analytics. Her highly rated technical tutorials on YouTube are known for being practical, entertaining, and visually engaging.

She writes about analytics and pop culture on her blog, *A Dash of Data*. Her work has been featured in *Huffington Post*, *Thrillist*, and *Working Mother*. She has spoken at a variety of conferences including Strata in New York City and ODSC in San Francisco on topics ranging from natural language processing to data visualization. She has her MS in analytics and BS in electrical engineering, both from Northwestern University.

Colophon

The animal on the cover of *SQL Pocket Guide* is an Alpine salamander (*salamandra atra*). Most commonly found in ravines high up in the Alps (upwards of 1,000m), the Alpine salamander stands out for its unusual ability to handle cold weather. The shiny black creatures prefer shady, moist places and the cracks and gaps in stone walls. It feeds on worms, spiders, snails, and small insect larvae.

Unlike other salamanders, the Alpine salamander gives birth to fully formed juveniles. A pregnancy lasts two years, but at even higher altitudes (1,400–1,700m), it can last up to three years. The species is generally protected throughout the Alps, but climate change has more recently impacted their preferred habitat of rocky, not-too-dry landscapes.

Many of the animals on O'Reilly covers are endangered; all of them are important to the world.

The cover illustration is by Karen Montgomery, based on a black and white engraving from *Lydekker's Royal Natural History*. The cover fonts are Gilroy Semibold and Guardian Sans. The text font is Adobe Minion Pro; the heading font is Adobe Myriad Condensed; and the code font is Dalton Maag's Ubuntu Mono.



O'REILLY®

There's much more where this came from.

Experience books, videos, live online training courses, and more from O'Reilly and our 200+ partners—all in one place.

Learn more at oreilly.com/online-learning

